

# Lecture 17: 3D Vision

# Reminder: A4

A4 due **Today, Wednesday, November 13, 11:59pm**

A4 covers:

- PyTorch autograd
- Residual networks
- Recurrent neural networks
- Attention
- Feature visualization
- Style transfer
- Adversarial examples

# Recall: Course Structure

We are here!



- First half: Fundamentals
  - Details of how to implement and train different types of networks
  - Fully-connected networks, convolutional networks, recurrent networks
  - How to train and debug, very detailed
- Second half: Applications and “Researchy” topics
  - Object detection, image segmentation, 3D vision, videos
  - Attention, Transformers
  - Vision and Language
  - Generative models: GANs, VAEs, etc
  - Less detailed: provide overview and references, but skip some details

# Last Time: Predicting 2D Shapes of Objects

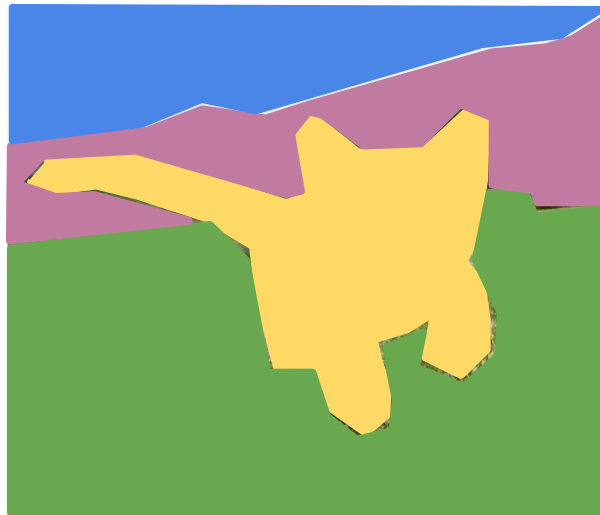
## Classification



CAT

No spatial extent

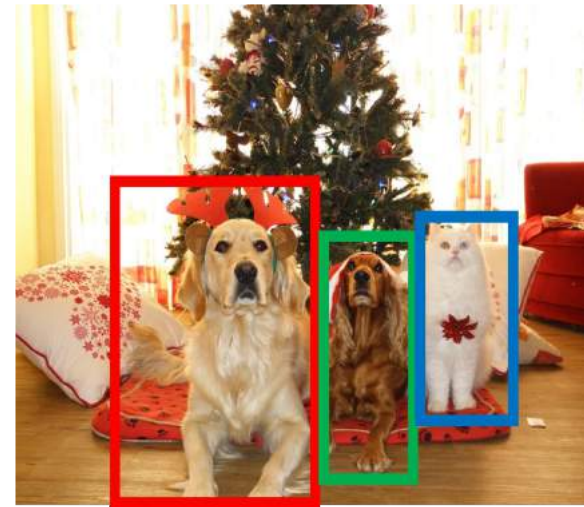
## Semantic Segmentation



GRASS, CAT, TREE,  
SKY

No objects, just pixels

## Object Detection



DOG, DOG, CAT

Multiple Objects

## Instance Segmentation

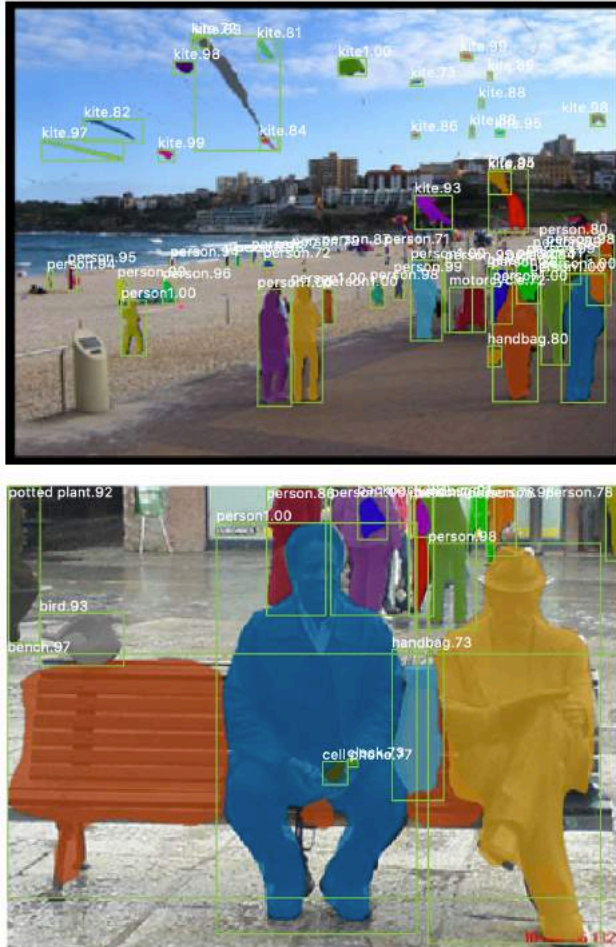


DOG, DOG, CAT

[This image is CC0 public domain](#)

# Today: Predicting 3D Shapes of Objects

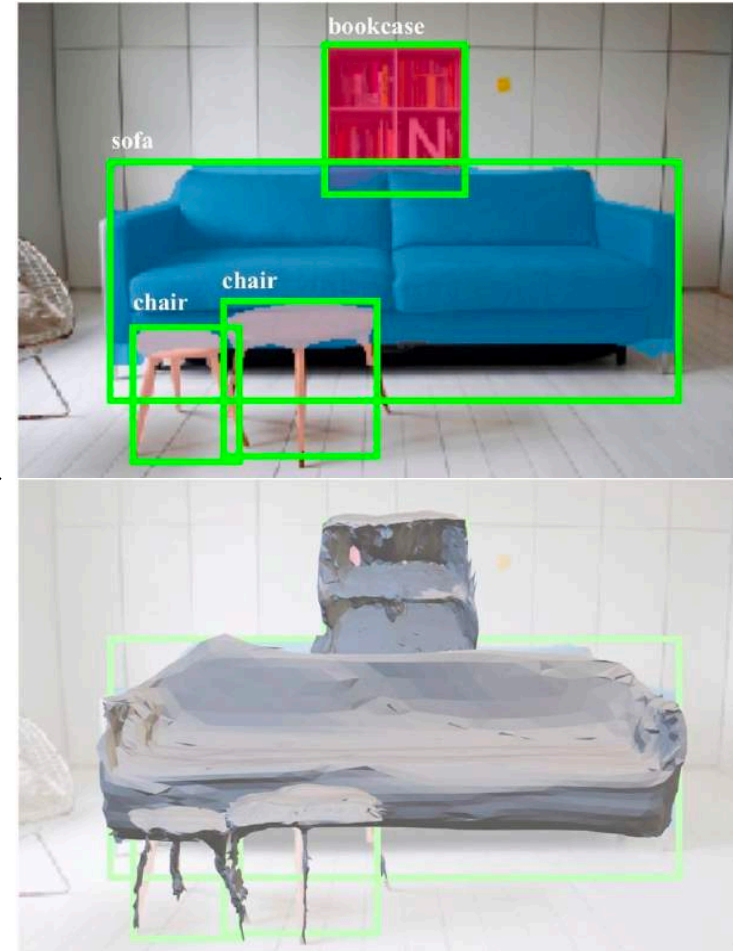
Mask R-CNN:  
2D Image -> 2D shapes



He, Gkioxari, Dollár, and Girshick, "Mask R-CNN", ICCV 2017

## Mesh R-CNN:

2D Image -> **3D** shapes



Gkioxari, Malik, and Johnson,  
“Mesh R-CNN”, ICCV 2019

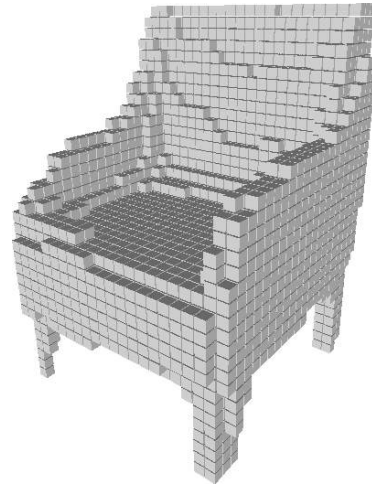
# Focus on Two Problems today

Predicting 3D Shapes  
from single image

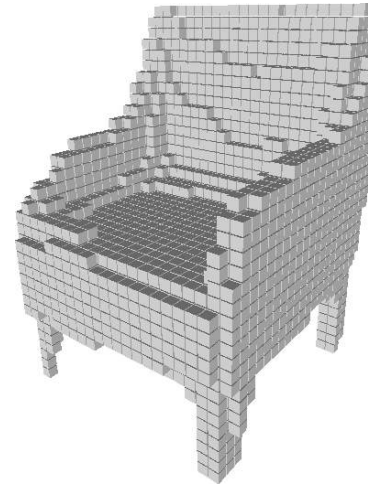
Processing 3D  
input data



Input Image



3D Shape



3D Shape



Chair

# Many more topics in 3D Vision!

Computing correspondences

Multi-view stereo

Structure from Motion

Simultaneous Localization and Mapping (SLAM)

Self-supervised learning

View Synthesis

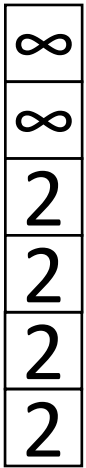
Differentiable graphics

3D Sensors

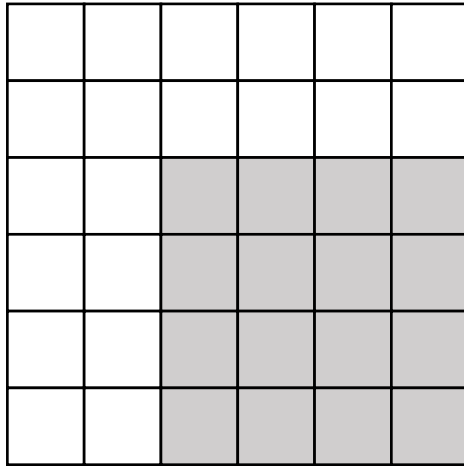
Many non-Deep Learning methods alive and well in 3D!



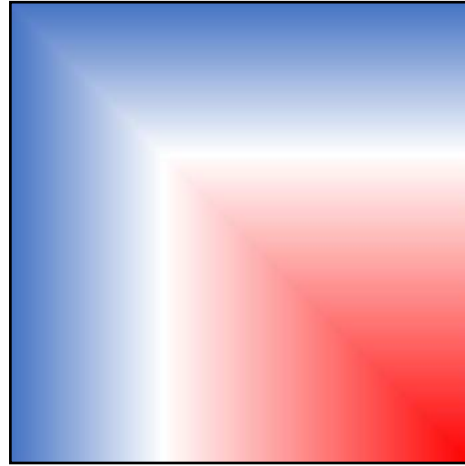
# 3D Shape Representations



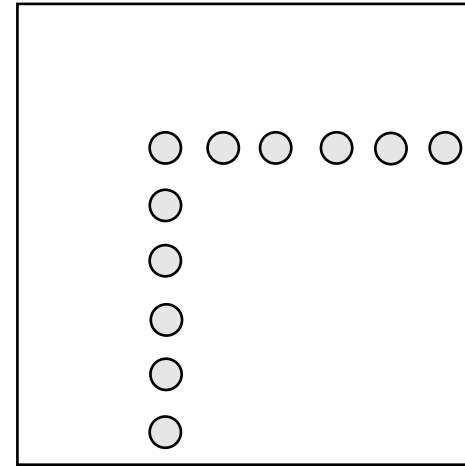
Depth  
Map



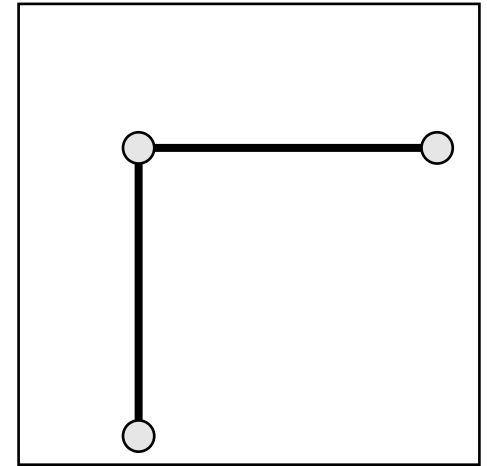
Voxel  
Grid



Implicit  
Surface



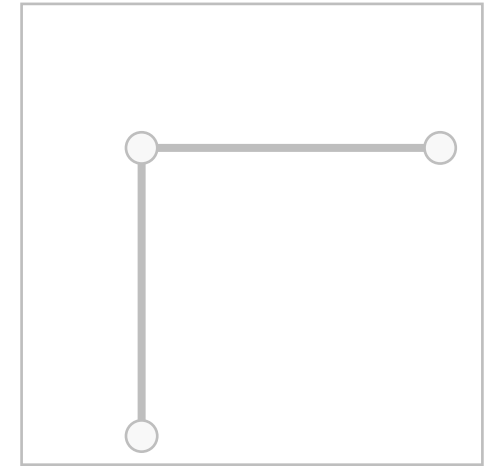
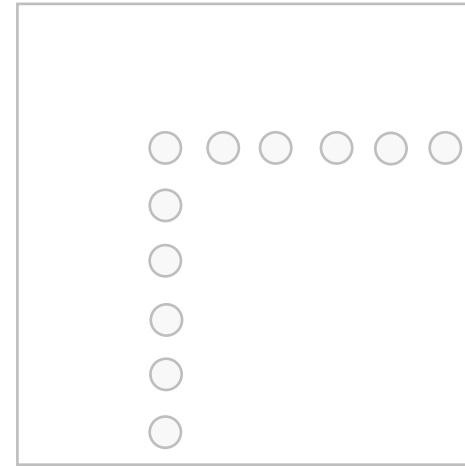
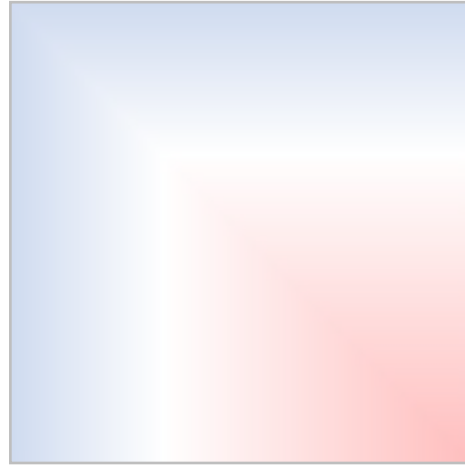
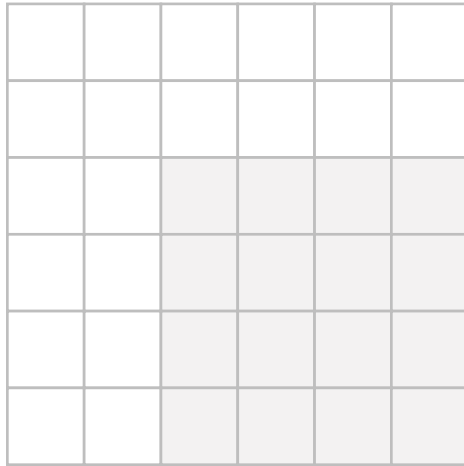
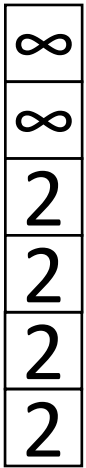
Pointcloud



Mesh



# 3D Shape Representations



Depth  
Map

Voxel  
Grid

Implicit  
Surface

Pointcloud

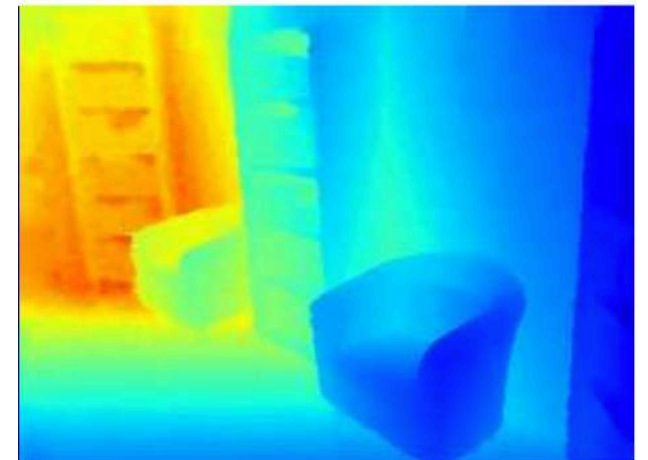
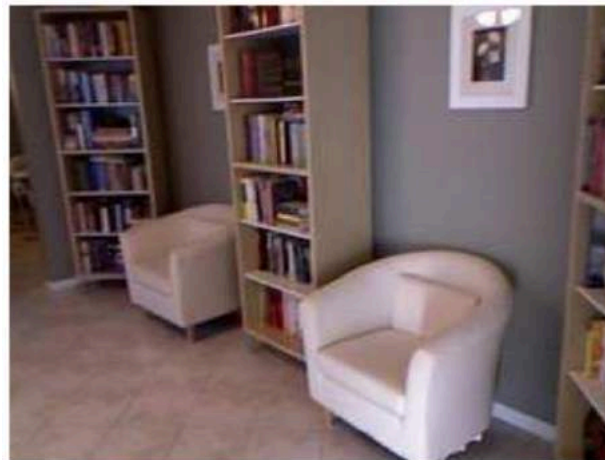
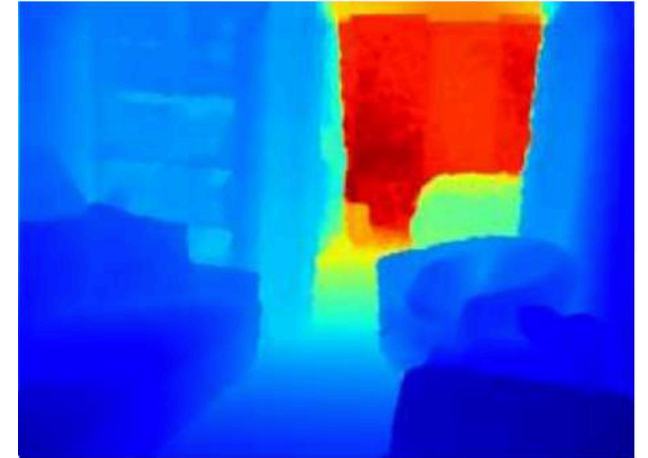
Mesh

# 3D Shape Representations: Depth Map

For each pixel, **depth map** gives distance from the camera to the object in the world at that pixel

RGB image + Depth image  
= RGB-D Image (2.5D)

This type of data can be recorded directly for some types of 3D sensors (e.g. Microsoft Kinect)

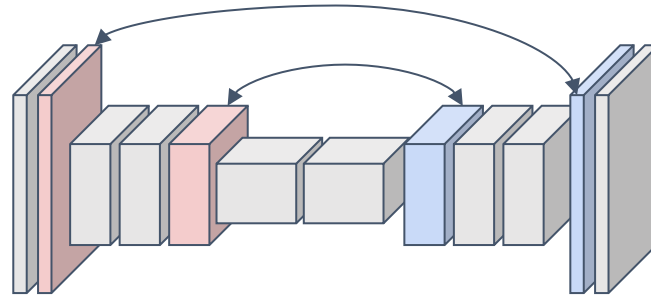


RGB Image:  $3 \times H \times W$     Depth Map:  $H \times W$

# Predicting Depth Maps

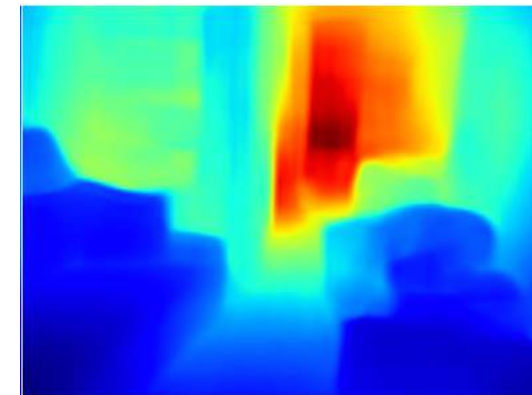
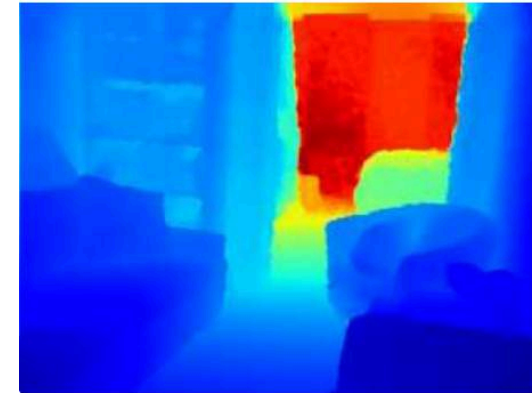


**RGB Input Image:**  
 $3 \times H \times W$



**Fully Convolutional  
network**

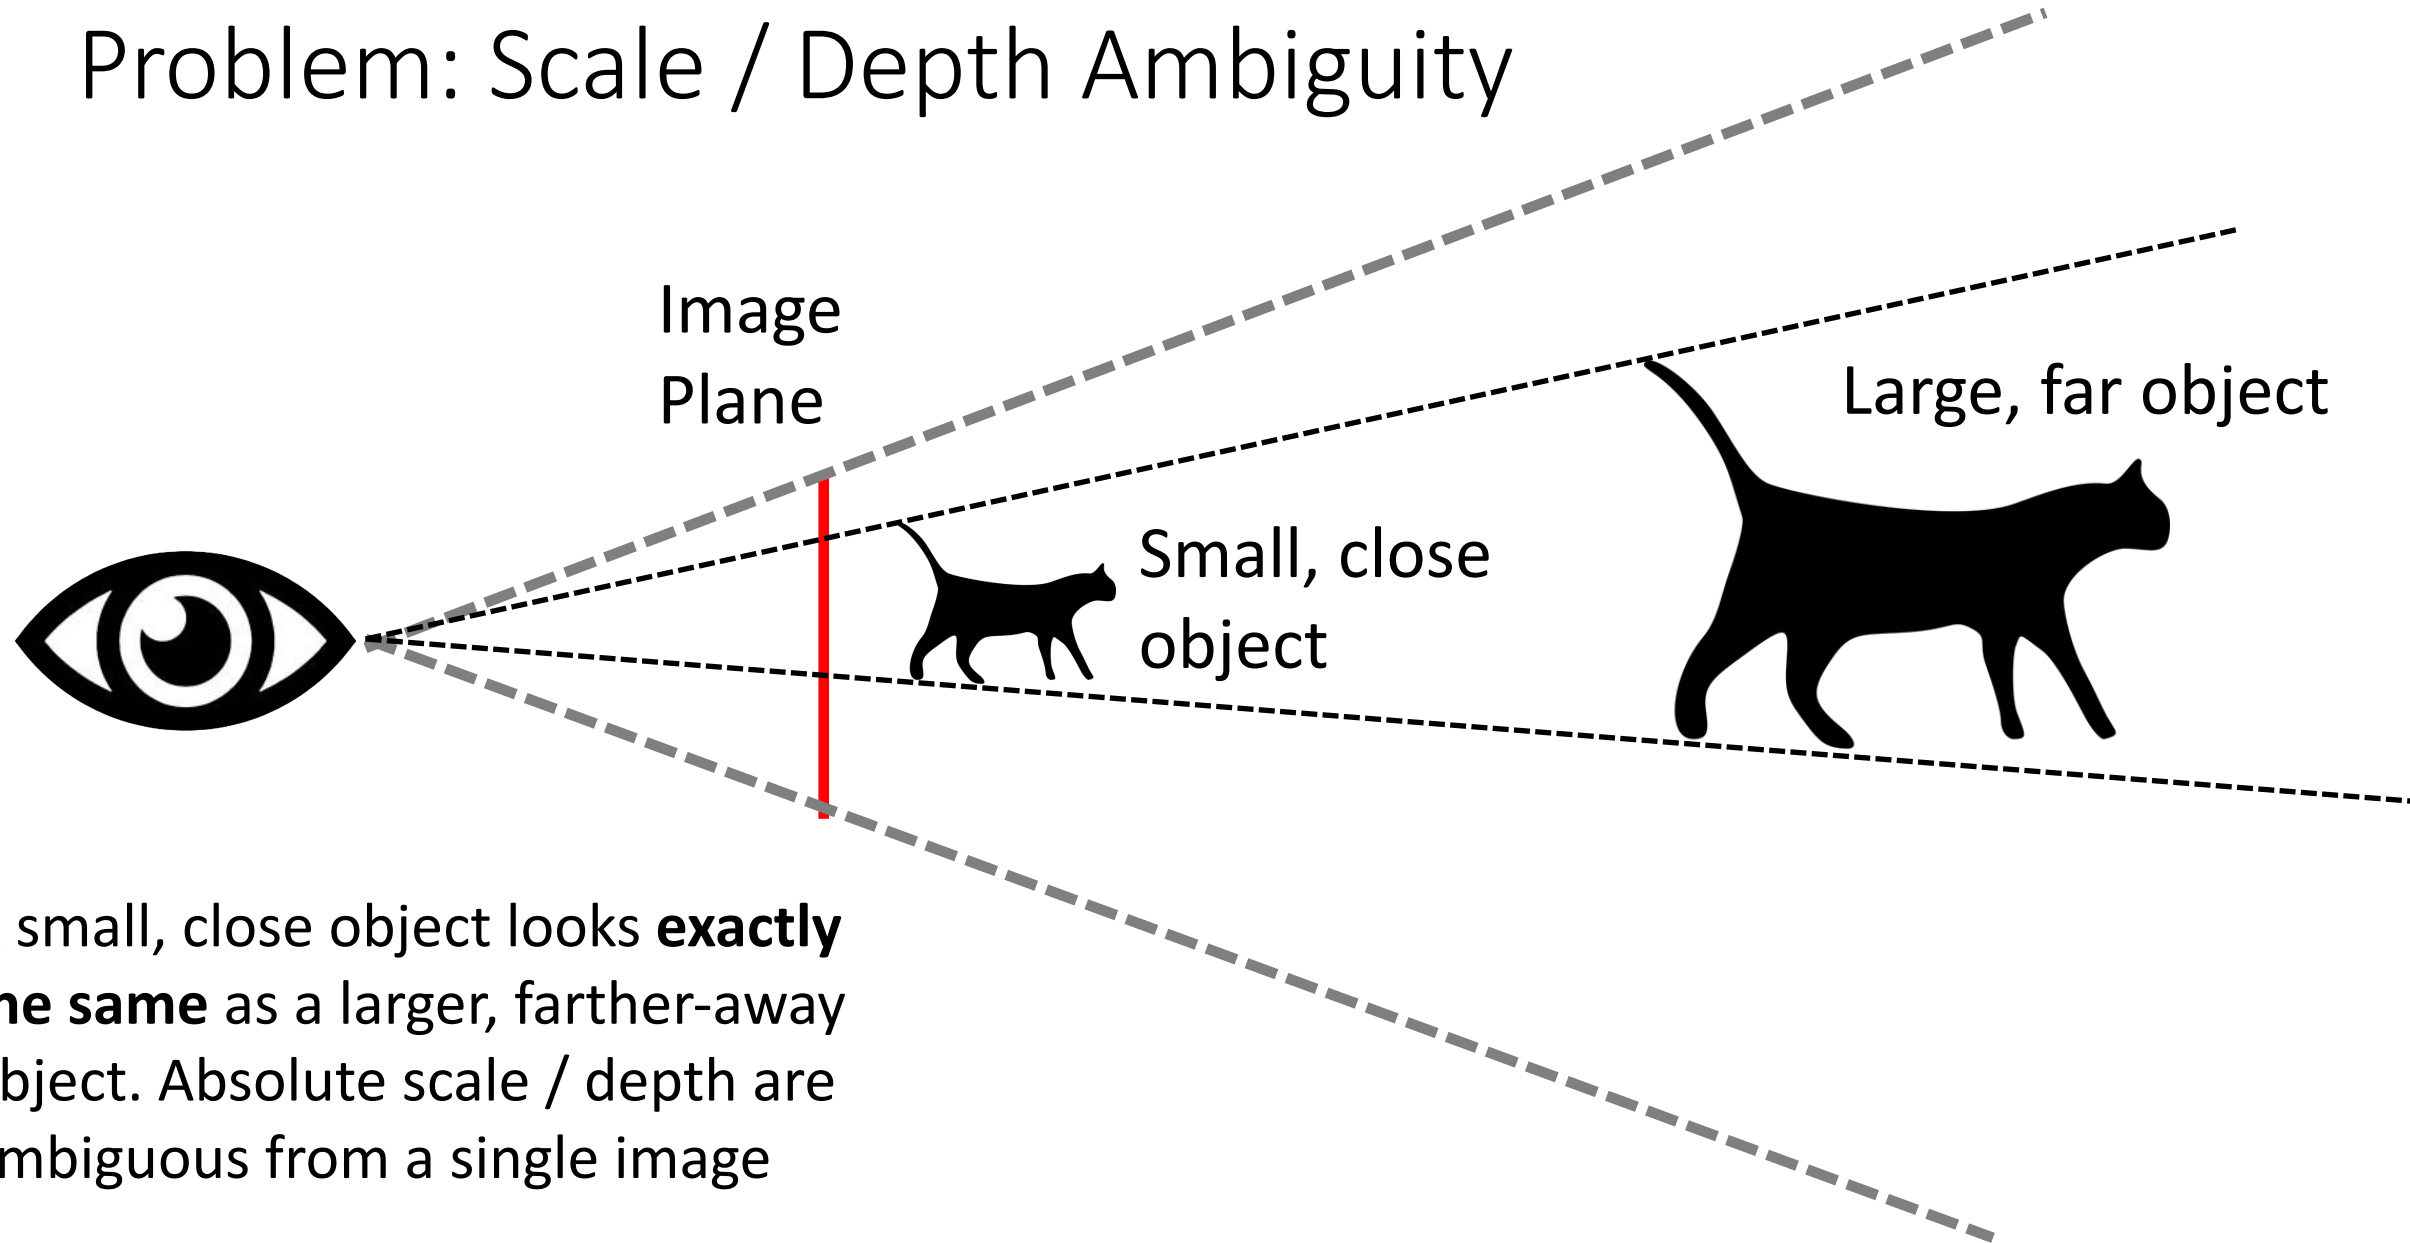
**Predicted Depth Image:**  
 $1 \times H \times W$



**Predicted Depth Image:**  
 $1 \times H \times W$

**Per-Pixel Loss  
(L2 Distance)**

# Problem: Scale / Depth Ambiguity



A small, close object looks **exactly the same** as a larger, farther-away object. Absolute scale / depth are ambiguous from a single image

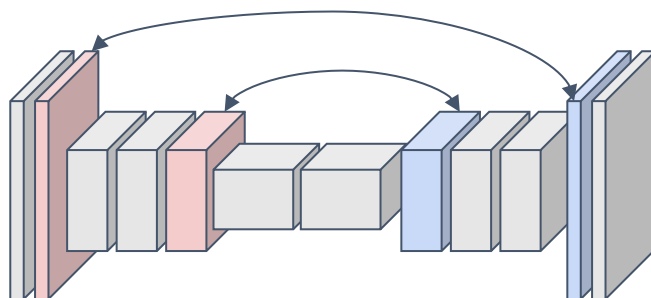
# Predicting Depth Maps

## Scale invariant loss

$$\begin{aligned} D(y, y^*) &= \frac{1}{2n^2} \sum_{i,j} ((\log y_i - \log y_j) - (\log y_i^* - \log y_j^*))^2 \\ &= \frac{1}{n} \sum_i d_i^2 - \frac{1}{n^2} \sum_{i,j} d_i d_j = \frac{1}{n} \sum_i d_i^2 - \frac{1}{n^2} \left( \sum_i d_i \right)^2 \end{aligned}$$

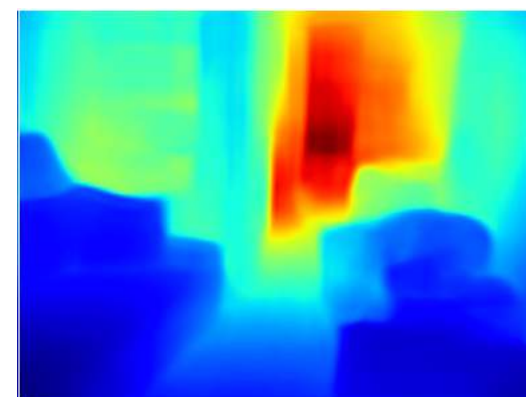
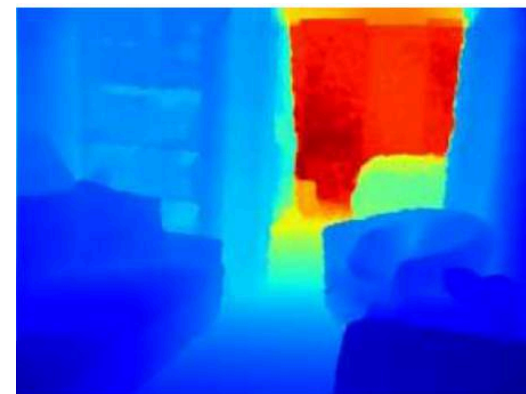


**RGB Input Image:**  
3 x H x W



**Fully Convolutional network**

**Predicted Depth Image:**  
1 x H x W



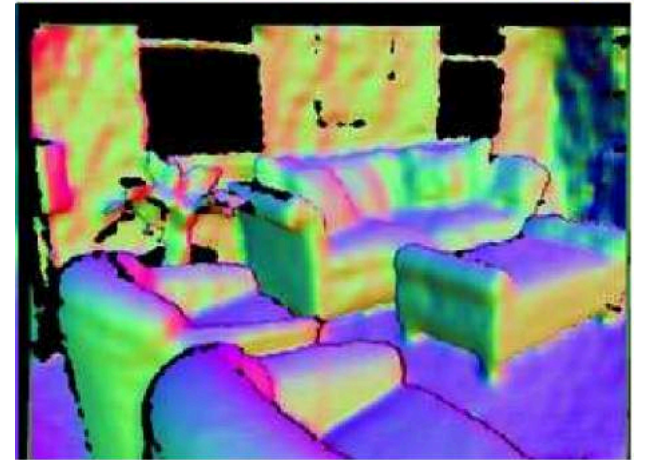
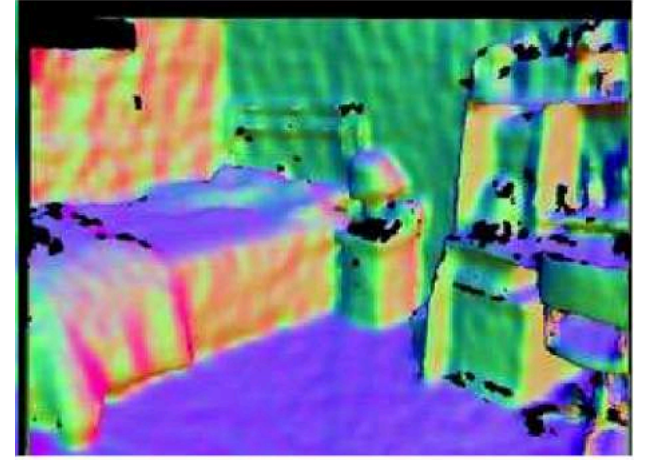
**Predicted Depth Image:**  
1 x H x W

**Per-Pixel Loss**  
(Scale invariant)



# 3D Shape Representations: Surface Normals

For each pixel, **surface normals** give a vector giving the normal vector to the object in the world for that pixel



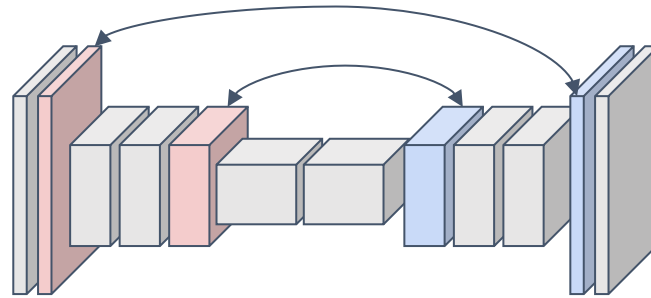
RGB Image:  $3 \times H \times W$

Normals:  $3 \times H \times W$

# Predicting Normals

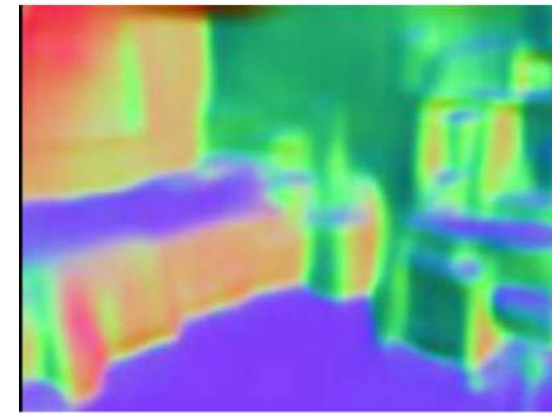
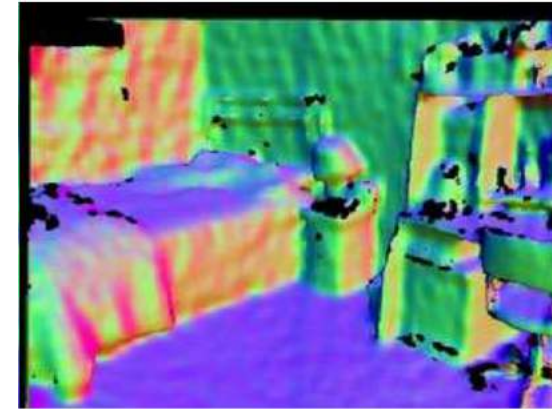


**RGB Input Image:**  
 $3 \times H \times W$



**Fully Convolutional network**

**Ground-truth Normals:**  
 $3 \times H \times W$



**Predicted Normals:**  
 $3 \times H \times W$

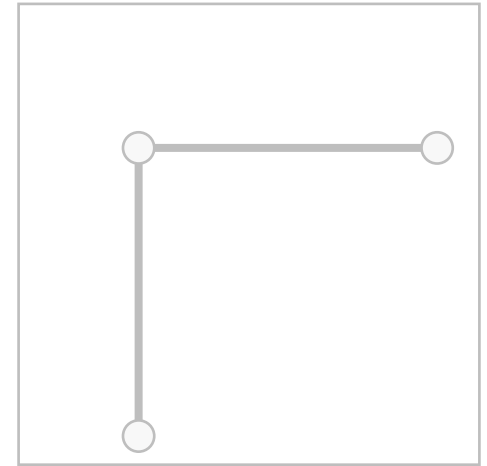
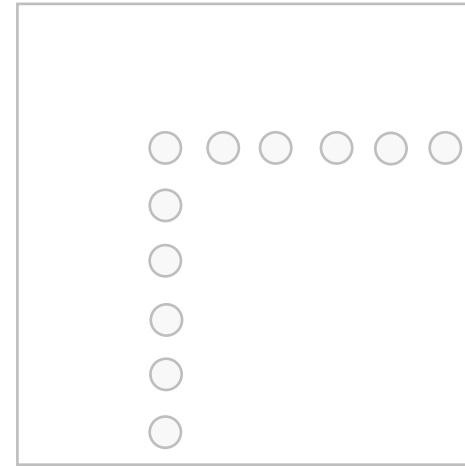
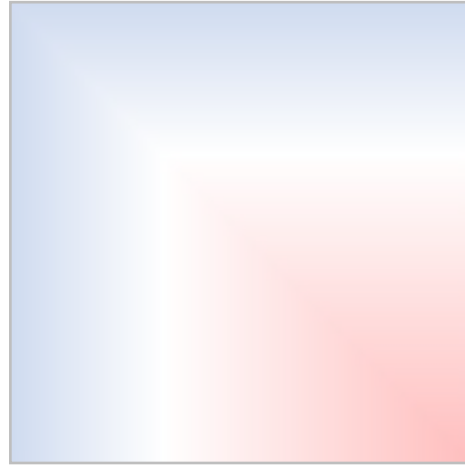
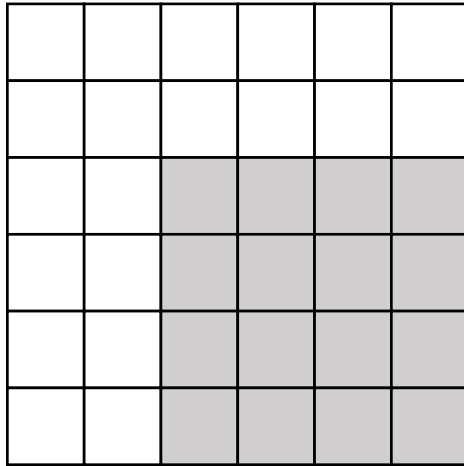
**Per-Pixel Loss:**  
 $(x \cdot y) / (|x| |y|)$

Recall:

$$\begin{aligned} & x \cdot y \\ &= |x| |y| \cos \theta \end{aligned}$$



# 3D Shape Representations



Depth  
Map

Voxel  
Grid

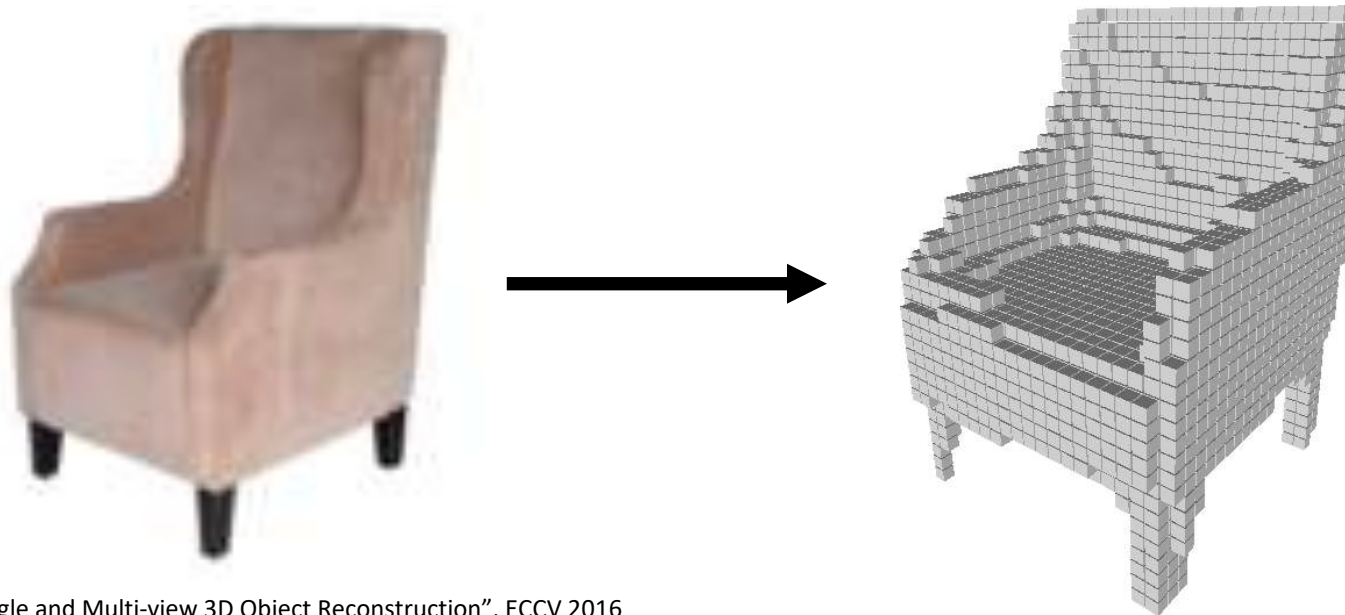
Implicit  
Surface

Pointcloud

Mesh

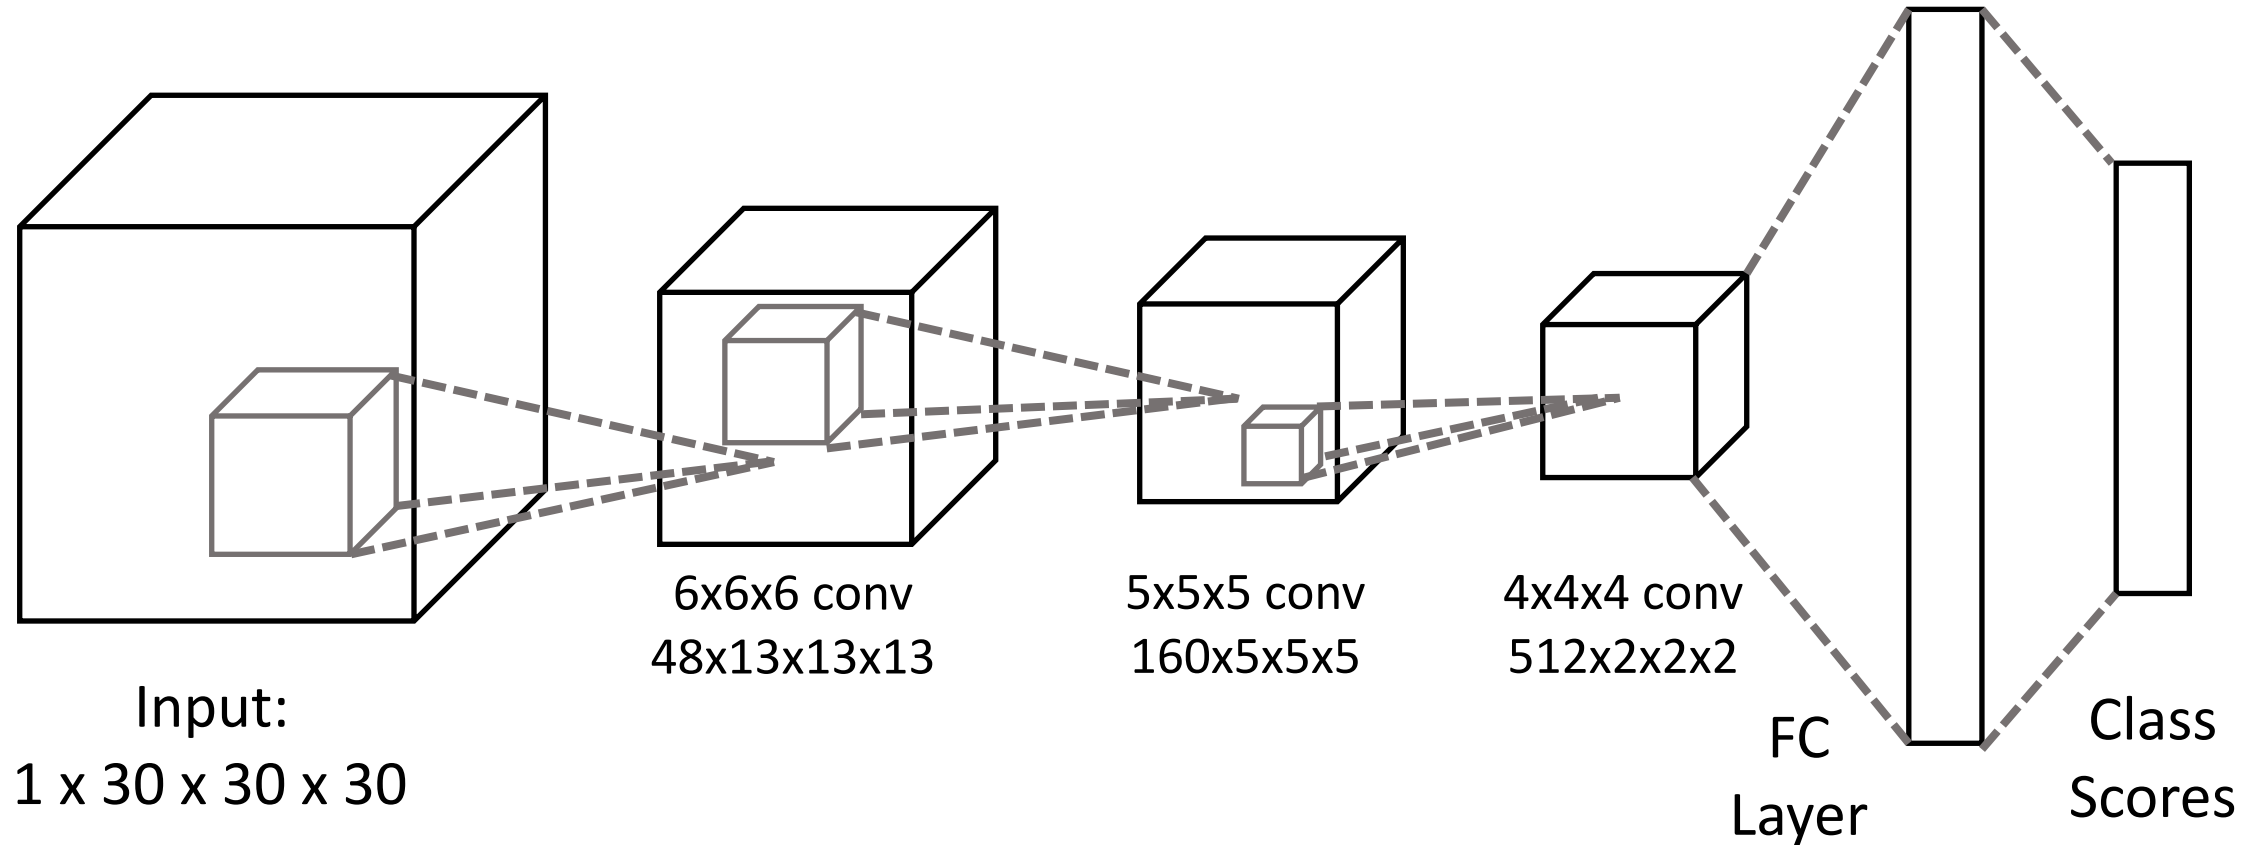
# 3D Shape Representations: Voxels

- Represent a shape with a  $V \times V \times V$  grid of occupancies
- Just like segmentation masks in Mask R-CNN, but in 3D!
- (+) Conceptually simple: just a 3D grid!
- (-) Need high spatial resolution to capture fine structures
- (-) Scaling to high resolutions is nontrivial!



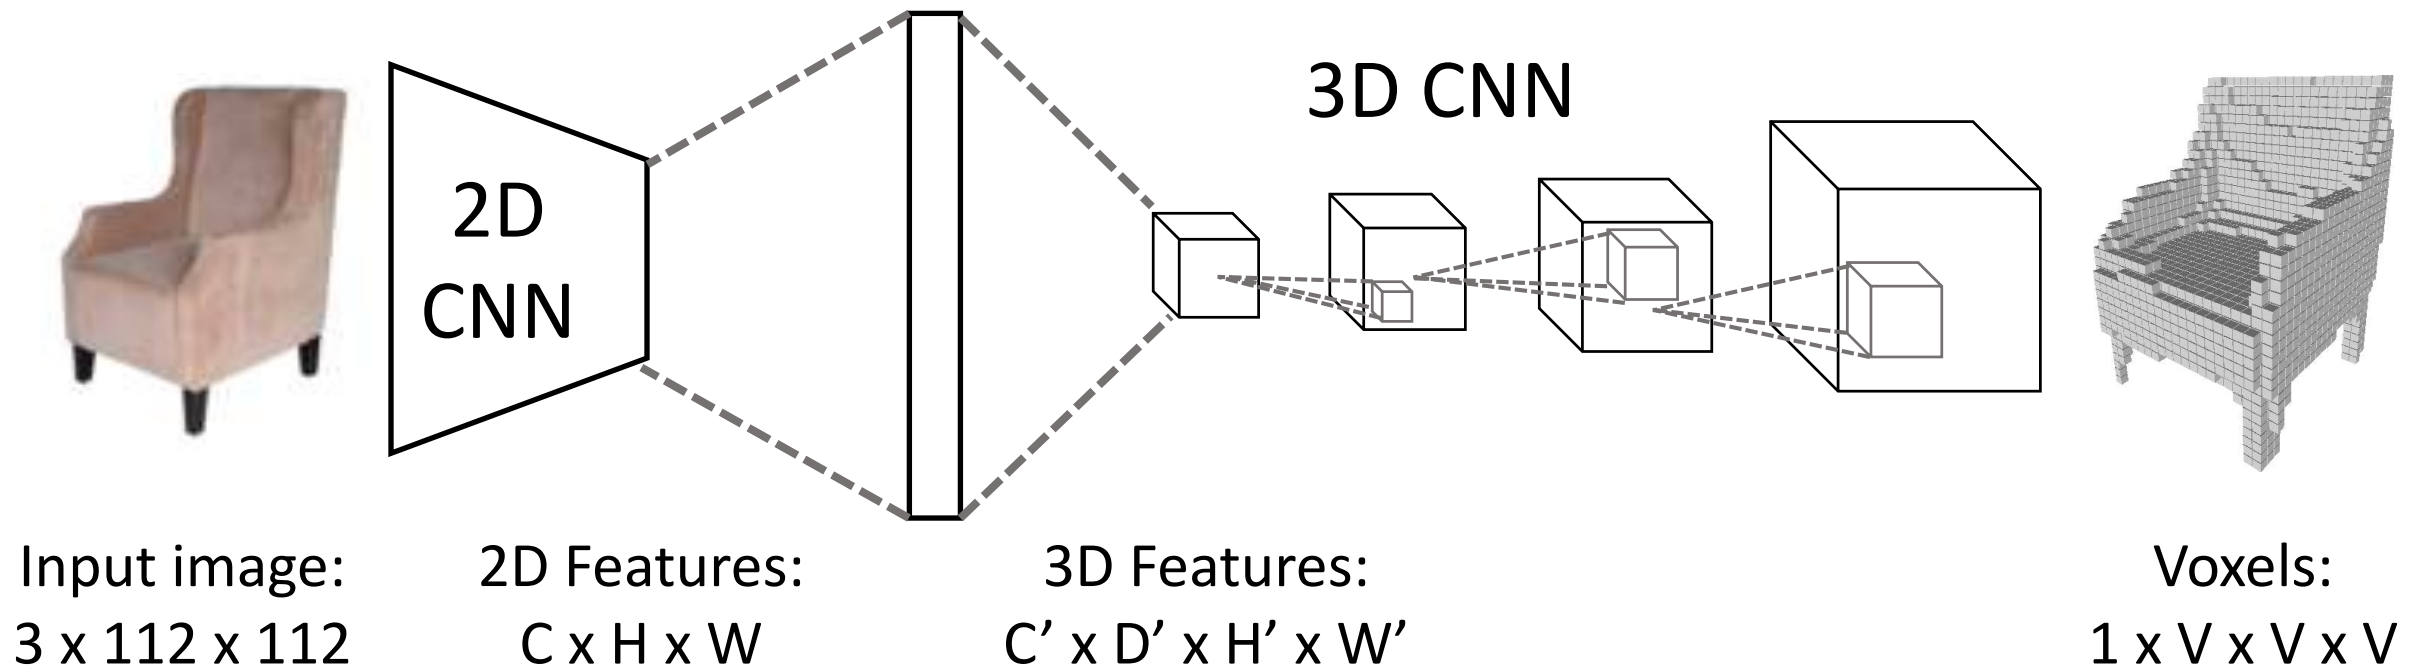
Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016

# Processing Voxel Inputs: 3D Convolution



Train with classification loss

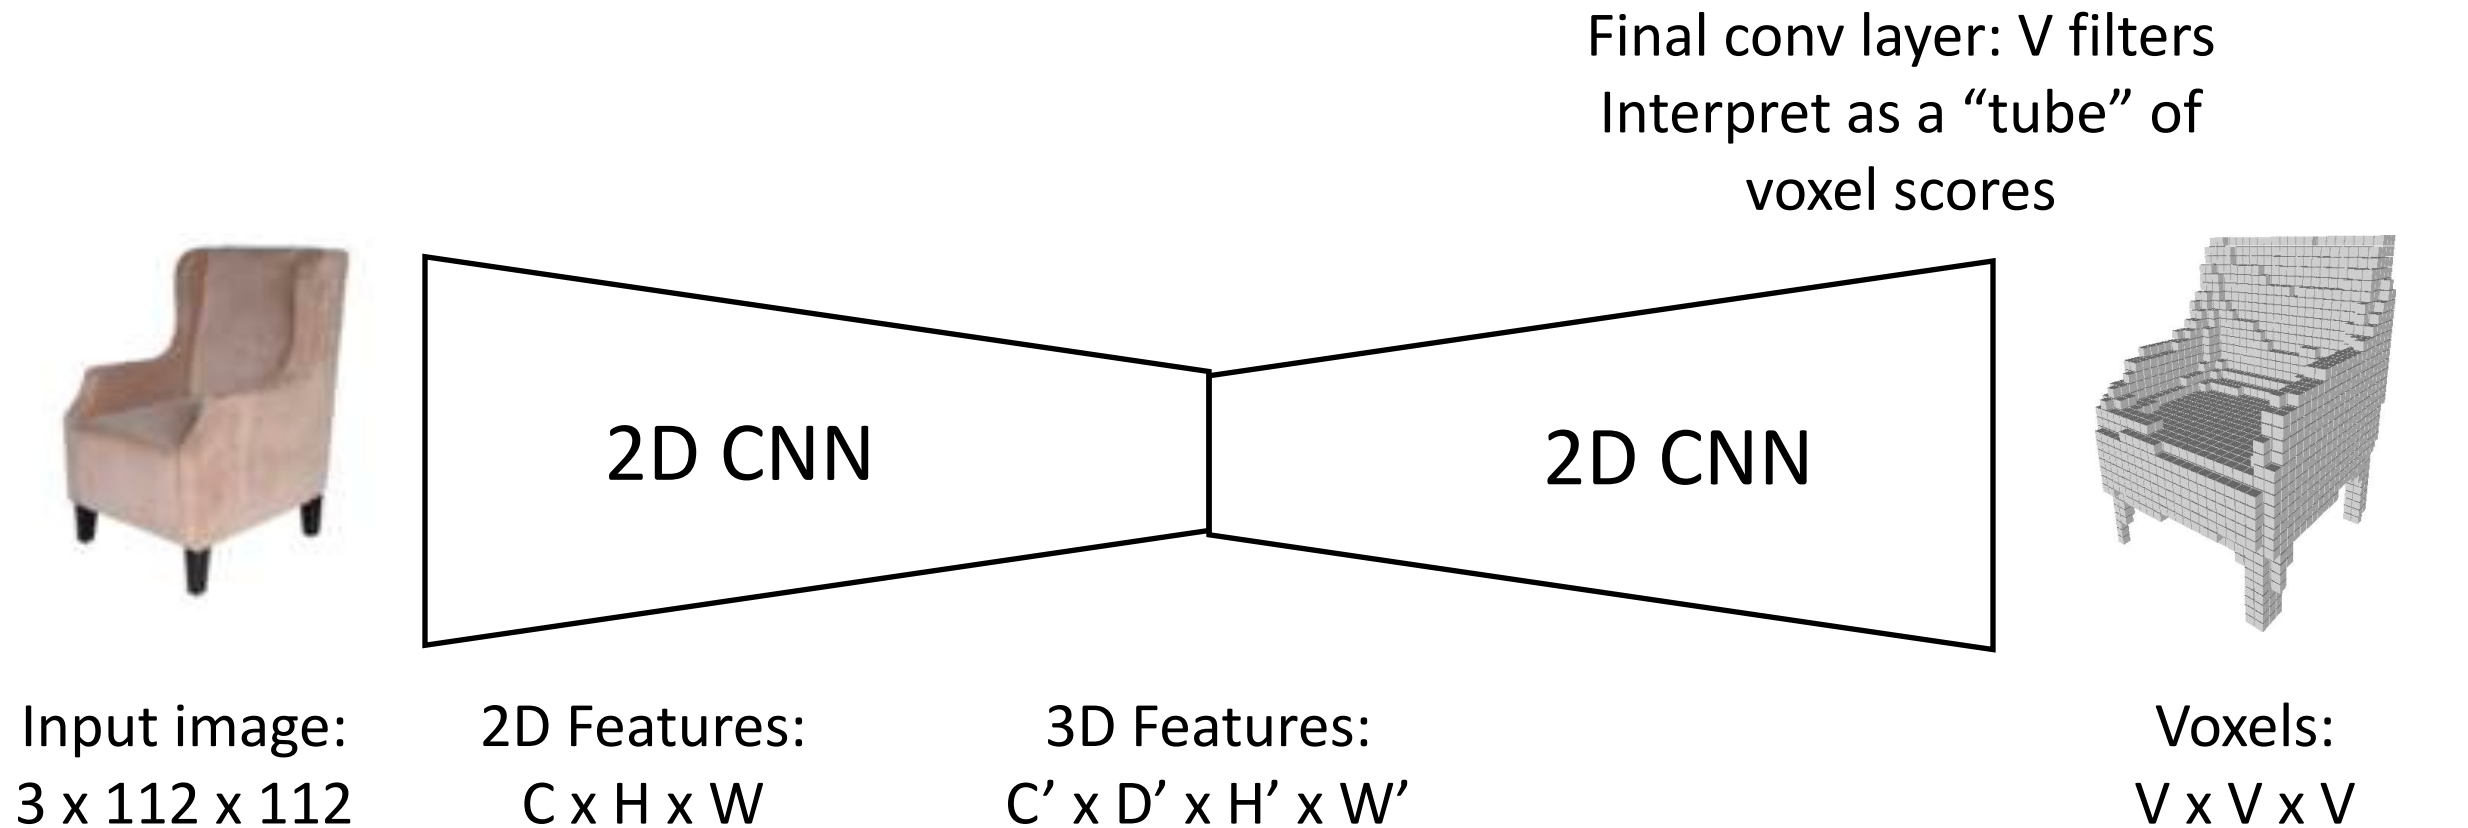
# Generating Voxel Shapes: 3D Convolution



Train with per-voxel cross-entropy loss

Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016

# Generating Voxel Shapes: "Voxel Tubes"

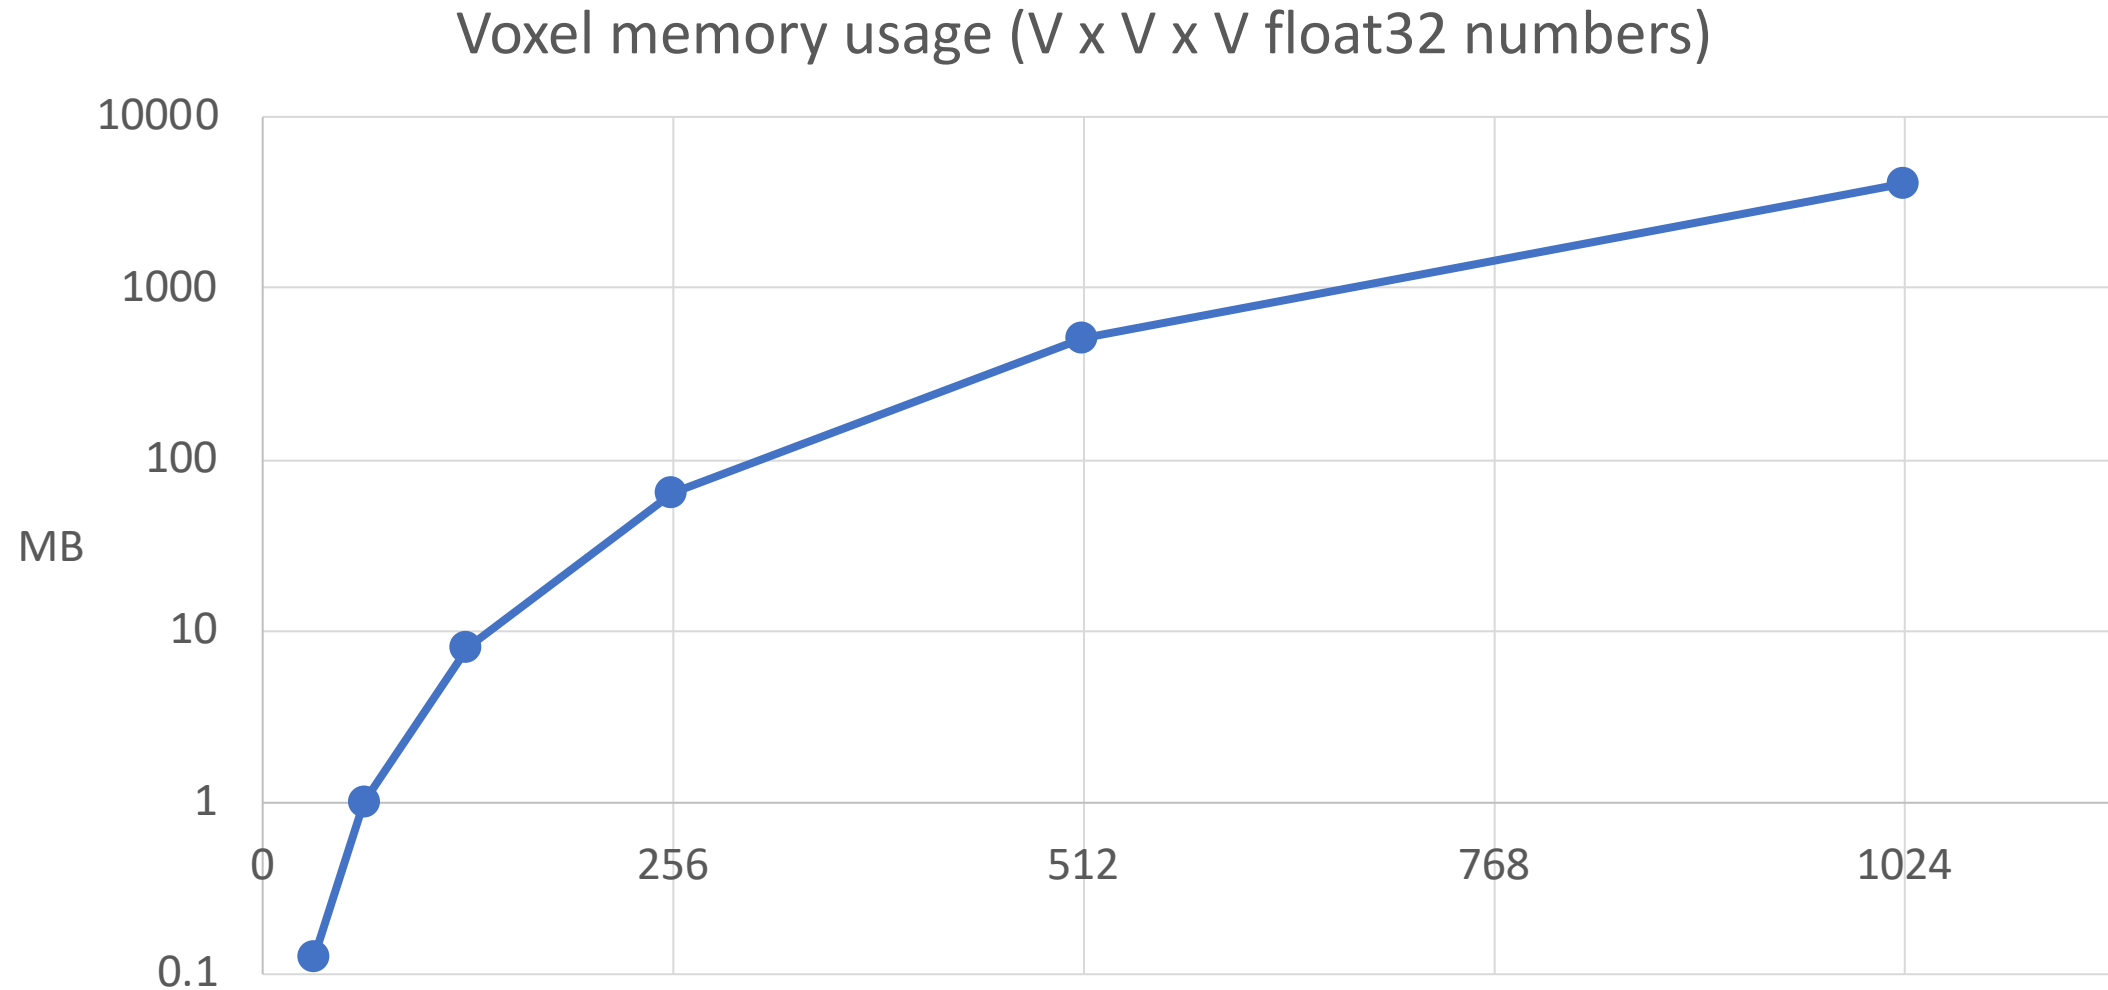


Train with per-voxel cross-entropy loss

Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016

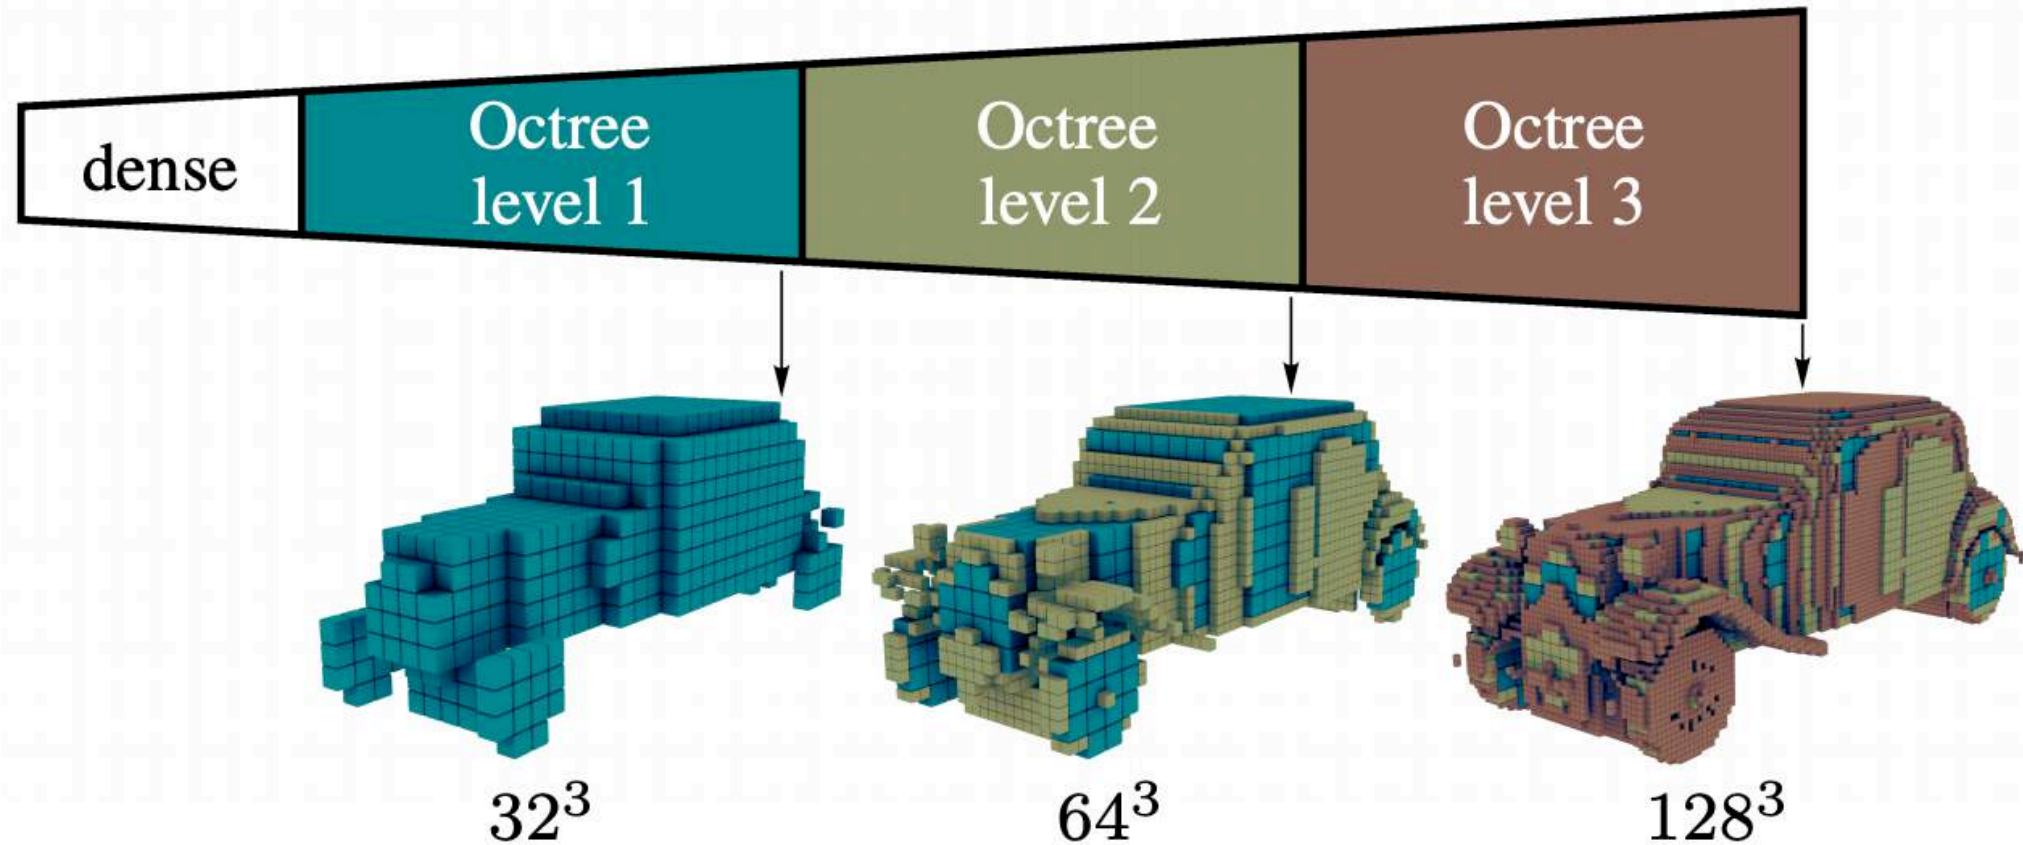
# Voxel Problems: Memory Usage

Storing  $1024^3$  voxel grid  
takes 4GB of memory!



# Scaling Voxels: Oct-Trees

Use voxel grids with heterogenous resolution!



Tatarchenko et al, "Octree Generating Networks: Efficient Convolutional Architectures for High-resolution 3D Outputs", ICCV 2017

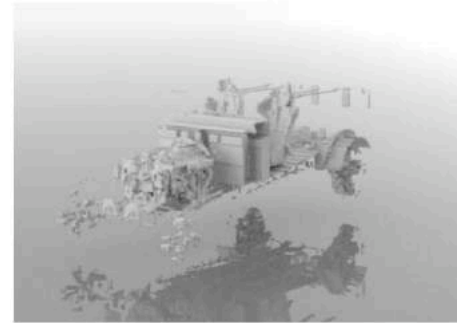
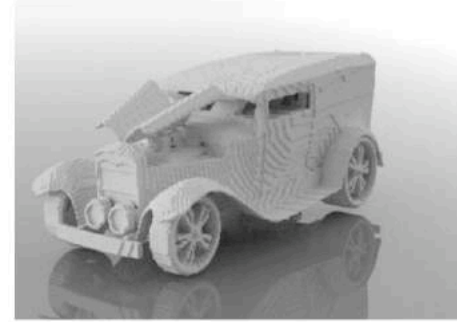


# Scaling Voxels: Nested Shape Layers

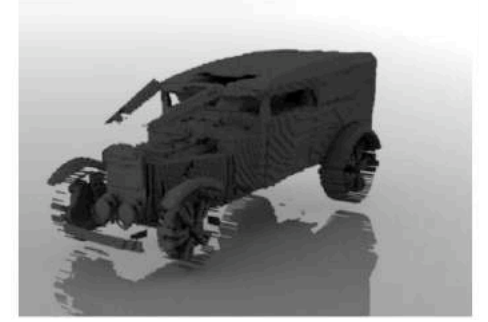
Predict shape as a composition of positive and negative spaces



= + +



-



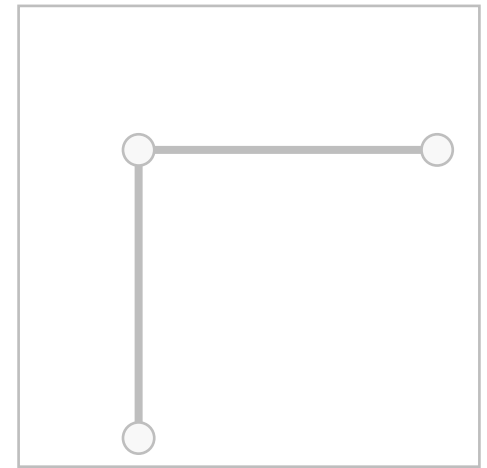
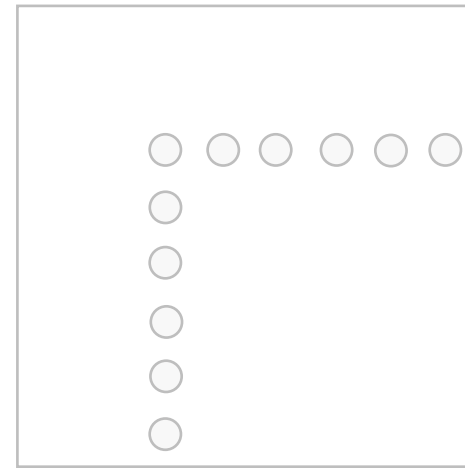
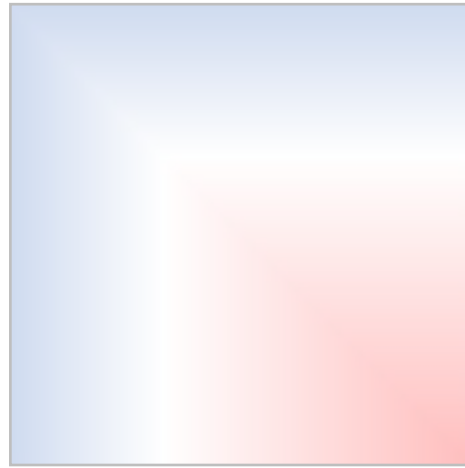
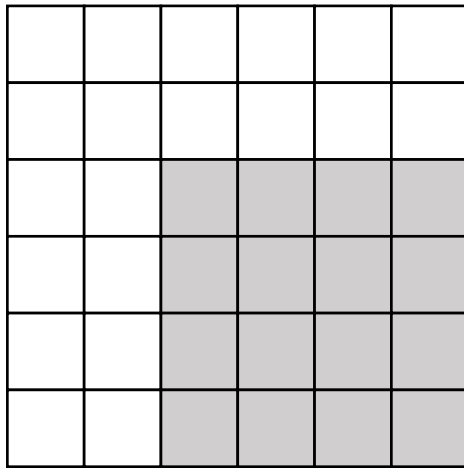
-



Richter and Roth, "Matryoshka Networks: Predicting 3D Geometry via Nested Shape Layers", CVPR 2018

[Doll image](#) is licensed under [CC-BY 2.0](#)

# 3D Shape Representations



Depth  
Map

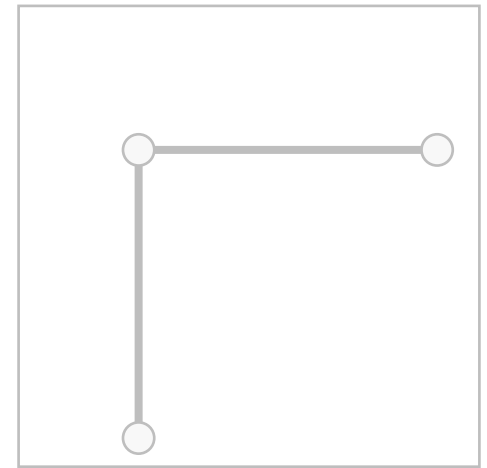
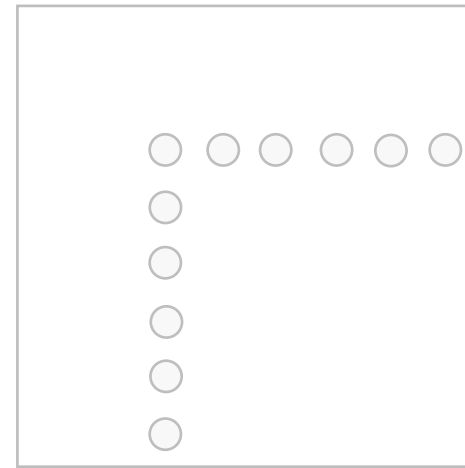
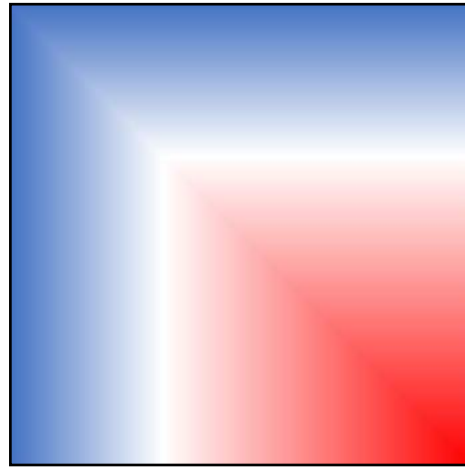
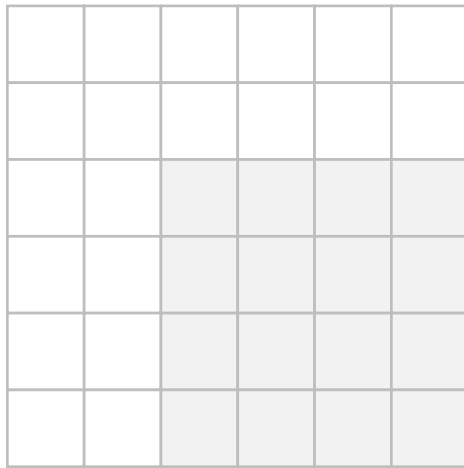
Voxel  
Grid

Implicit  
Surface

Pointcloud

Mesh

# 3D Shape Representations



Depth  
Map

Voxel  
Grid

Implicit  
Surface

Pointcloud

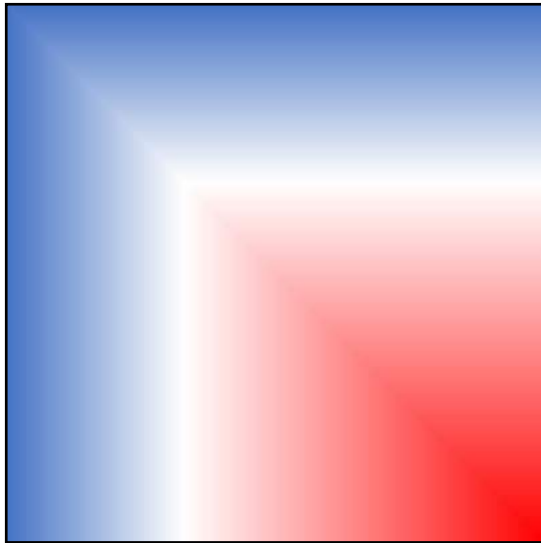
Mesh

# 3D Shape Representations: Implicit Functions

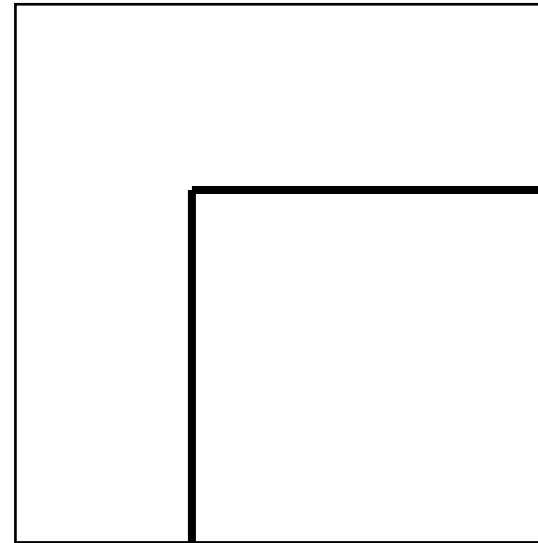
Learn a function to classify arbitrary 3D points as inside / outside the shape

$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set  $\{\mathbf{x} : o(\mathbf{x}) = \frac{1}{2}\}$



Implicit function



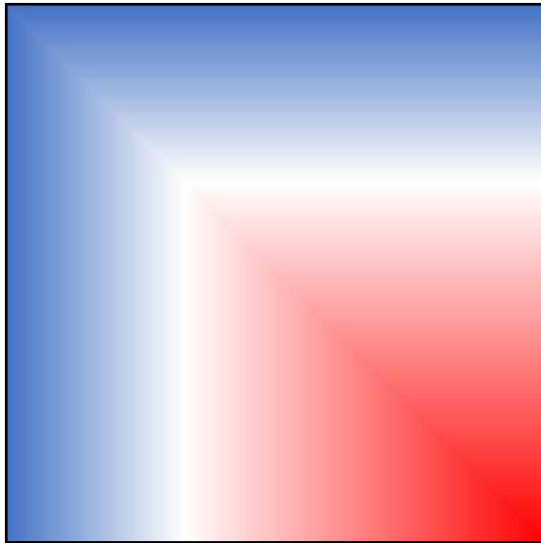
Explicit Shape

# 3D Shape Representations: Implicit Functions

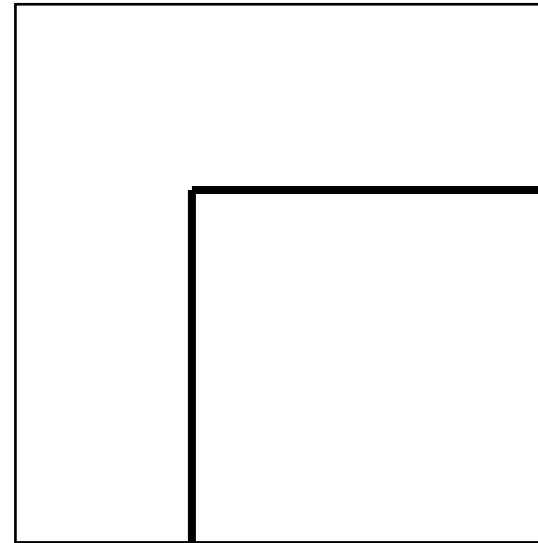
Learn a function to classify arbitrary 3D points as inside / outside the shape

$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set  $\{x : o(x) = \frac{1}{2}\}$



Implicit function



Explicit Shape

Same idea: **signed distance function (SDF)** gives the Euclidean distance to the surface of the shape; sign gives inside / outside

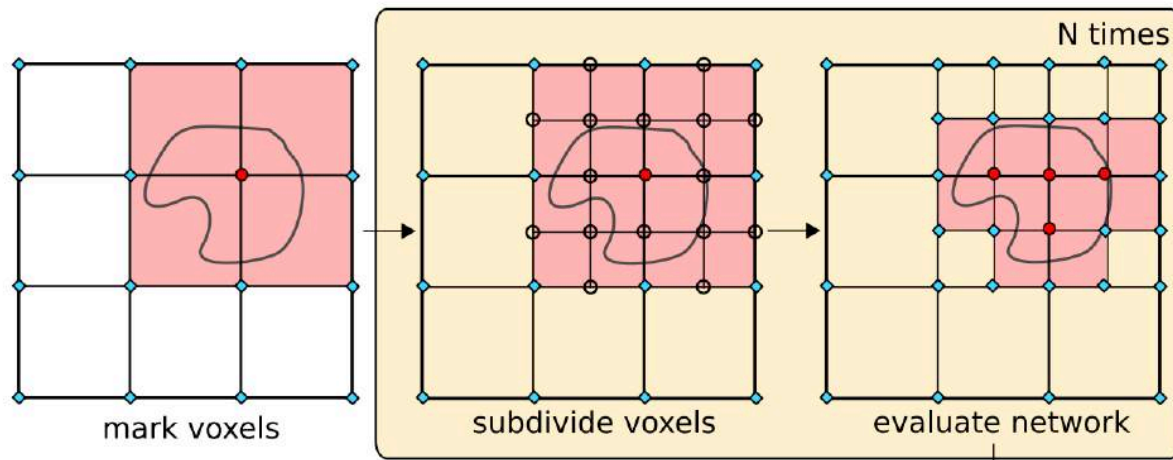
# 3D Shape Representations: Implicit Functions

Learn a function to classify arbitrary 3D points as inside / outside the shape

$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set

$$\{\mathbf{x} : o(\mathbf{x}) = \frac{1}{2}\}$$



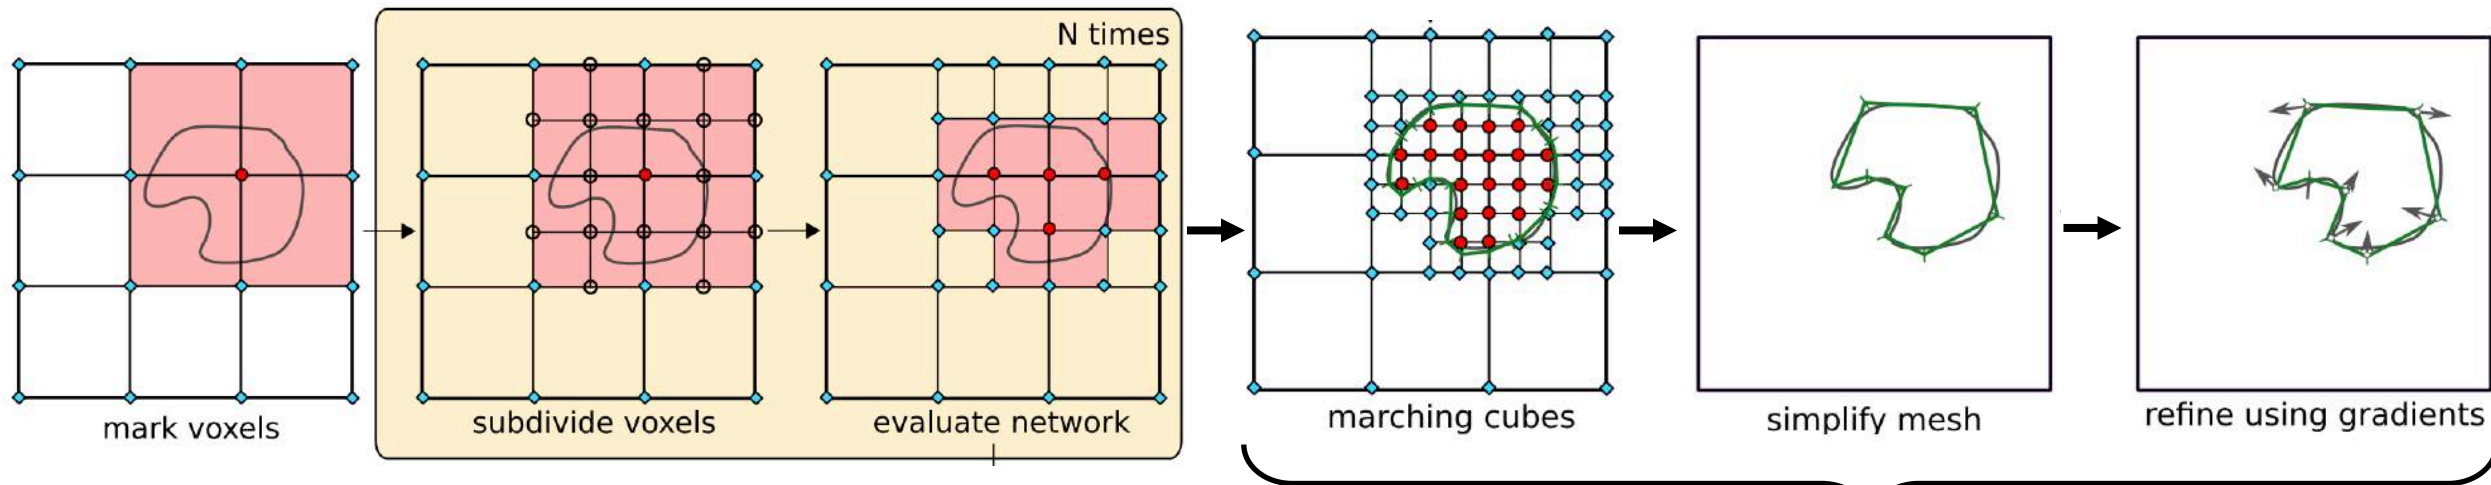
Allows for multiscale outputs like Oct-Trees

# 3D Shape Representations: Implicit Functions

Learn a function to classify arbitrary 3D points as inside / outside the shape

$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

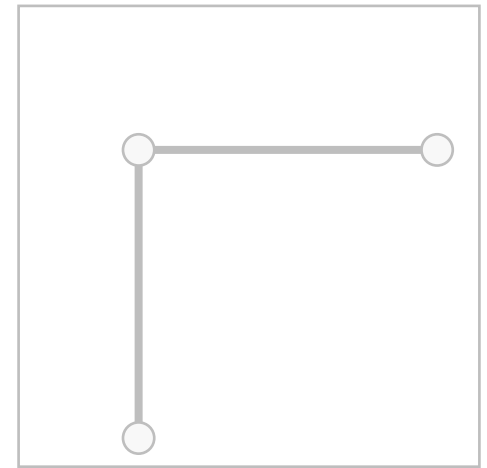
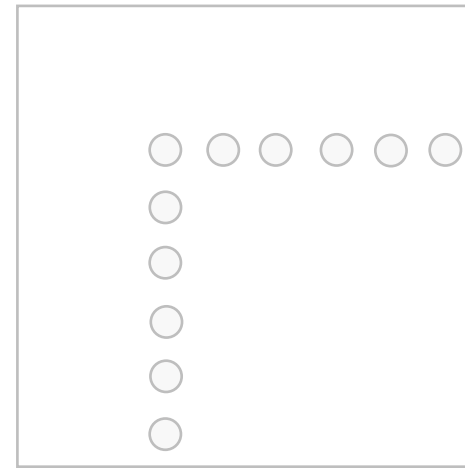
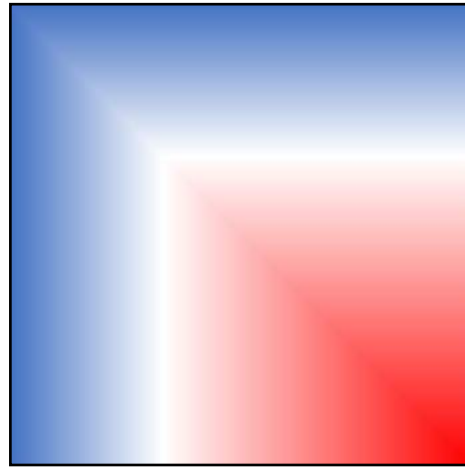
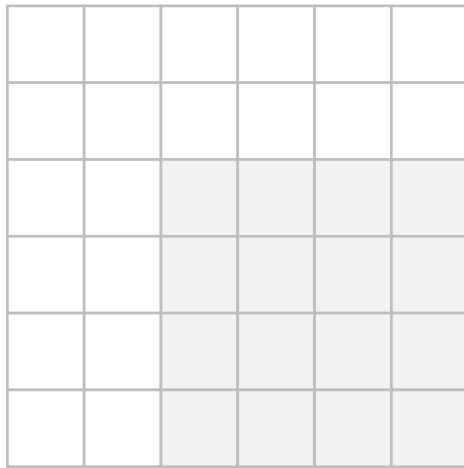
The surface of the 3D object is the level set  $\{\mathbf{x} : o(\mathbf{x}) = \frac{1}{2}\}$



Extracting explicit shape outputs  
requires post-processing



# 3D Shape Representations



Depth  
Map

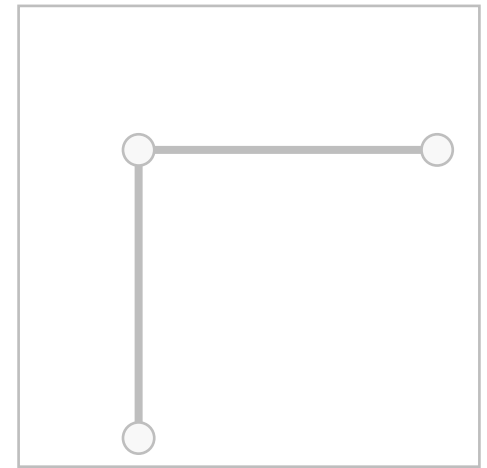
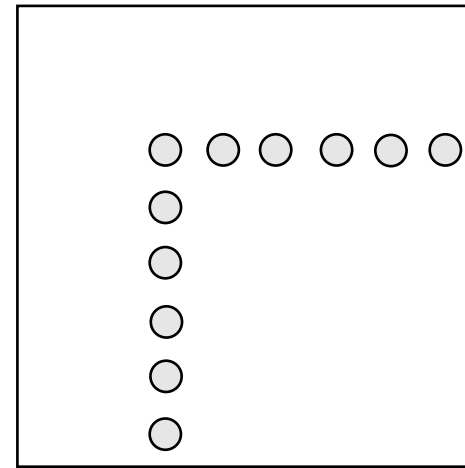
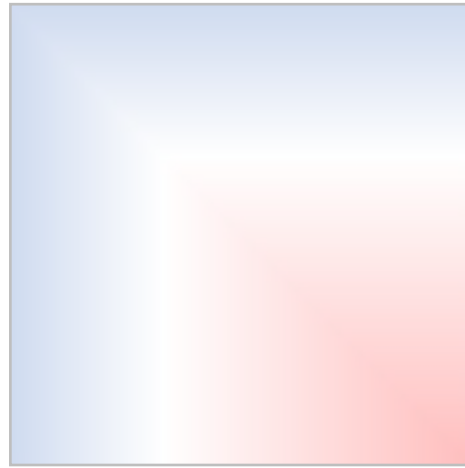
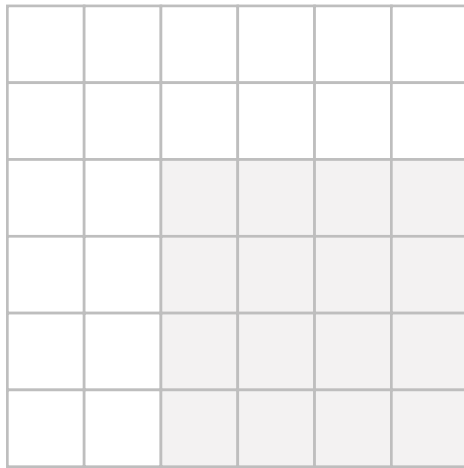
Voxel  
Grid

Implicit  
Surface

Pointcloud

Mesh

# 3D Shape Representations



Depth  
Map

Voxel  
Grid

Implicit  
Surface

Pointcloud

Mesh

# 3D Shape Representations: Point Cloud

- Represent shape as a set of  $P$  points in 3D space
- (+) Can represent fine structures without huge numbers of points
- ( ) Requires new architecture, losses, etc
- (-) Doesn't explicitly represent the surface of the shape: extracting a mesh for rendering or other applications requires post-processing



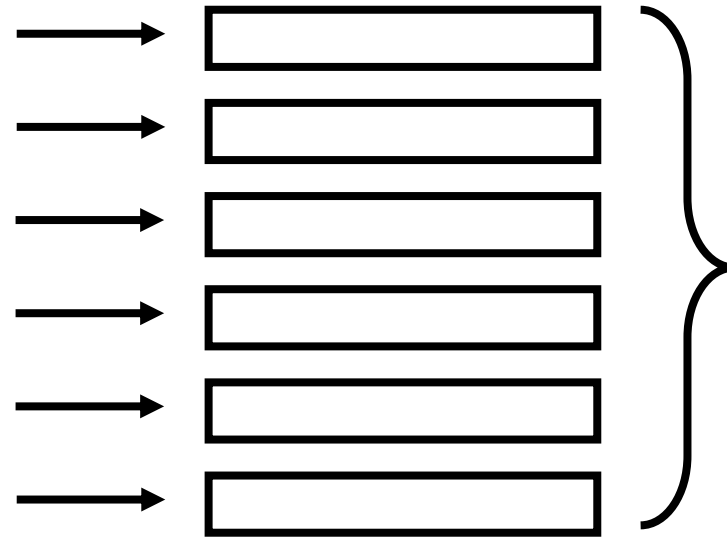
Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

# Processing Pointcloud Inputs: PointNet

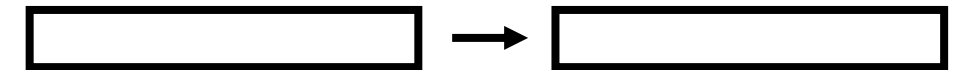
Want to process  
pointclouds as **sets**:  
order should not matter

Run MLP on  
each point

Max-Pool



Fully  
Connected



**Input pointcloud:**

$P \times 3$

**Point features:**

$P \times D$

**Pooled vector:**

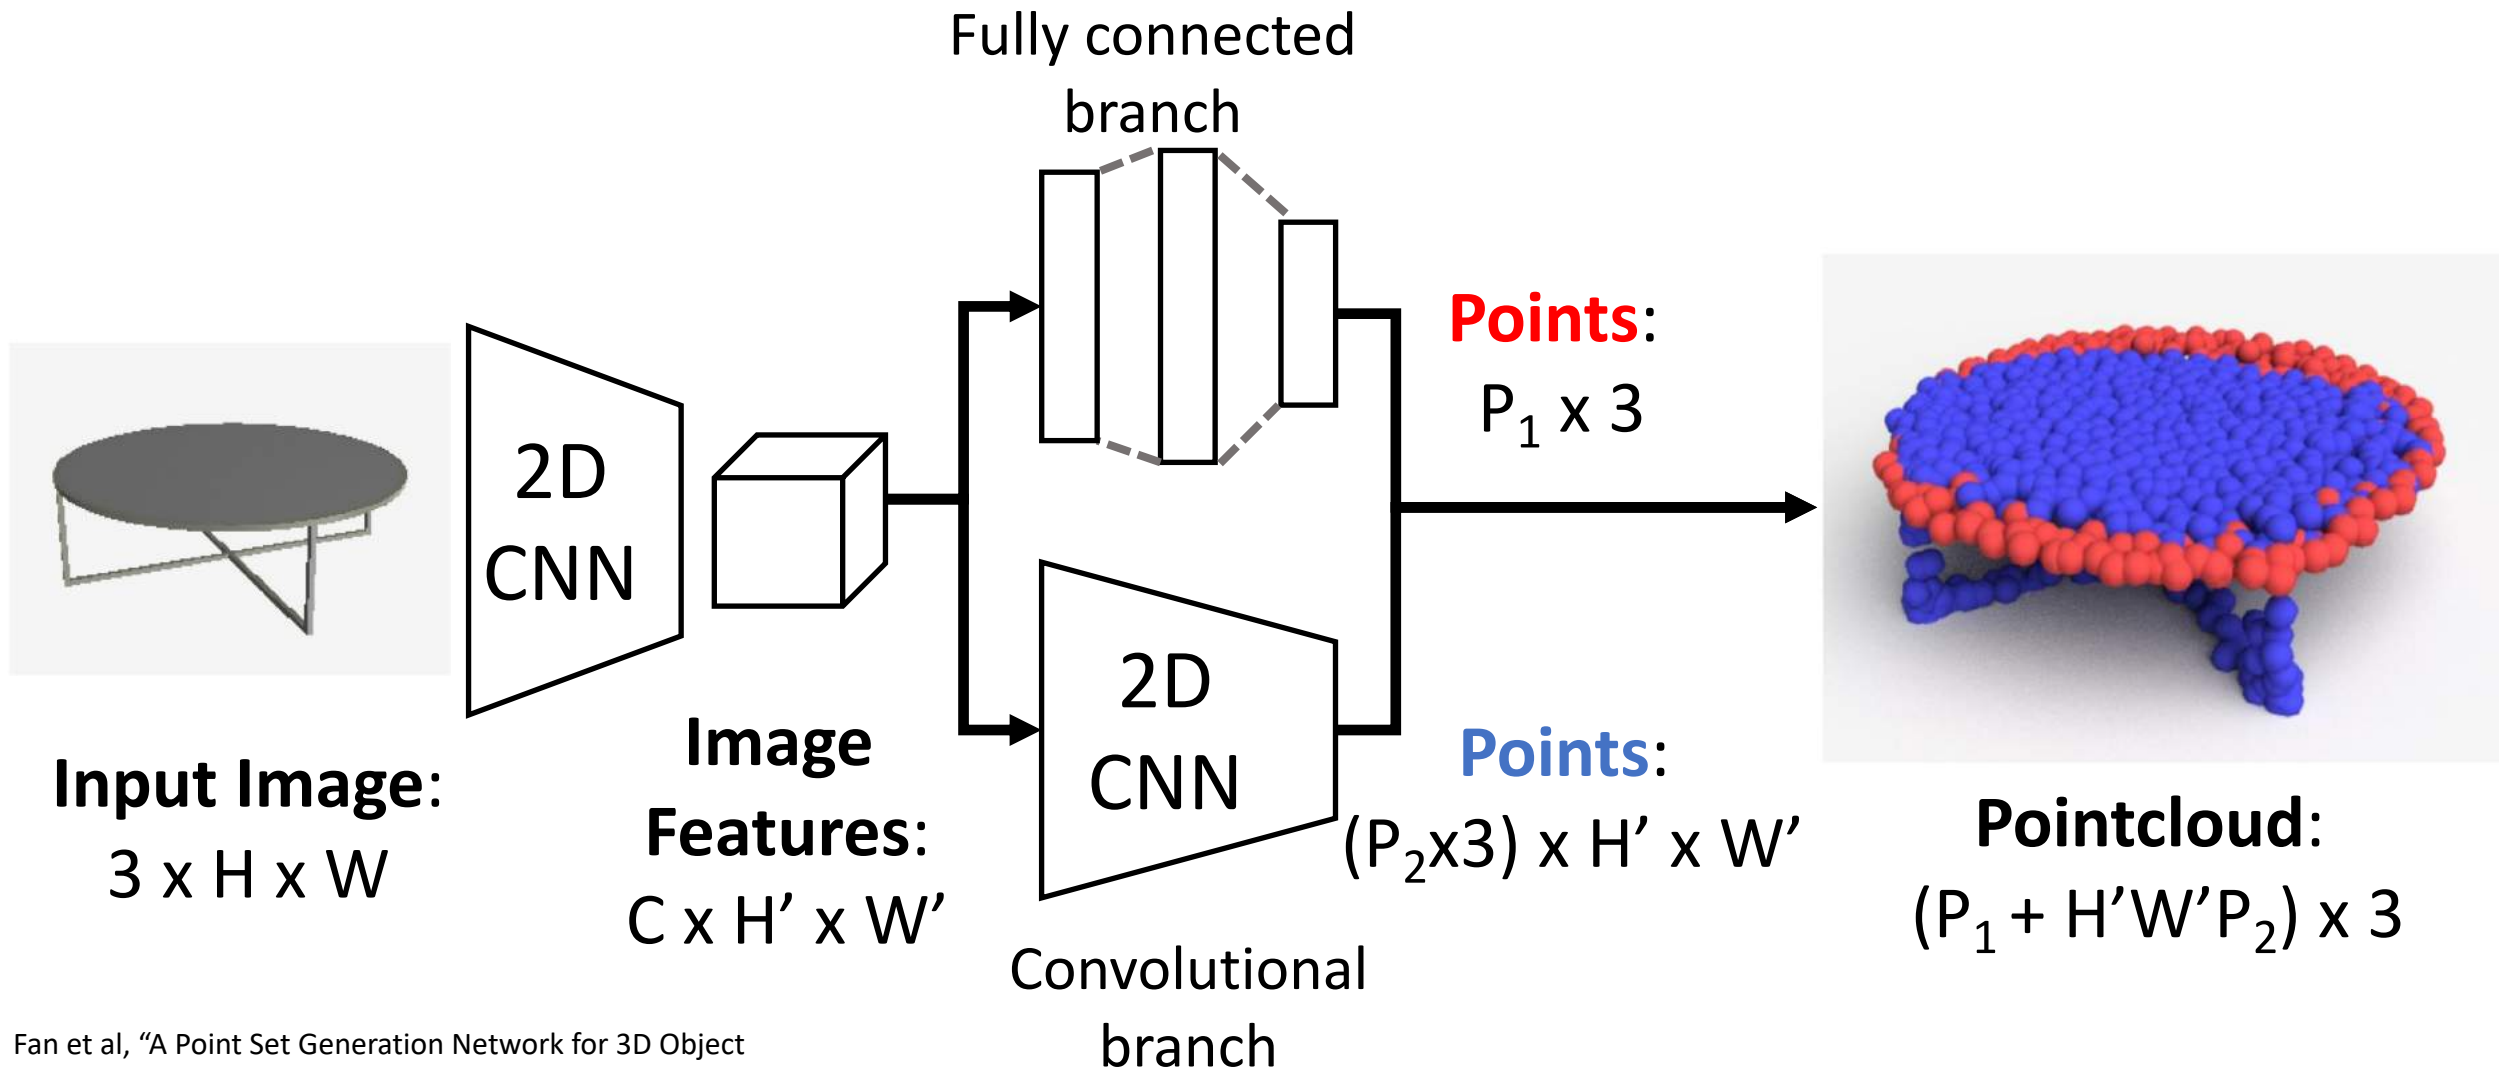
$D$

**Class score:**

$C$

Qi et al, "PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation", CVPR 2017  
Qi et al, "PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space", NeurIPS 2017

# Generating Pointcloud Outputs



Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

# Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds **as sets!**

Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

# Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds **as sets**!

**Chamfer distance** is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$

Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

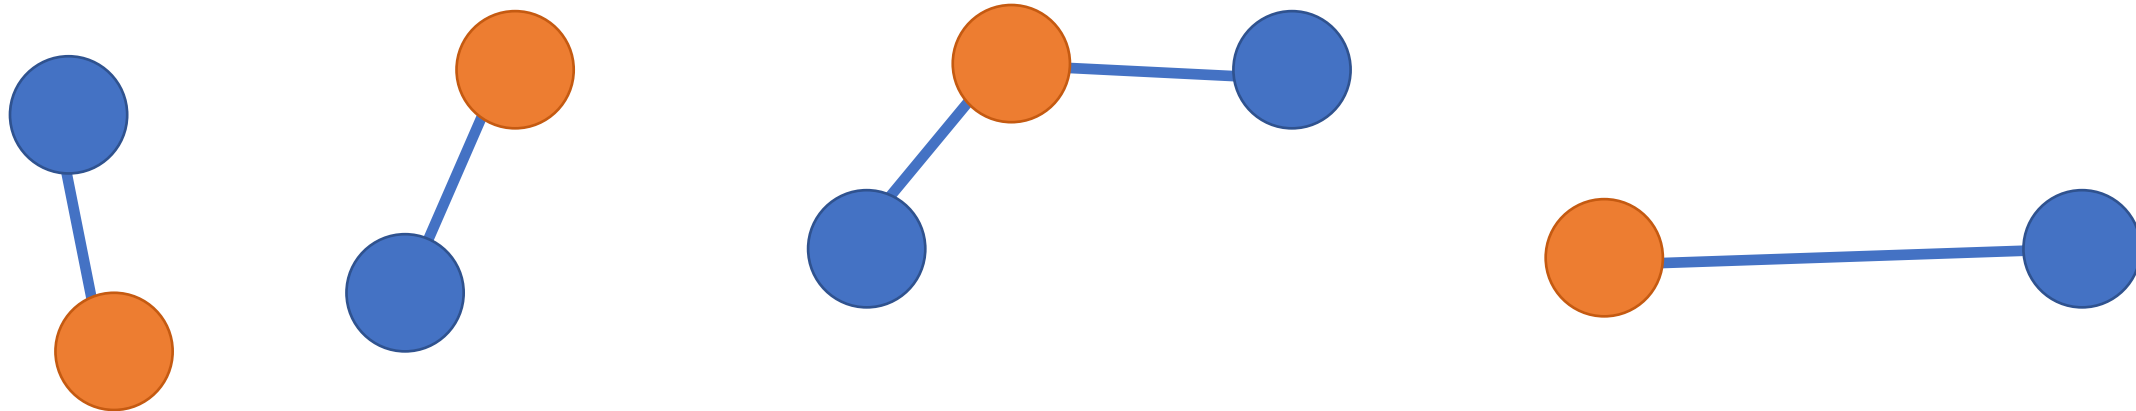


# Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds **as sets**!

**Chamfer distance** is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$



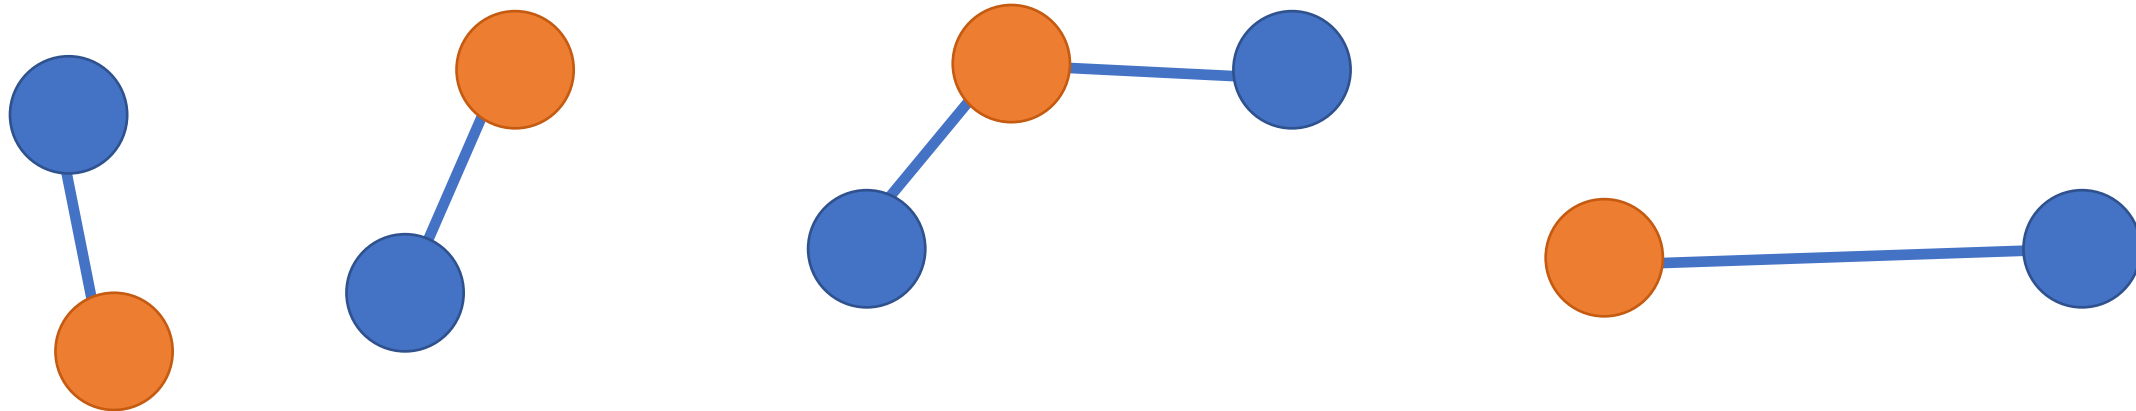
Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

# Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds **as sets**!

**Chamfer distance** is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$



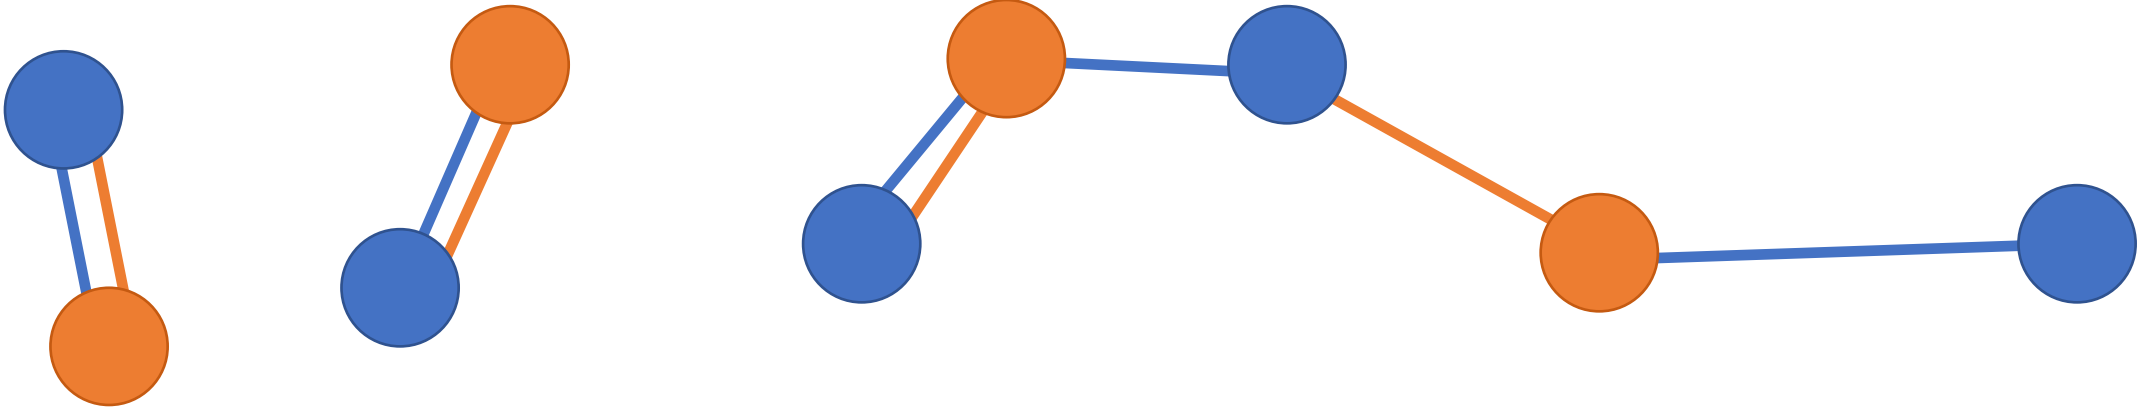
Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

# Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds **as sets**!

**Chamfer distance** is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$

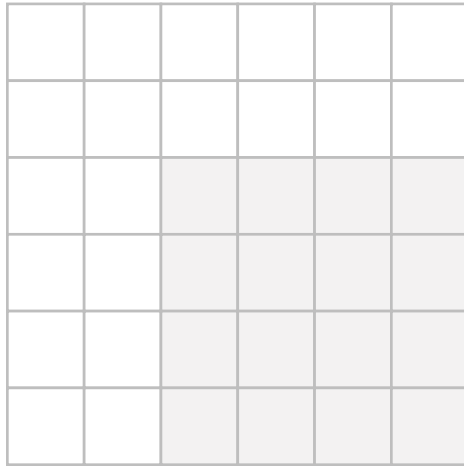


Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

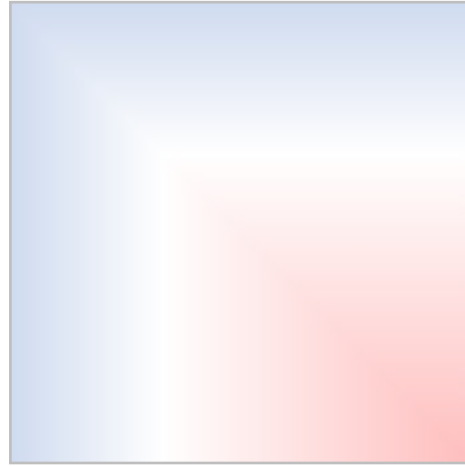
# 3D Shape Representations



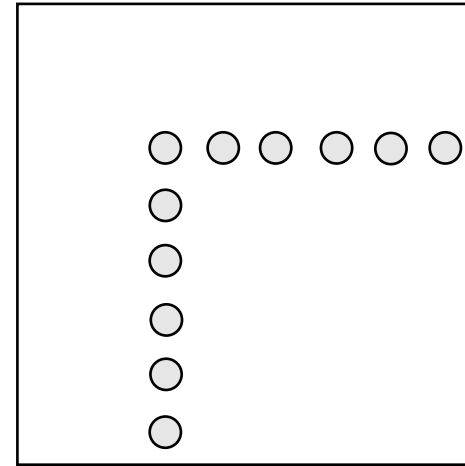
Depth  
Map



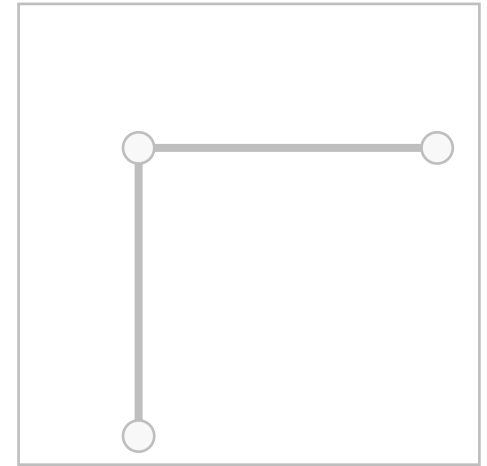
Voxel  
Grid



Implicit  
Surface

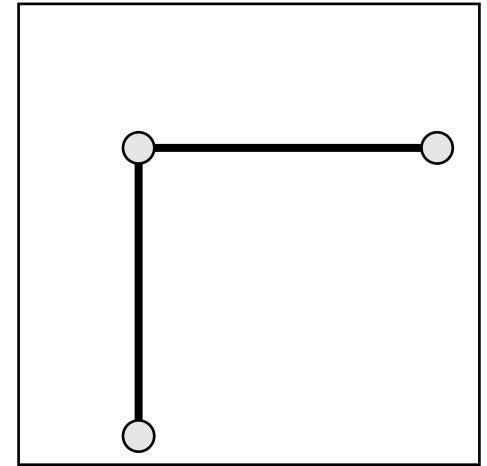
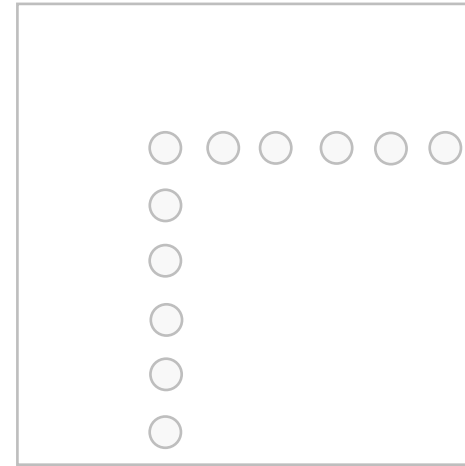
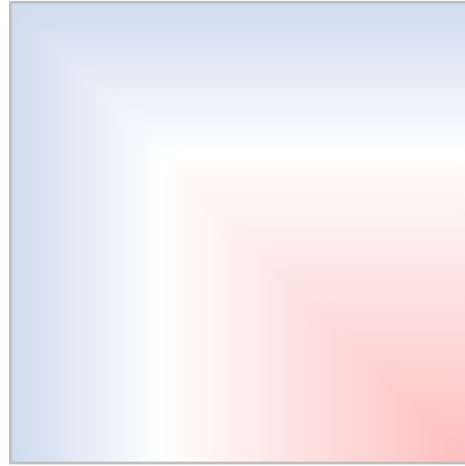
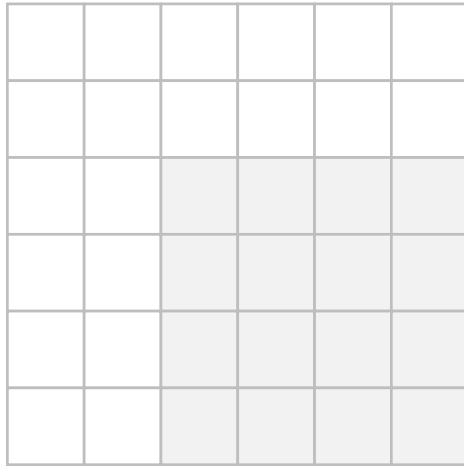


Pointcloud



Mesh

# 3D Shape Representations



Depth  
Map

Voxel  
Grid

Implicit  
Surface

Pointcloud

Mesh

# 3D Shape Representations: Triangle Mesh

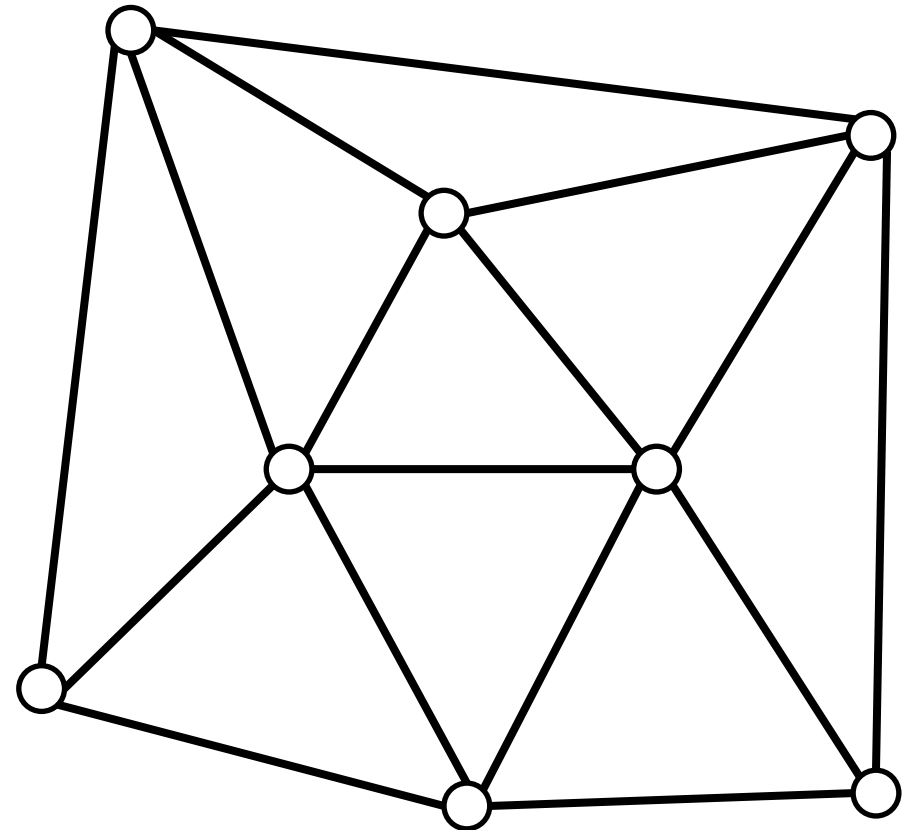
Represent a 3D shape as a set of triangles

**Vertices:** Set of  $V$  points in 3D space

**Faces:** Set of triangles over the vertices

(+) Standard representation for graphics

(+) Explicitly represents 3D shapes



# 3D Shape Representations: Triangle Mesh

Represent a 3D shape as a set of triangles

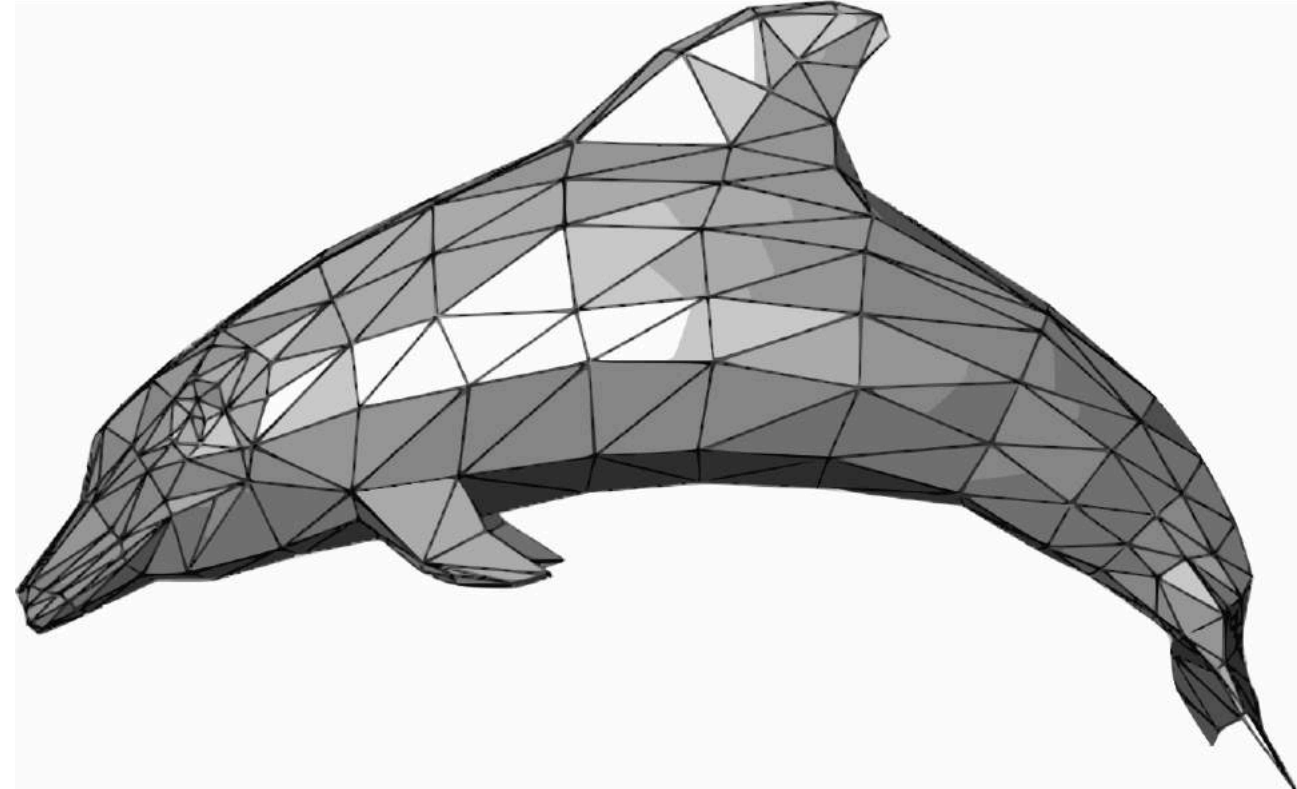
**Vertices:** Set of  $V$  points in 3D space

**Faces:** Set of triangles over the vertices

(+) Standard representation for graphics

(+) Explicitly represents 3D shapes

(+) Adaptive: Can represent flat surfaces very efficiently, can allocate more faces to areas with fine detail



[Dolphin image](#) is in the public domain

# 3D Shape Representations: Triangle Mesh

Represent a 3D shape as a set of triangles

**Vertices:** Set of  $V$  points in 3D space

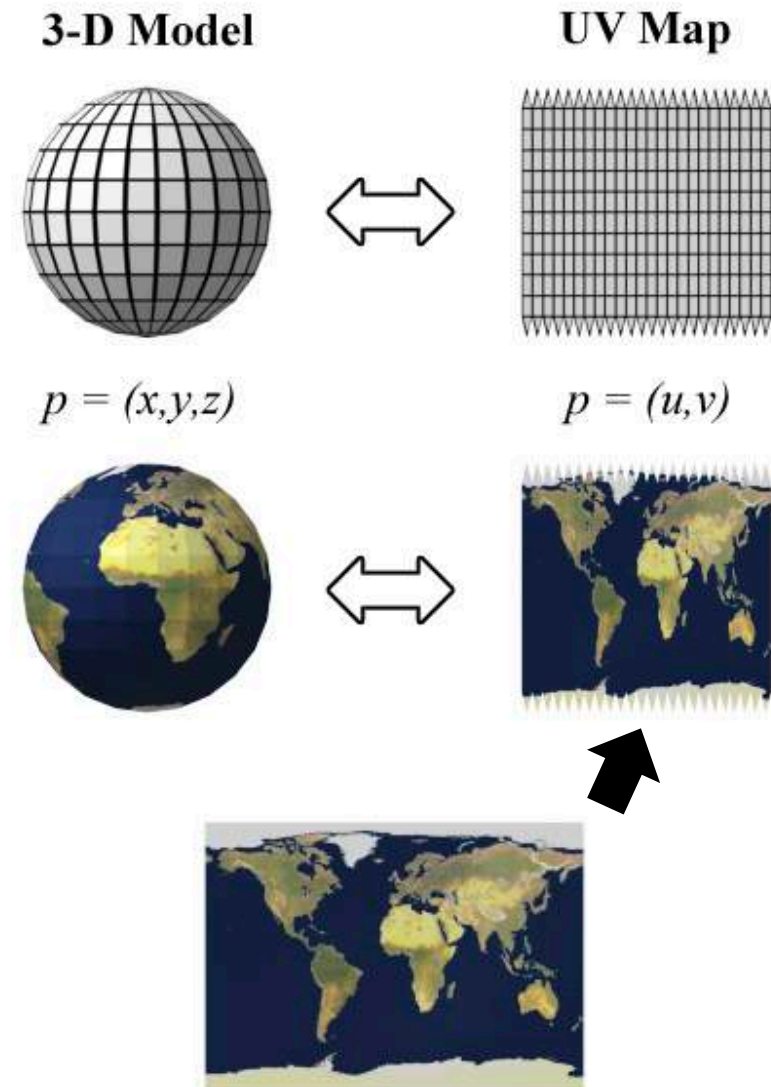
**Faces:** Set of triangles over the vertices

(+) Standard representation for graphics

(+) Explicitly represents 3D shapes

(+) Adaptive: Can represent flat surfaces very efficiently, can allocate more faces to areas with fine detail

(+) Can attach data on verts and interpolate over the whole surface: RGB colors, texture coordinates, normal vectors, etc.



UV mapping figure is licensed under CC BY-SA 3.0. Figure slightly reorganized.



# 3D Shape Representations: Triangle Mesh

Represent a 3D shape as a set of triangles

**Vertices:** Set of  $V$  points in 3D space

**Faces:** Set of triangles over the vertices

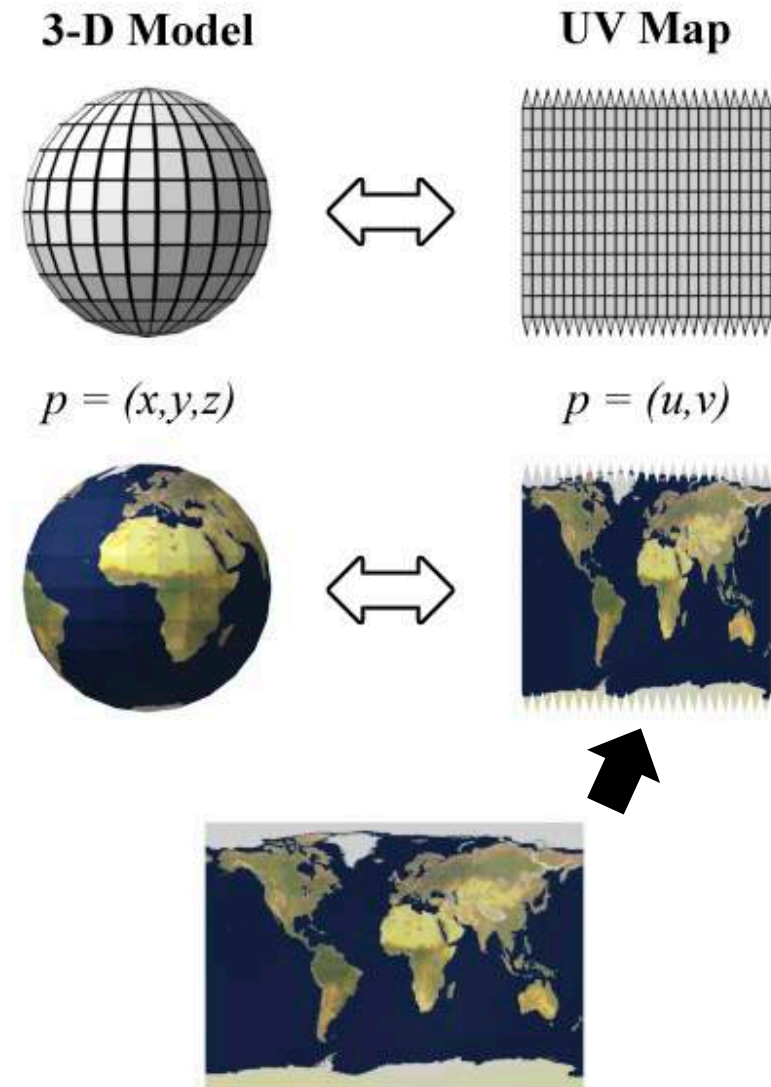
(+) Standard representation for graphics

(+) Explicitly represents 3D shapes

(+) Adaptive: Can represent flat surfaces very efficiently, can allocate more faces to areas with fine detail

(+) Can attach data on verts and interpolate over the whole surface: RGB colors, texture coordinates, normal vectors, etc.

(-) Nontrivial to process with neural nets!

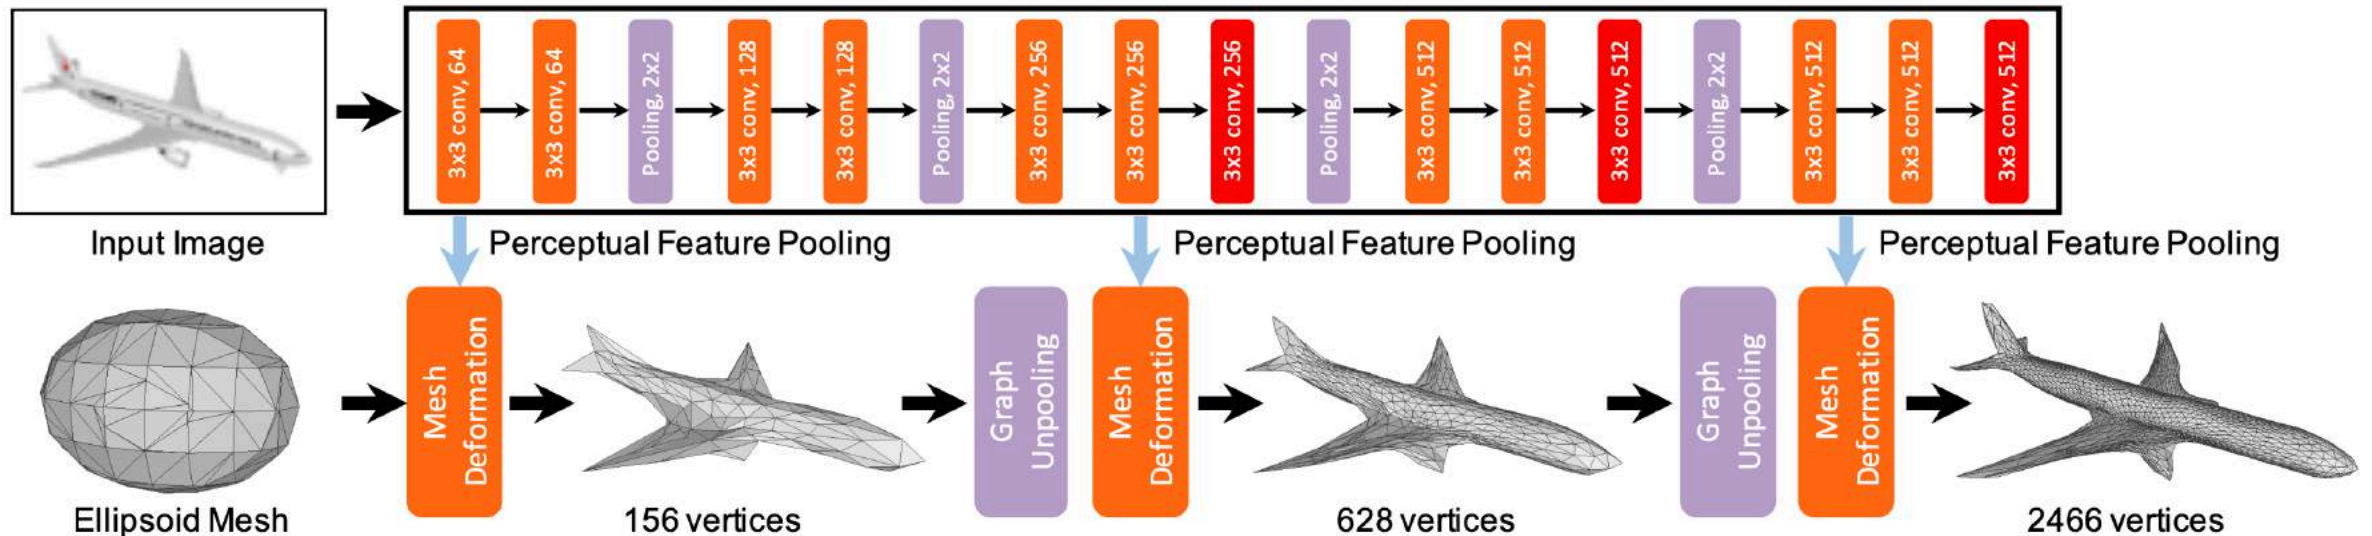


UV mapping figure is licensed under CC BY-SA 3.0. Figure slightly reorganized.

# Predicting Meshes: Pixel2Mesh

**Input:** Single RGB  
Image of an object

**Output:** Triangle  
mesh for the object



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

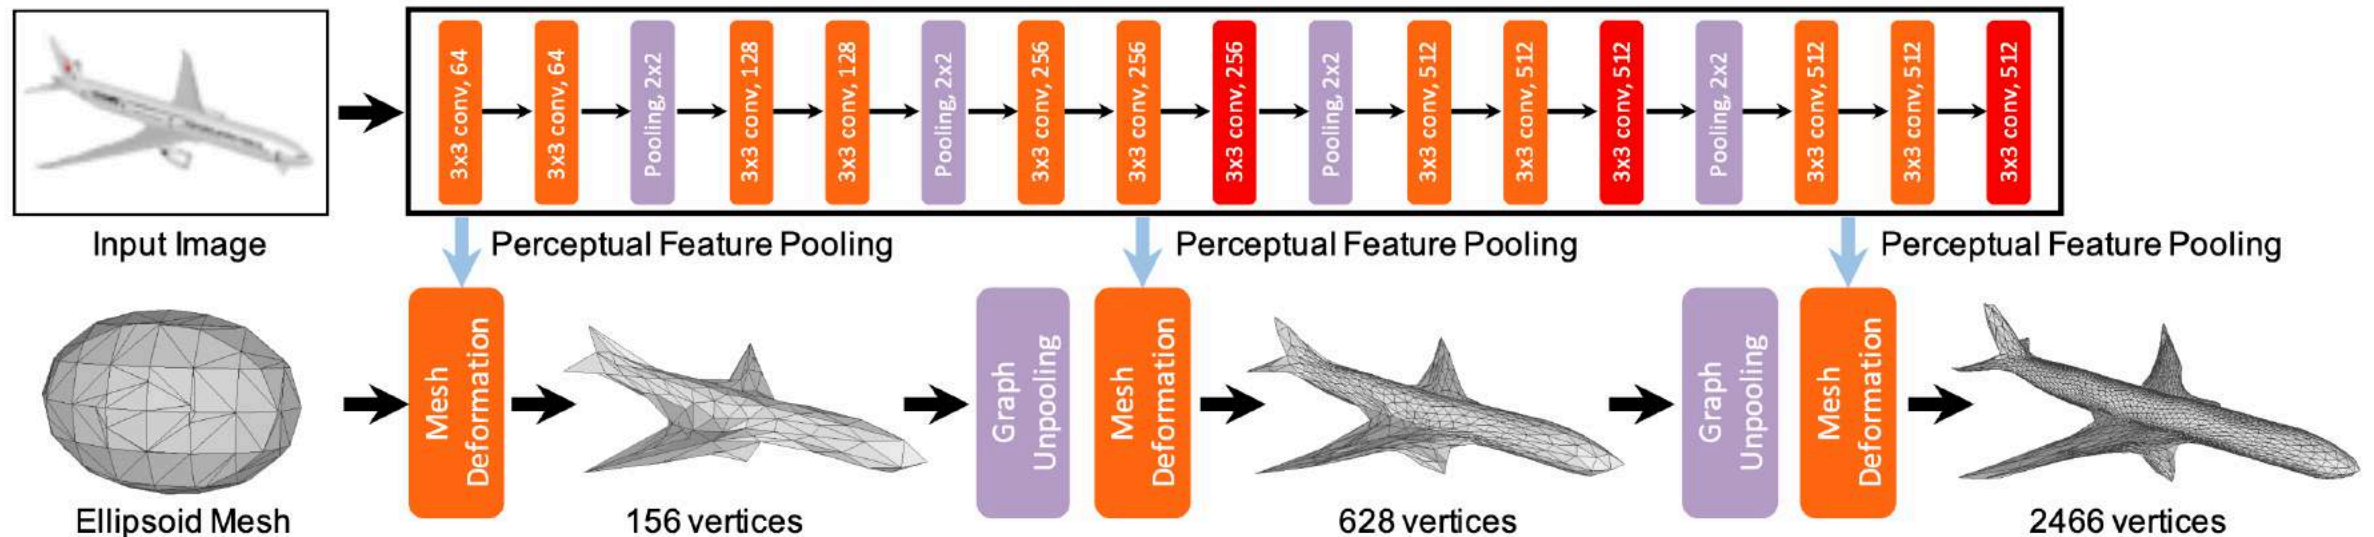
# Predicting Meshes: Pixel2Mesh

**Input:** Single RGB  
Image of an object

**Key ideas:**

Iterative Refinement  
Graph Convolution  
Vertex Aligned-Features  
Chamfer Loss Function

**Output:** Triangle  
mesh for the object



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

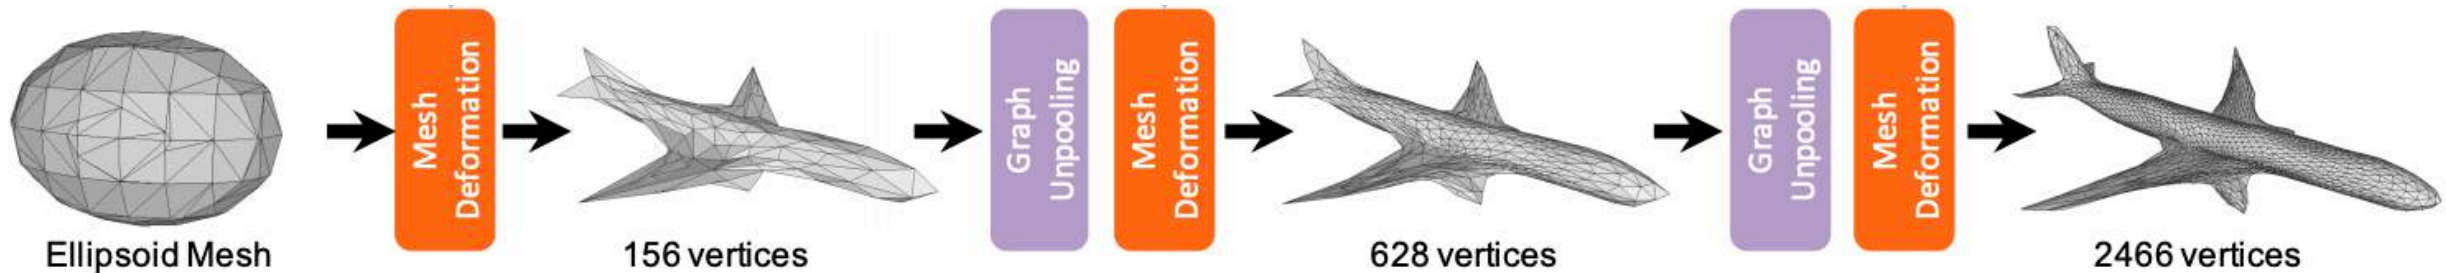
# Predicting Triangle Meshes: Iterative Refinement

## Idea #1: Iterative mesh refinement

Start from initial ellipsoid mesh

Network predicts offsets for each vertex

Repeat.



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

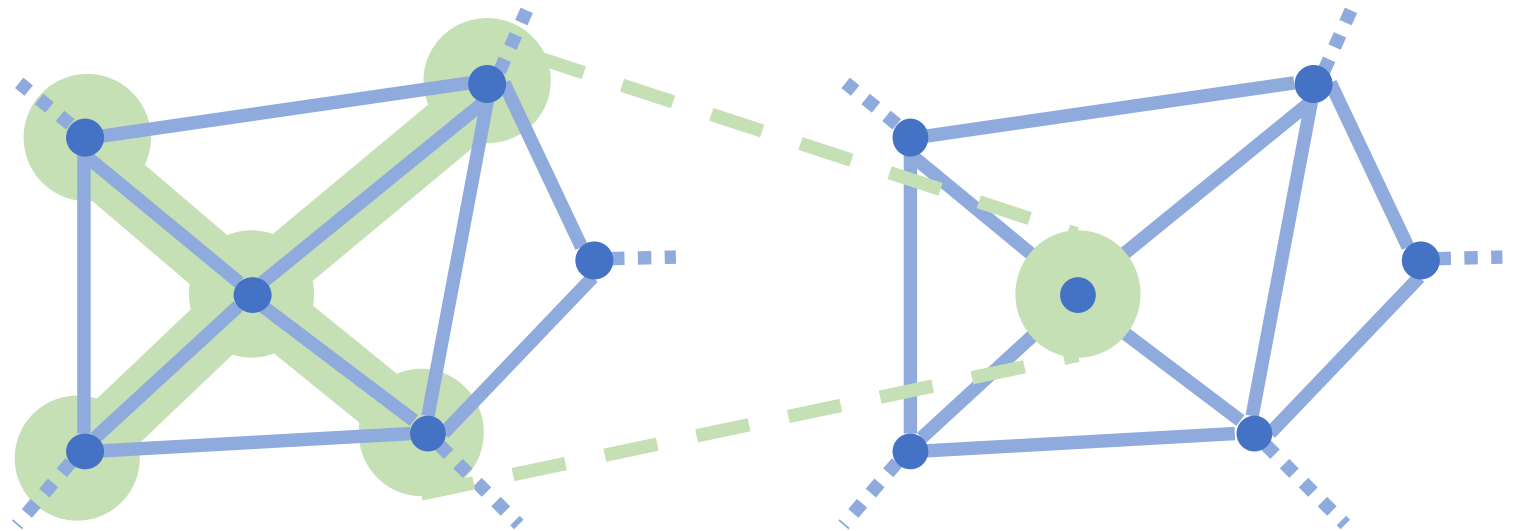
# Predicting Triangle Meshes: Graph Convolution

$$f'_i = W_0 f_i + \sum_{j \in \mathcal{N}(i)} W_1 f_j$$

Vertex  $v_i$  has feature  $f_i$

New feature  $f'_i$  for vertex  $v_i$  depends on feature of neighboring vertices  $\mathcal{N}(i)$

Use same weights  $W_0$  and  $W_1$  to compute all outputs



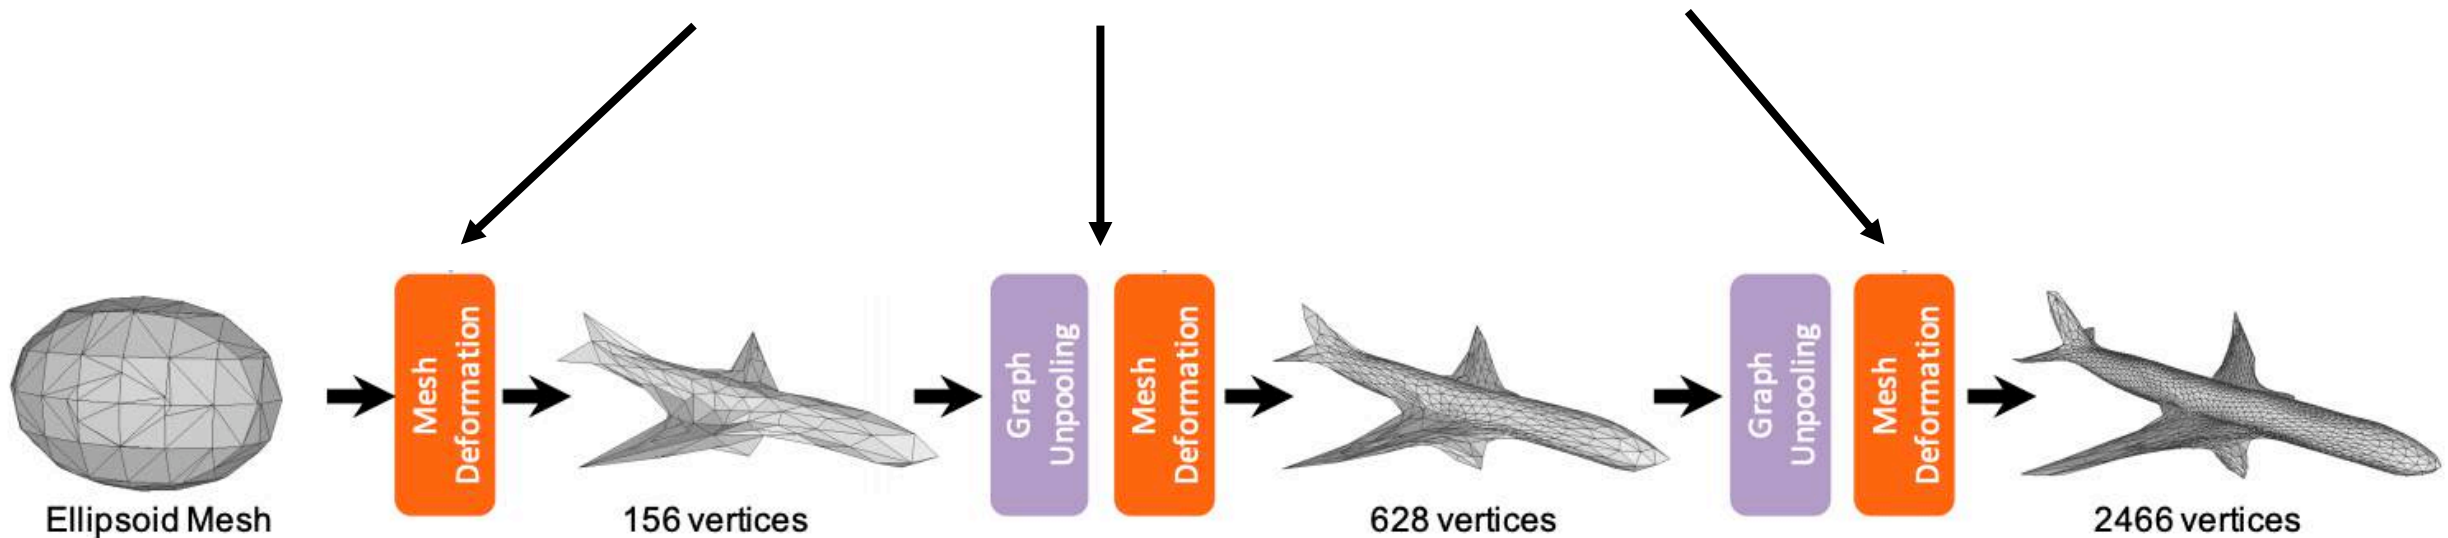
**Input:** Graph with a feature vector at each vertex

**Output:** New feature vector for each vertex



# Predicting Triangle Meshes: Graph Convolution

Each of these blocks consists of a stack of **graph convolution layers** operating on edges of the mesh

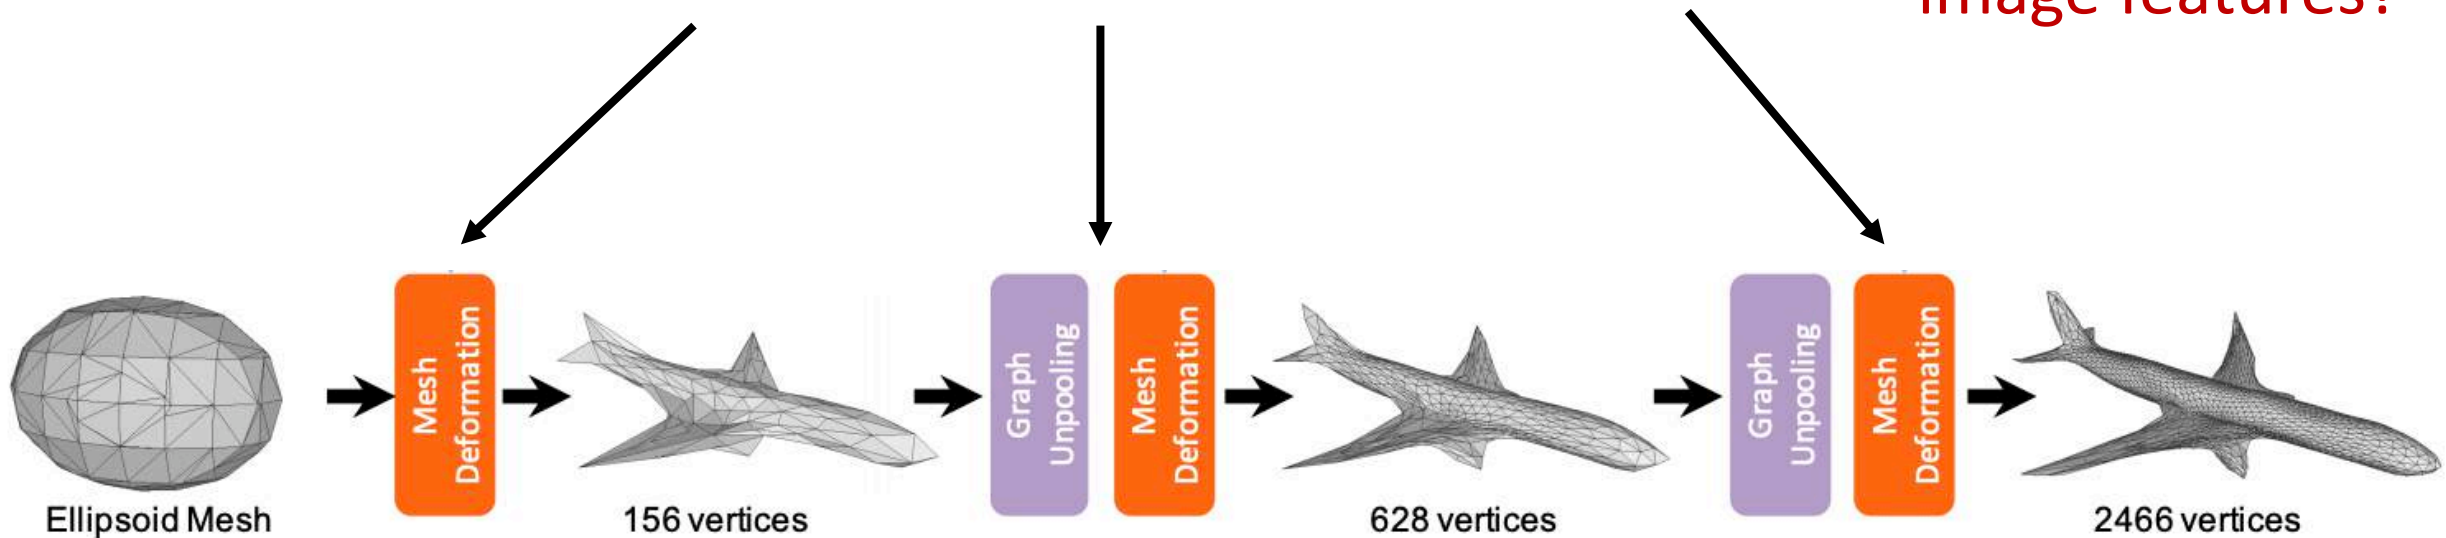


Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

# Predicting Triangle Meshes: Graph Convolution

Each of these blocks consists of a stack of **graph convolution layers** operating on edges of the mesh

Problem: How to incorporate image features?



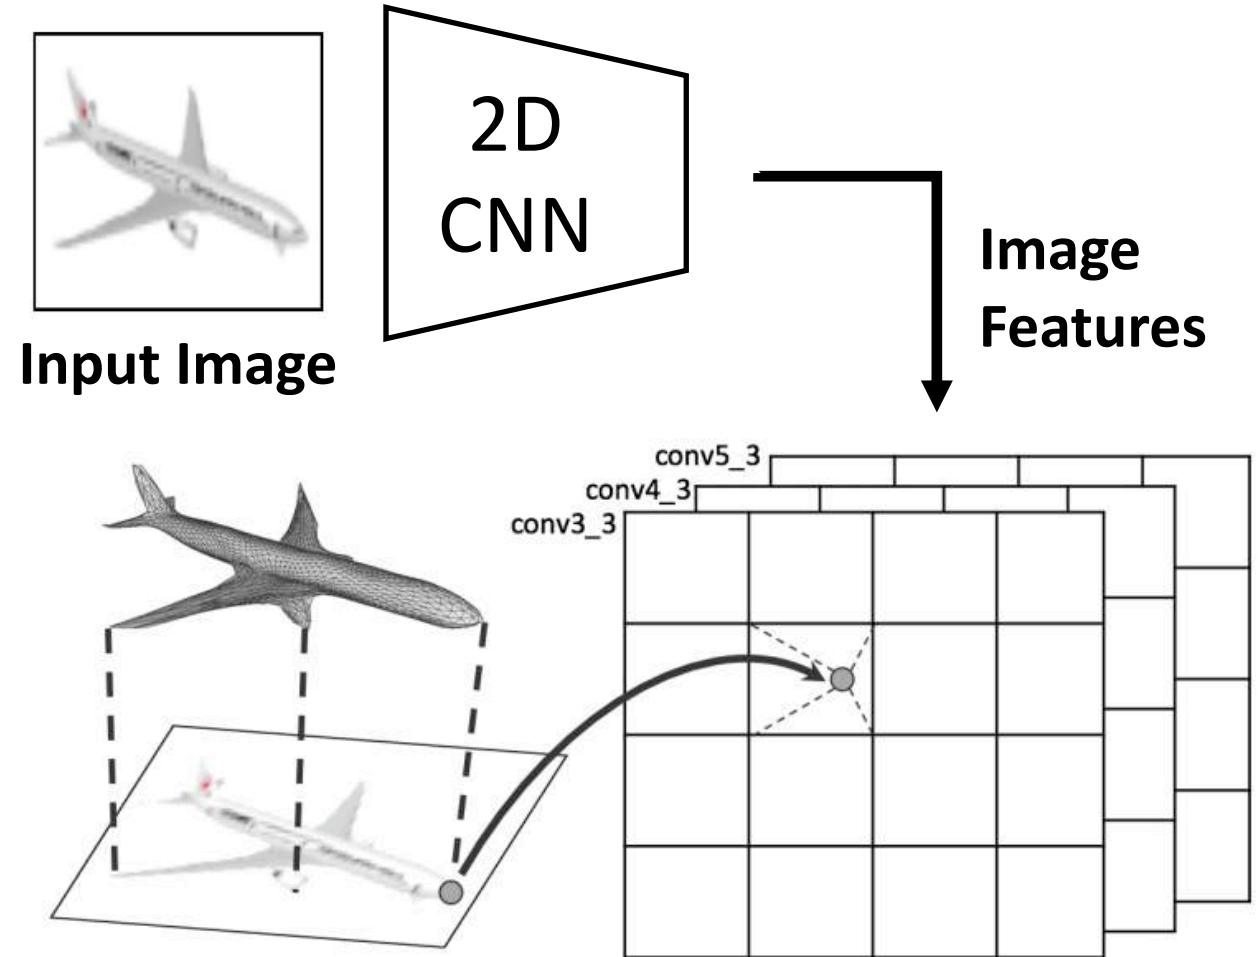


# Predicting Triangle Meshes: Vertex-Aligned Features

## Idea #2: Aligned vertex features

For each vertex of the mesh:

- Use camera information to project onto image plane
- Use bilinear interpolation to sample a CNN feature



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

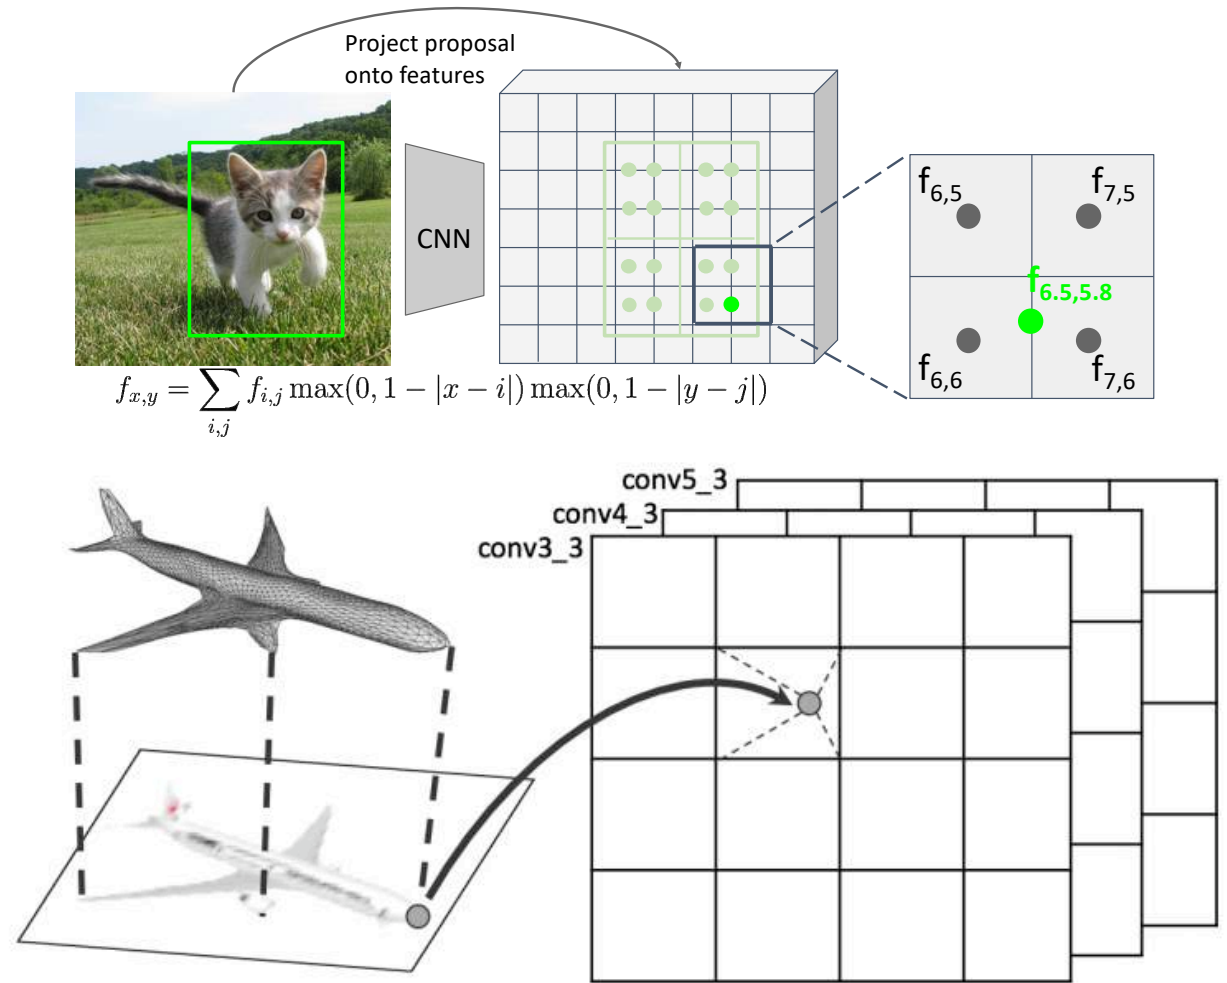
# Predicting Triangle Meshes: Vertex-Aligned Features

## Idea #2: Aligned vertex features

For each vertex of the mesh:

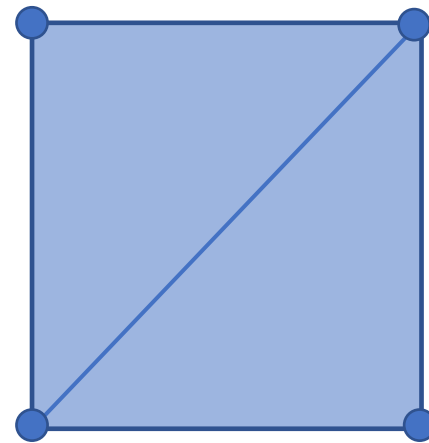
- Use camera information to project onto image plane
- Use bilinear interpolation to sample a CNN feature

Similar to RoI-Align operation from last time: maintains alignment between input image and feature vectors



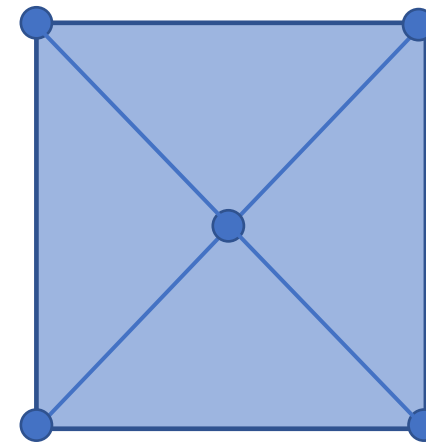
# Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?



Prediction

vs



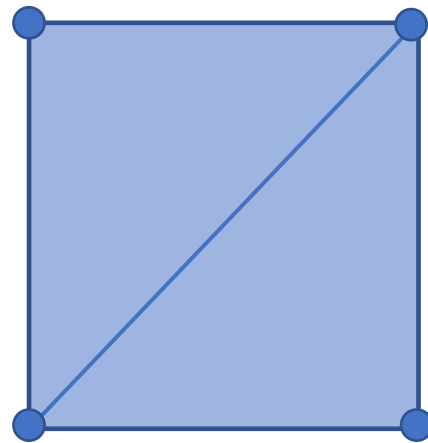
Ground-Truth

Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

# Predicting Meshes: Loss Function

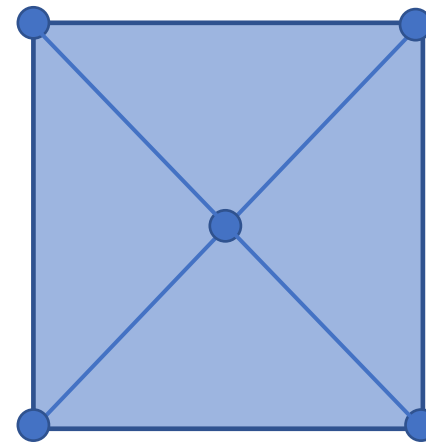
The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

**Idea:** Convert meshes to pointclouds, then compute loss



Prediction

vs

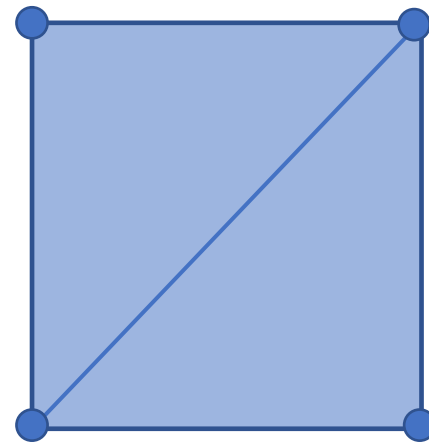


Ground-Truth

# Predicting Meshes: Loss Function

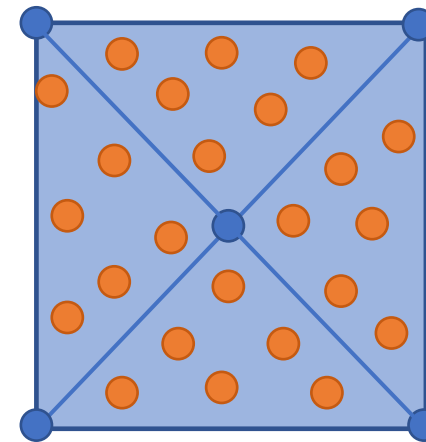
The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

**Idea:** Convert meshes to pointclouds, then compute loss



Prediction

vs



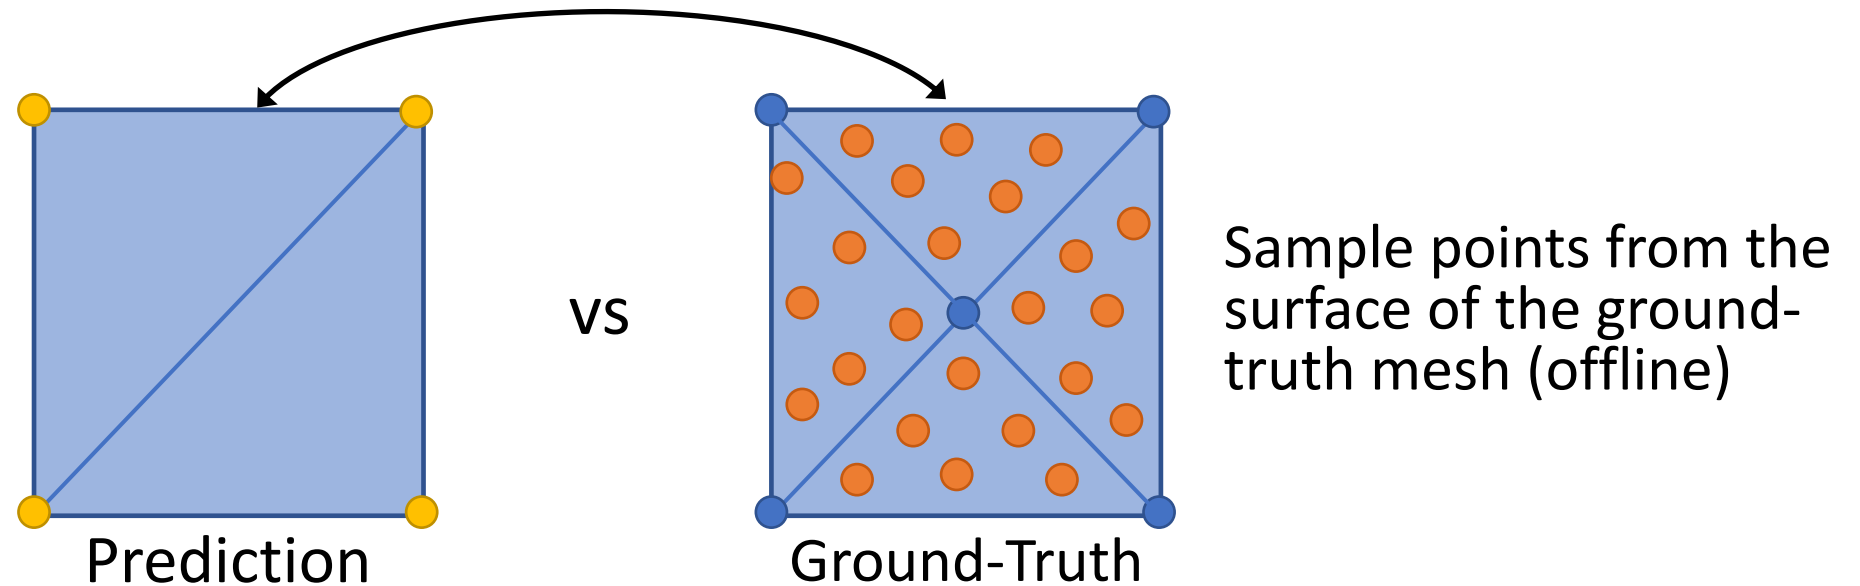
Ground-Truth

Sample points from the surface of the ground-truth mesh (offline)

# Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

Loss = Chamfer distance between **predicted verts** and **ground-truth samples**

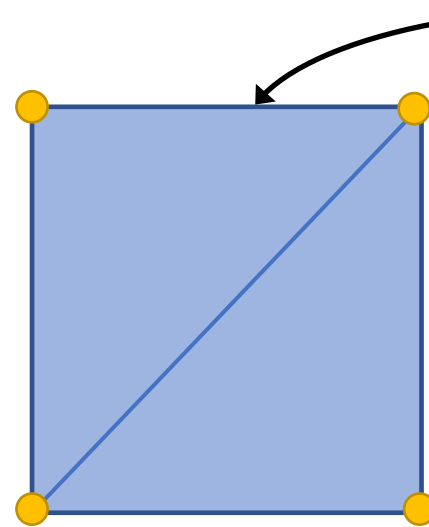


# Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

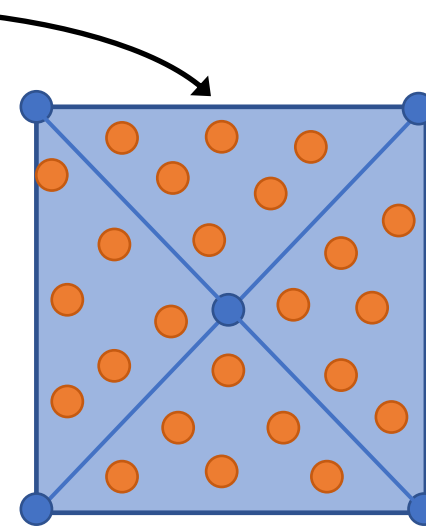
Loss = Chamfer distance between **predicted verts** and **ground-truth samples**

Problem: Doesn't take the interior of predicted faces into account!



Prediction

vs



Ground-Truth

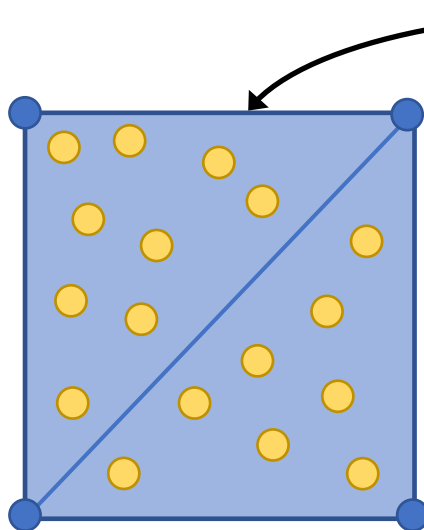
Sample points from the surface of the ground-truth mesh (offline)

# Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

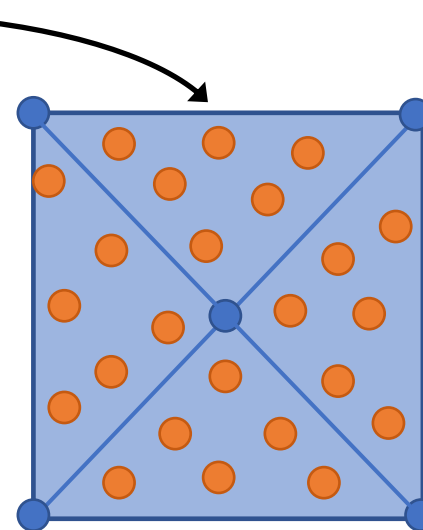
Loss = Chamfer distance between **predicted samples** and **ground-truth samples**

Sample points from the surface of the predicted mesh (online!)



Prediction

vs



Ground-Truth

Sample points from the surface of the ground-truth mesh (offline)



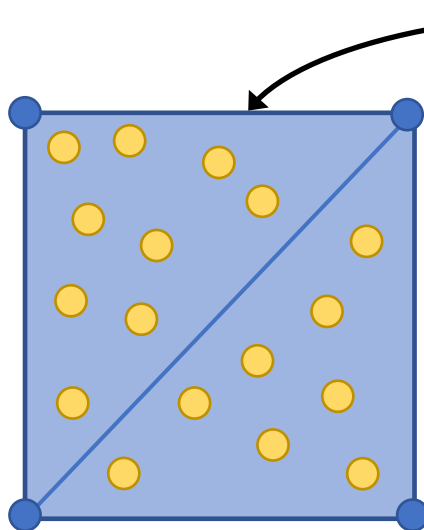
# Predicting Meshes: Loss Function

Problem: Need to sample online! Must be efficient!

Problem: Need to backprop through sampling!

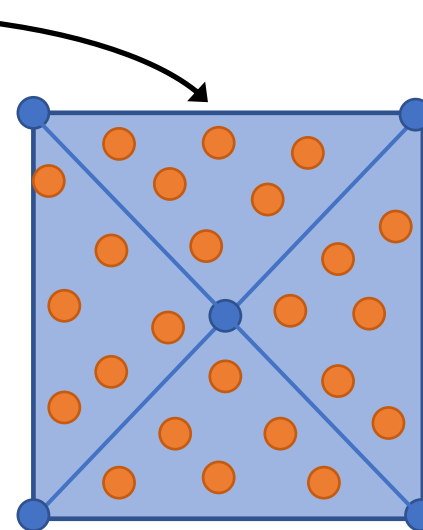
Loss = Chamfer distance between **predicted samples** and **ground-truth samples**

Sample points  
from the surface  
of the predicted  
mesh (online!)



Prediction

vs



Ground-Truth

Sample points from the  
surface of the ground-  
truth mesh (offline)

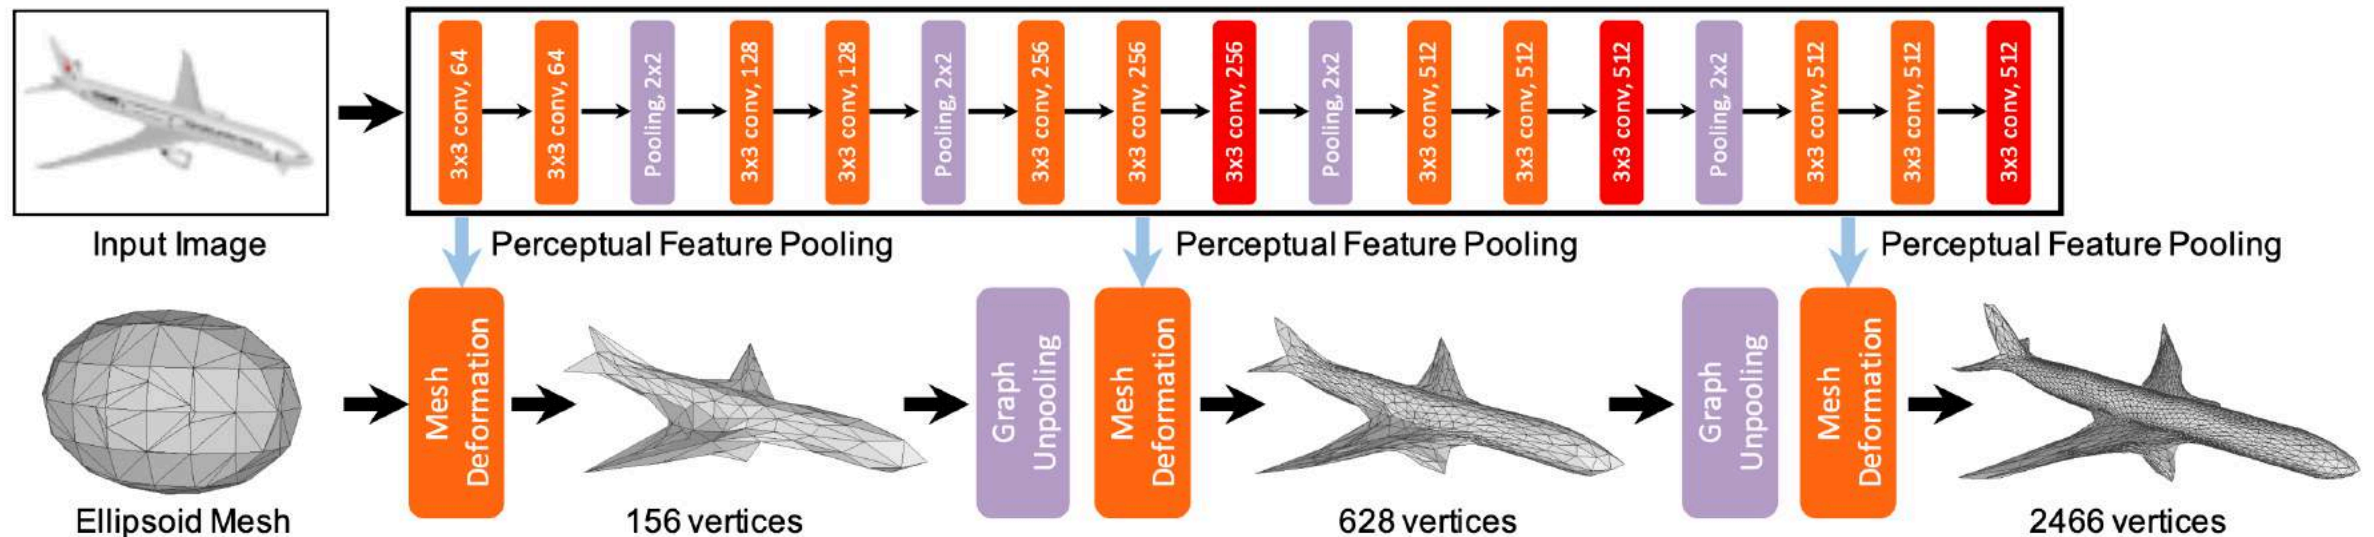
# Predicting Meshes: Pixel2Mesh

**Input:** Single RGB  
Image of an object

**Key ideas:**

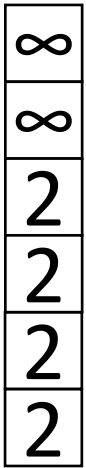
Iterative Refinement  
Graph Convolution  
Vertex Aligned-Features  
Chamfer Loss Function

**Output:** Triangle  
mesh for the object

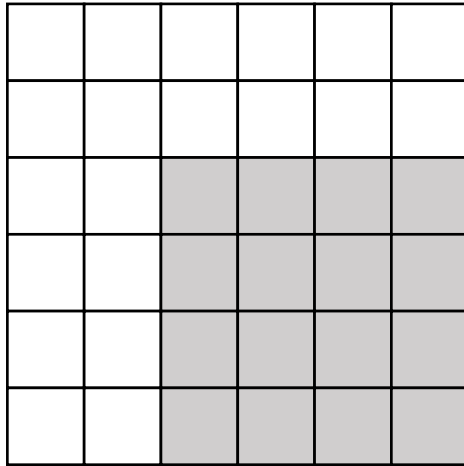


Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

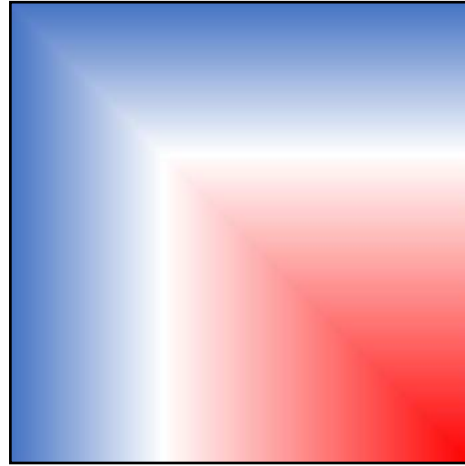
# 3D Shape Representations



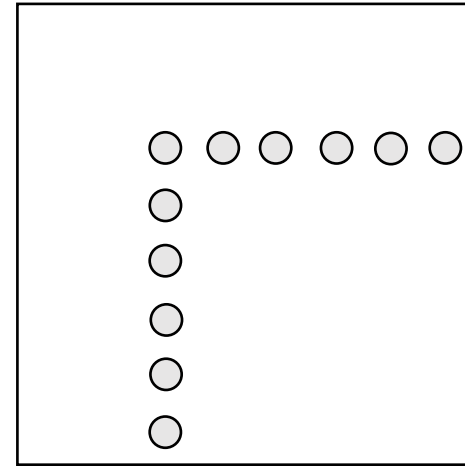
Depth  
Map



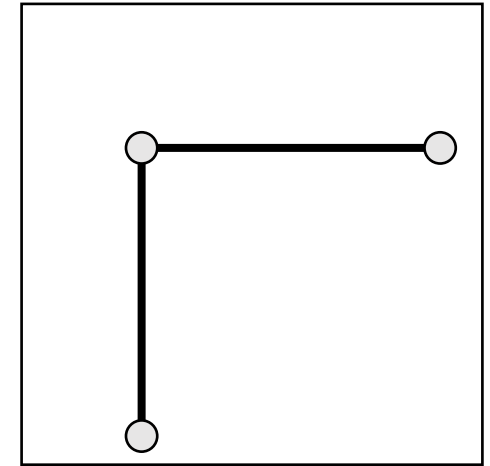
Voxel  
Grid



Implicit  
Surface



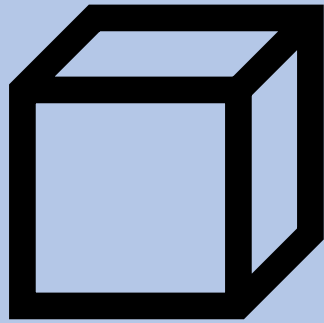
Pointcloud



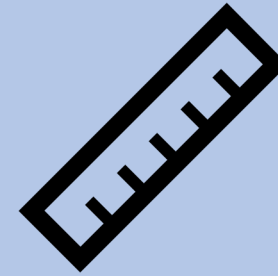
Mesh

# 3D Shape Prediction

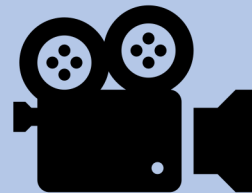
Shape Representations



Metrics



Camera Systems

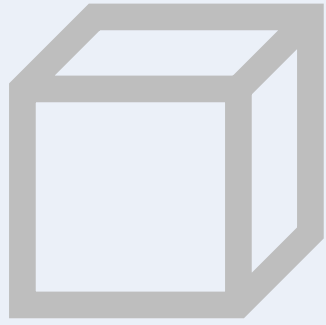


Datasets



# 3D Shape Prediction

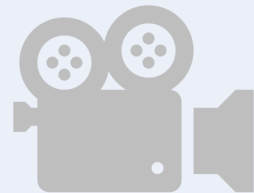
Shape Representations



Metrics



Camera Systems

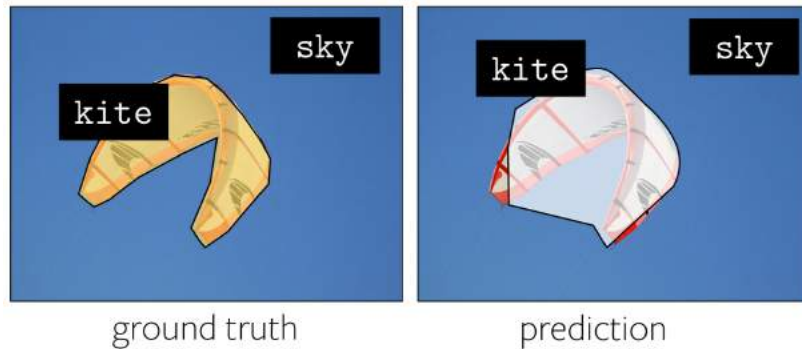


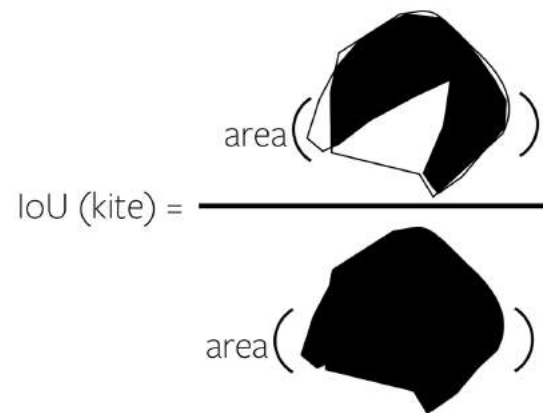
Datasets



# Shape Comparison Metrics: Intersection over Union

In 2D, we evaluate boxes and segmentation masks with intersection over union (IoU):



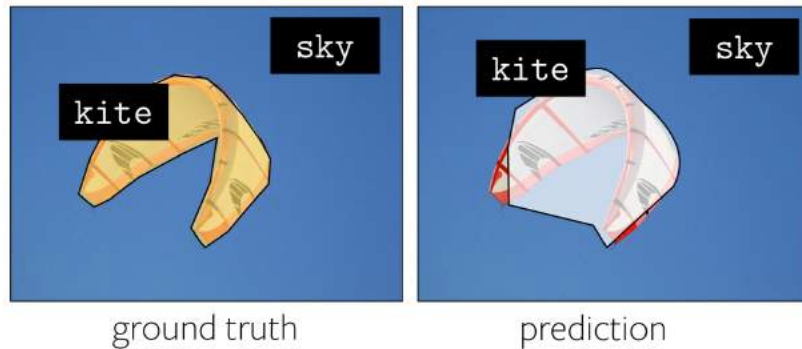
$$\text{IoU (kite)} = \frac{\text{area}(\text{intersection})}{\text{area}(\text{union})}$$


The diagram shows the formula for IoU (kite) as the ratio of the area of the intersection of the ground truth and prediction masks to the area of the union of the two masks. The ground truth mask is shown as a yellow shape, and the prediction mask is shown as a red shape. The intersection is the area where they overlap, and the union is the total area covered by both masks.

Figure credit: Alexander Kirillov

# Shape Comparison Metrics: Intersection over Union

In 2D, we evaluate boxes and segmentation masks with intersection over union (IoU):



$$\text{IoU (kite)} = \frac{\text{area}(\text{intersection})}{\text{area}(\text{union})}$$

In 3D: **Voxel IoU**

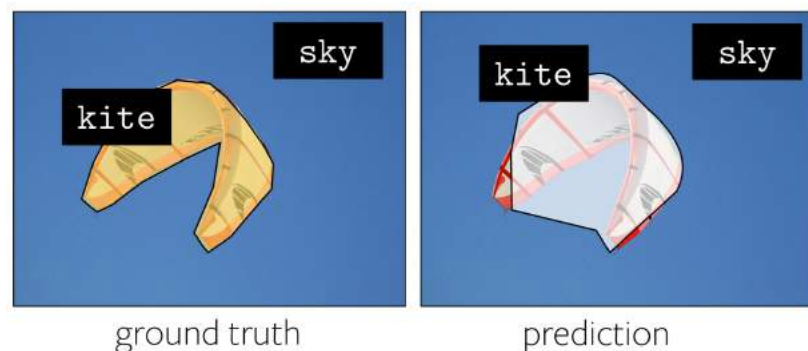
**Problem:** Cannot capture thin structures

**Problem:** Cannot be applied to pointclouds

**Problem:** For meshes, need to voxelize or sample

# Shape Comparison Metrics: Intersection over Union

In 2D, we evaluate boxes and segmentation masks with intersection over union (IoU):



$$\text{IoU (kite)} = \frac{\text{area}(\text{intersection})}{\text{area}(\text{union})}$$

In 3D: **Voxel IoU**

**Problem:** Cannot capture thin structures

**Problem:** Cannot be applied to pointclouds

**Problem:** For meshes, need to voxelize or sample

**Problem:** Not very meaningful at low values!

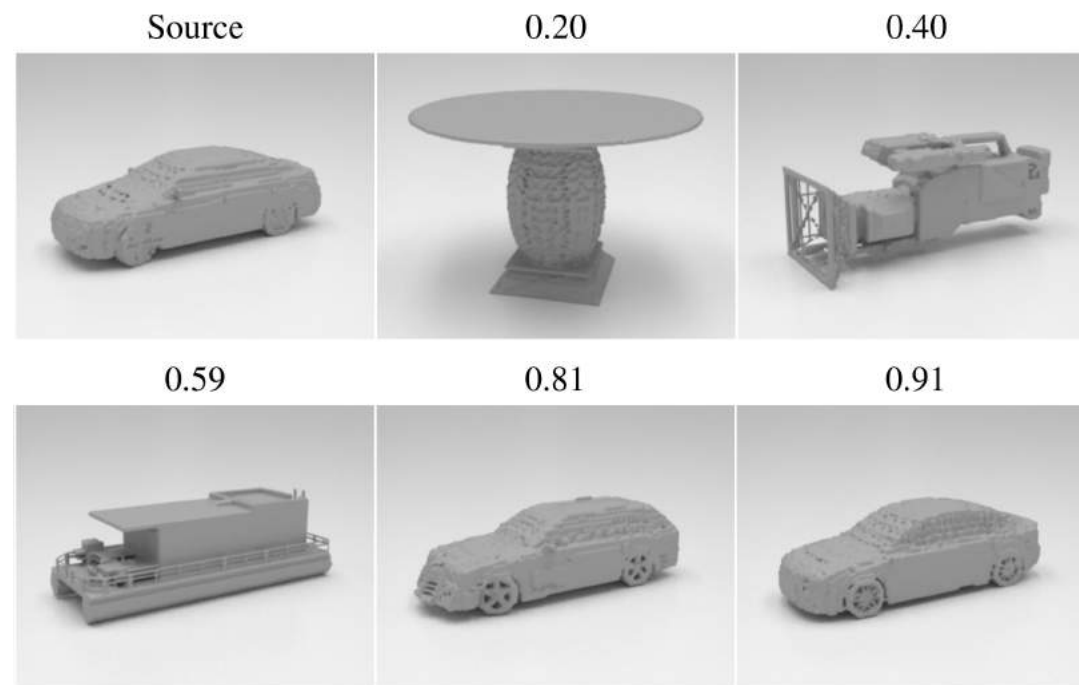


Figure credit: Alexander Kirillov

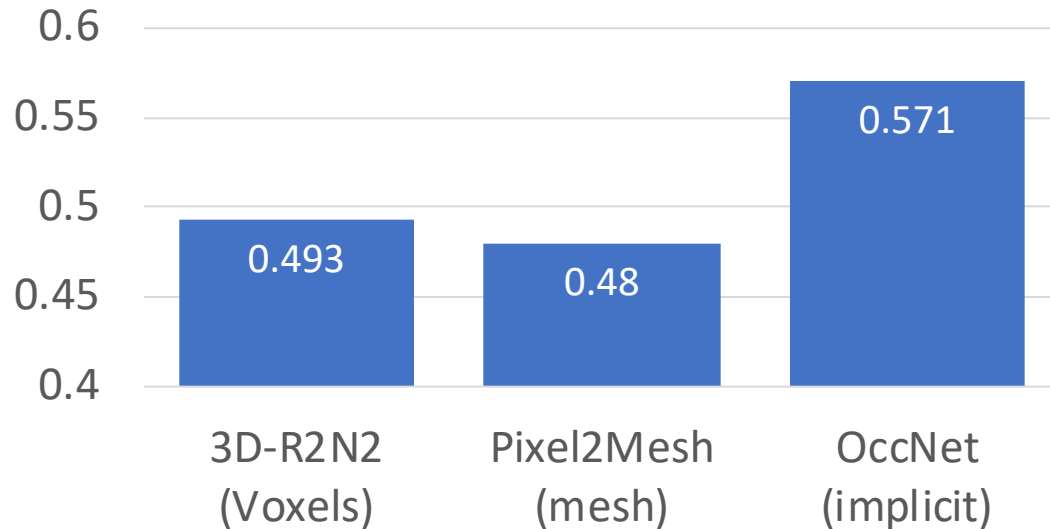
Figure credit: Tatarchenko et al, "What Do Single-view 3D Reconstruction Networks Learn?", CVPR 2019



# Shape Comparison Metrics: Intersection over Union

State-of-the-art methods  
achieve low IoU

IoU



Results from Mescheder et al, "Occupancy Networks: Learning 3D Reconstruction in Function Space", CVPR 2019

**Conclusion:** Voxel IoU not a good metric

In 3D: **Voxel IoU**

**Problem:** Cannot capture thin structures

**Problem:** Cannot be applied to pointclouds

**Problem:** For meshes, need to voxelize or sample

**Problem:** Not very meaningful at low values!

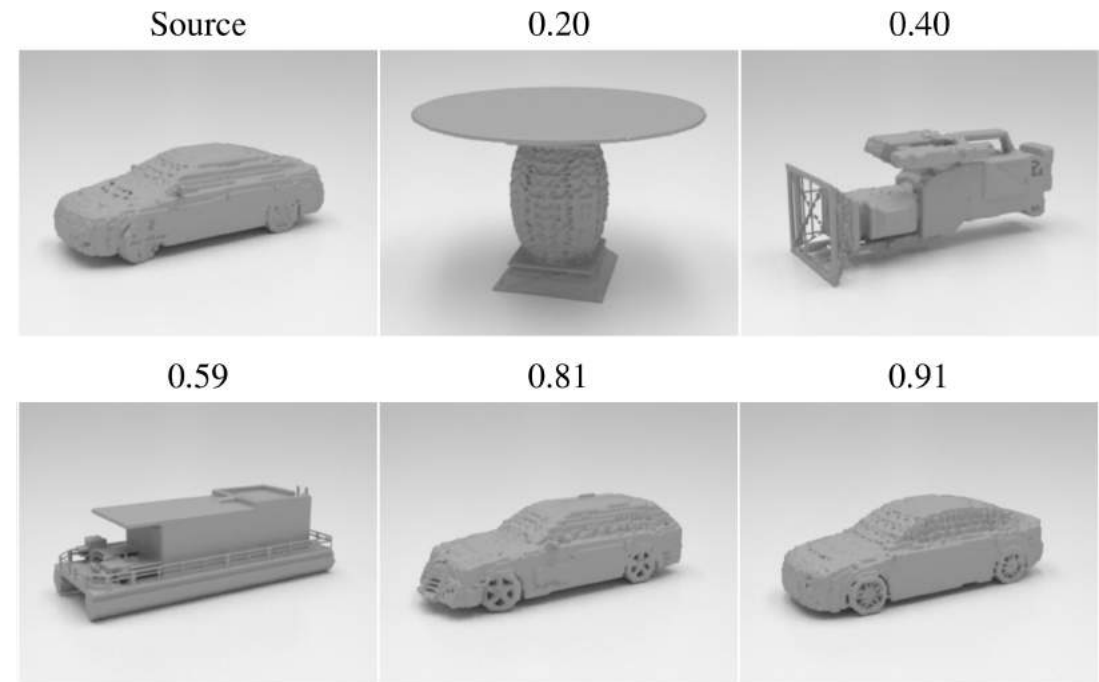


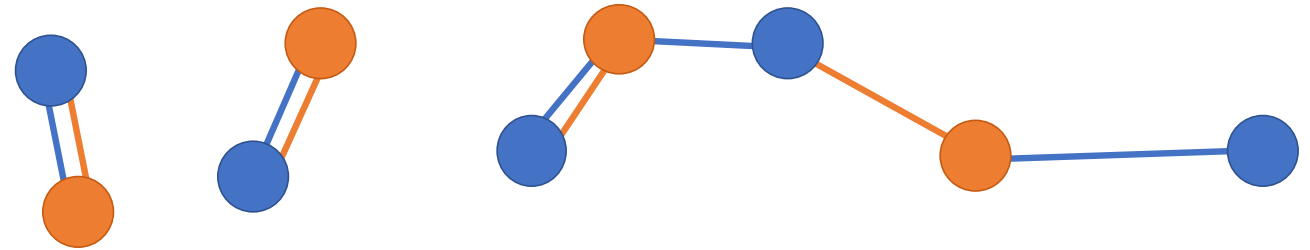
Figure credit: Tatarchenko et al, "What Do Single-view 3D Reconstruction Networks Learn?", CVPR 2019

# Shape Comparison Metrics: Chamfer Distance

We've already seen another shape comparison metric:  
**Chamfer distance**

1. Convert your prediction and ground-truth into pointclouds via sampling
2. Compare with Chamfer distance

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$



# Shape Comparison Metrics: Chamfer Distance

We've already seen another shape comparison metric:

## Chamfer distance

1. Convert your prediction and ground-truth into pointclouds via sampling
2. Compare with Chamfer distance

**Problem:** Chamfer is very sensitive to outliers

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2$$

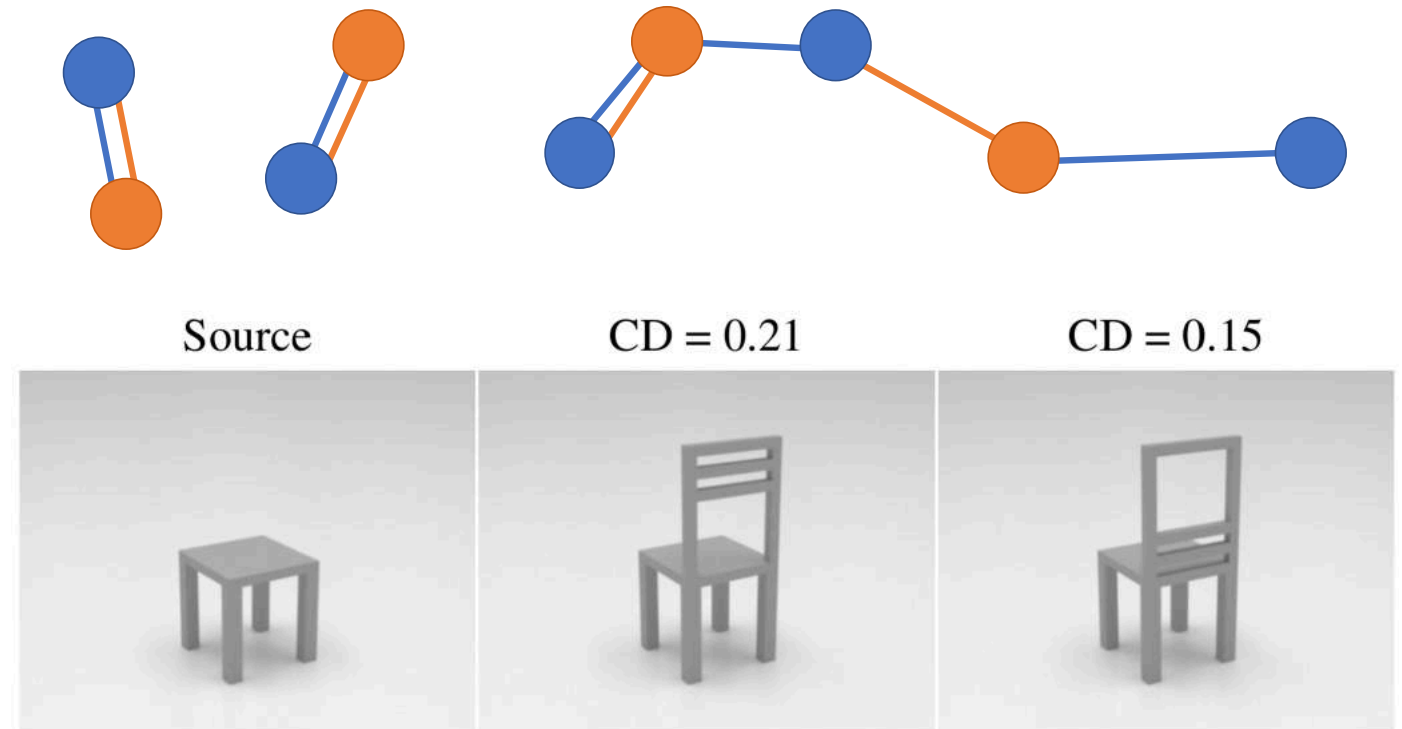
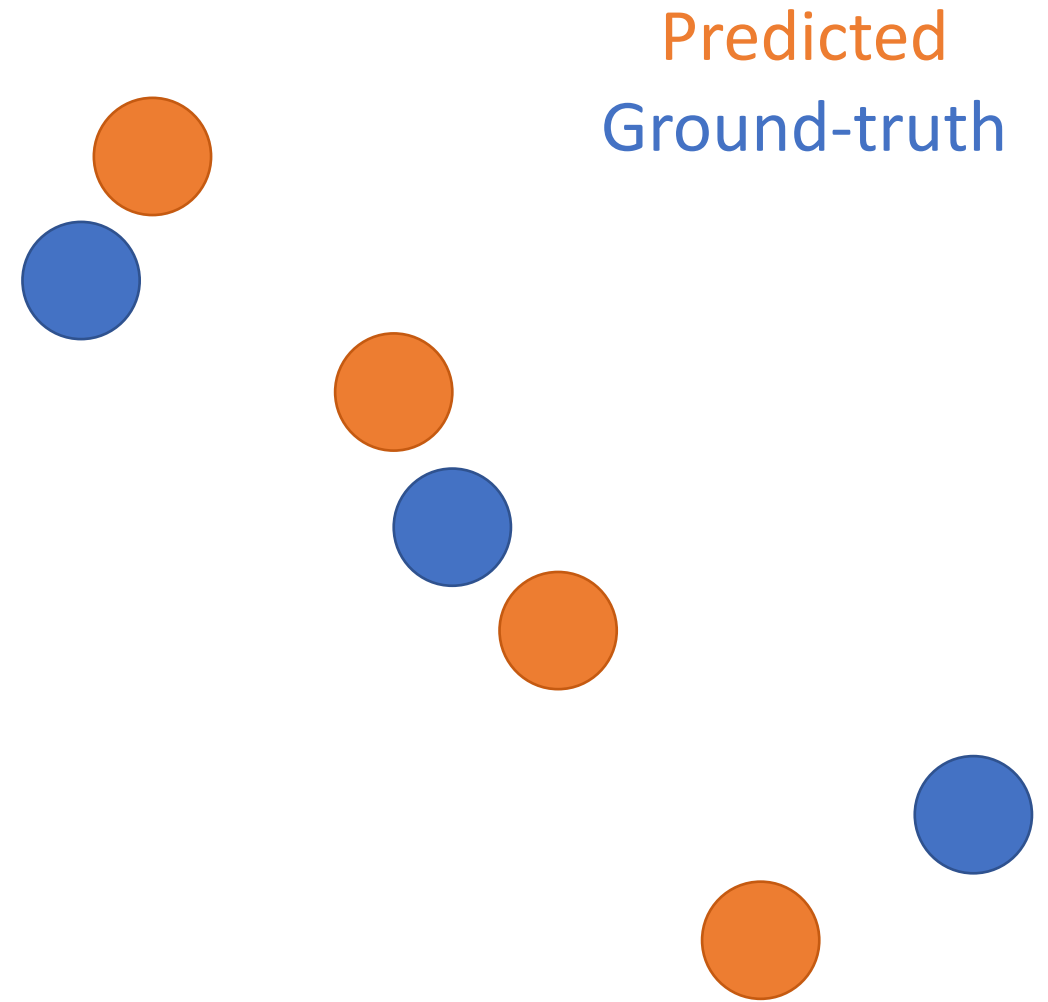


Figure credit: Tatarchenko et al, "What Do Single-view 3D Reconstruction Networks Learn?", CVPR 2019

# Shape Comparison Metrics: F1 Score

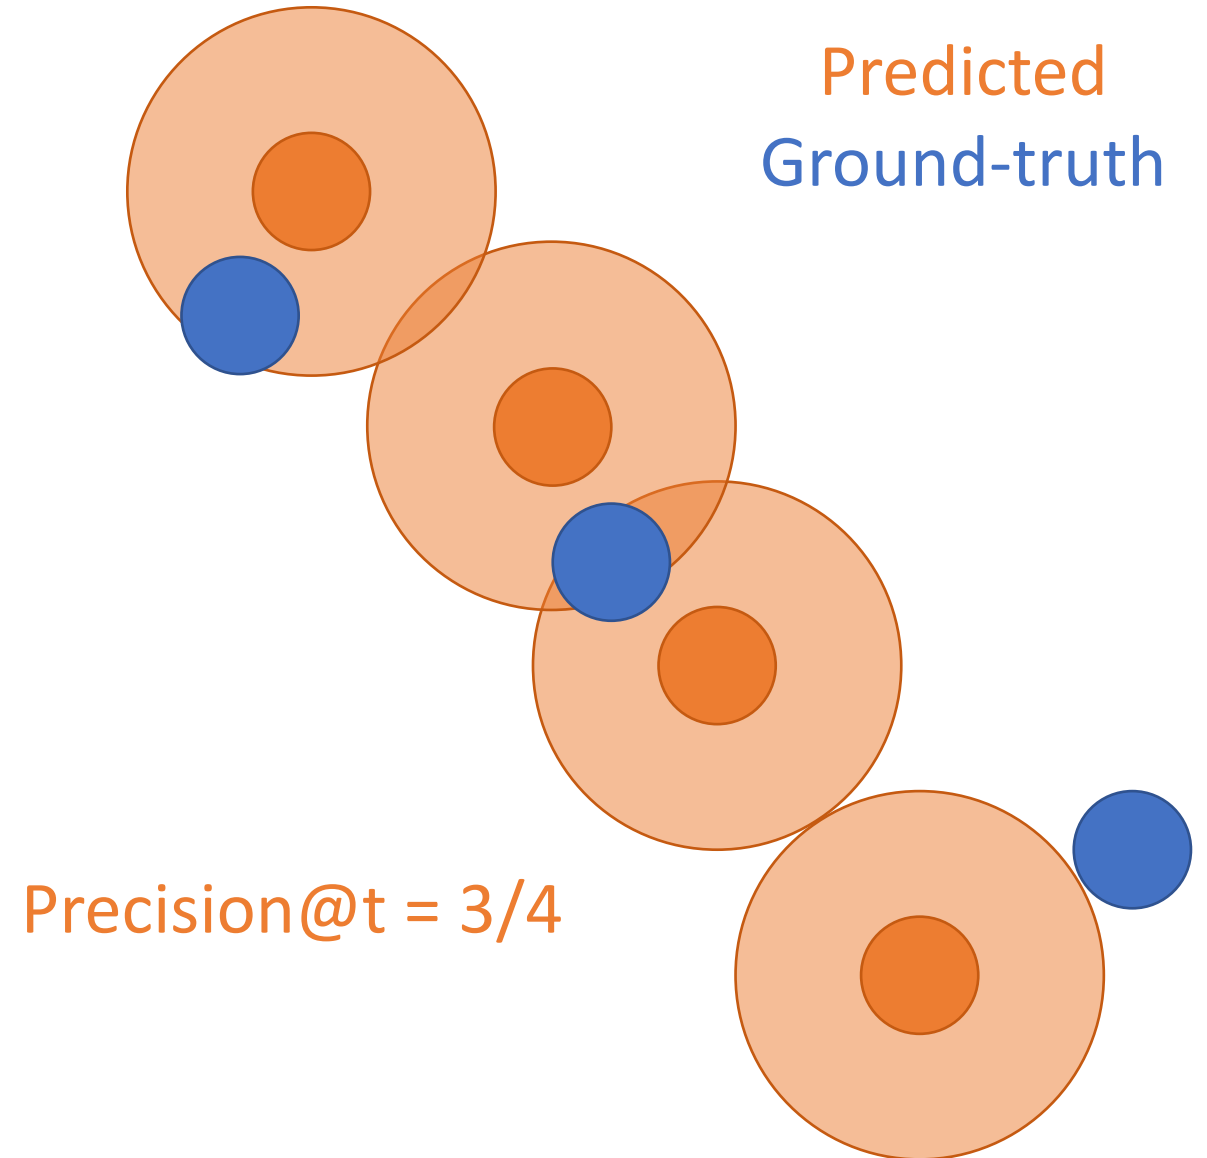
Similar to Chamfer, sample points from the surface of the prediction and the ground-truth



# Shape Comparison Metrics: F1 Score

Similar to Chamfer, sample points from the surface of the prediction and the ground-truth

Precision@t = fraction of predicted points within t of some ground-truth point

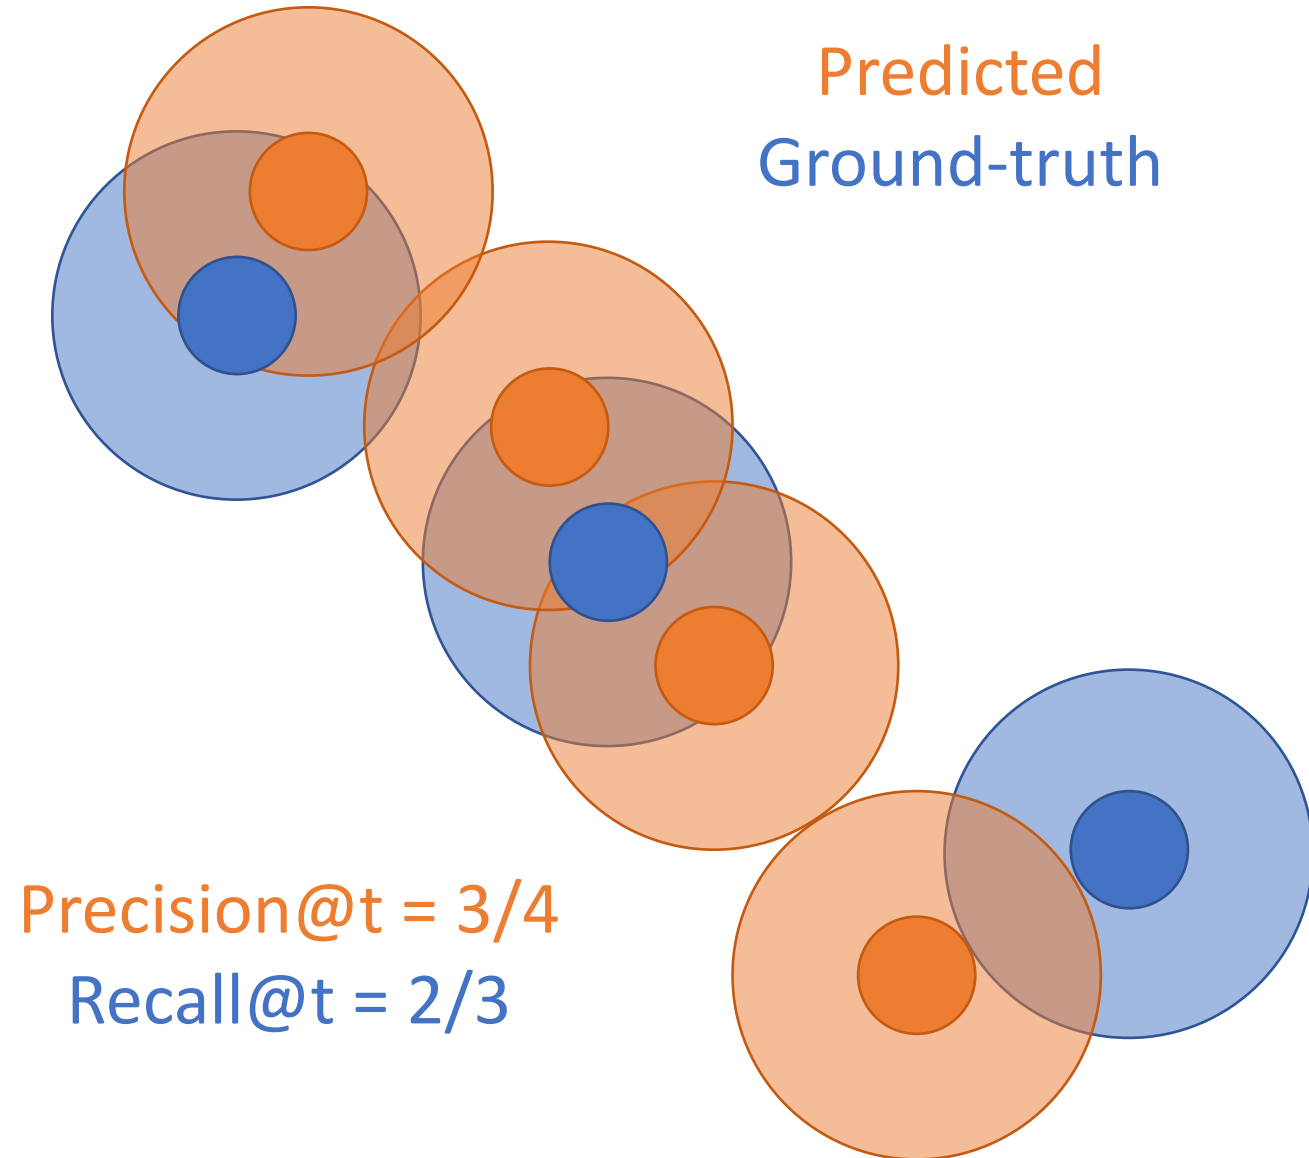


# Shape Comparison Metrics: F1 Score

Similar to Chamfer, sample points from the surface of the prediction and the ground-truth

Precision@t = fraction of predicted points within t of some ground-truth point

Recall@t = fraction of ground-truth points within t of some predicted point



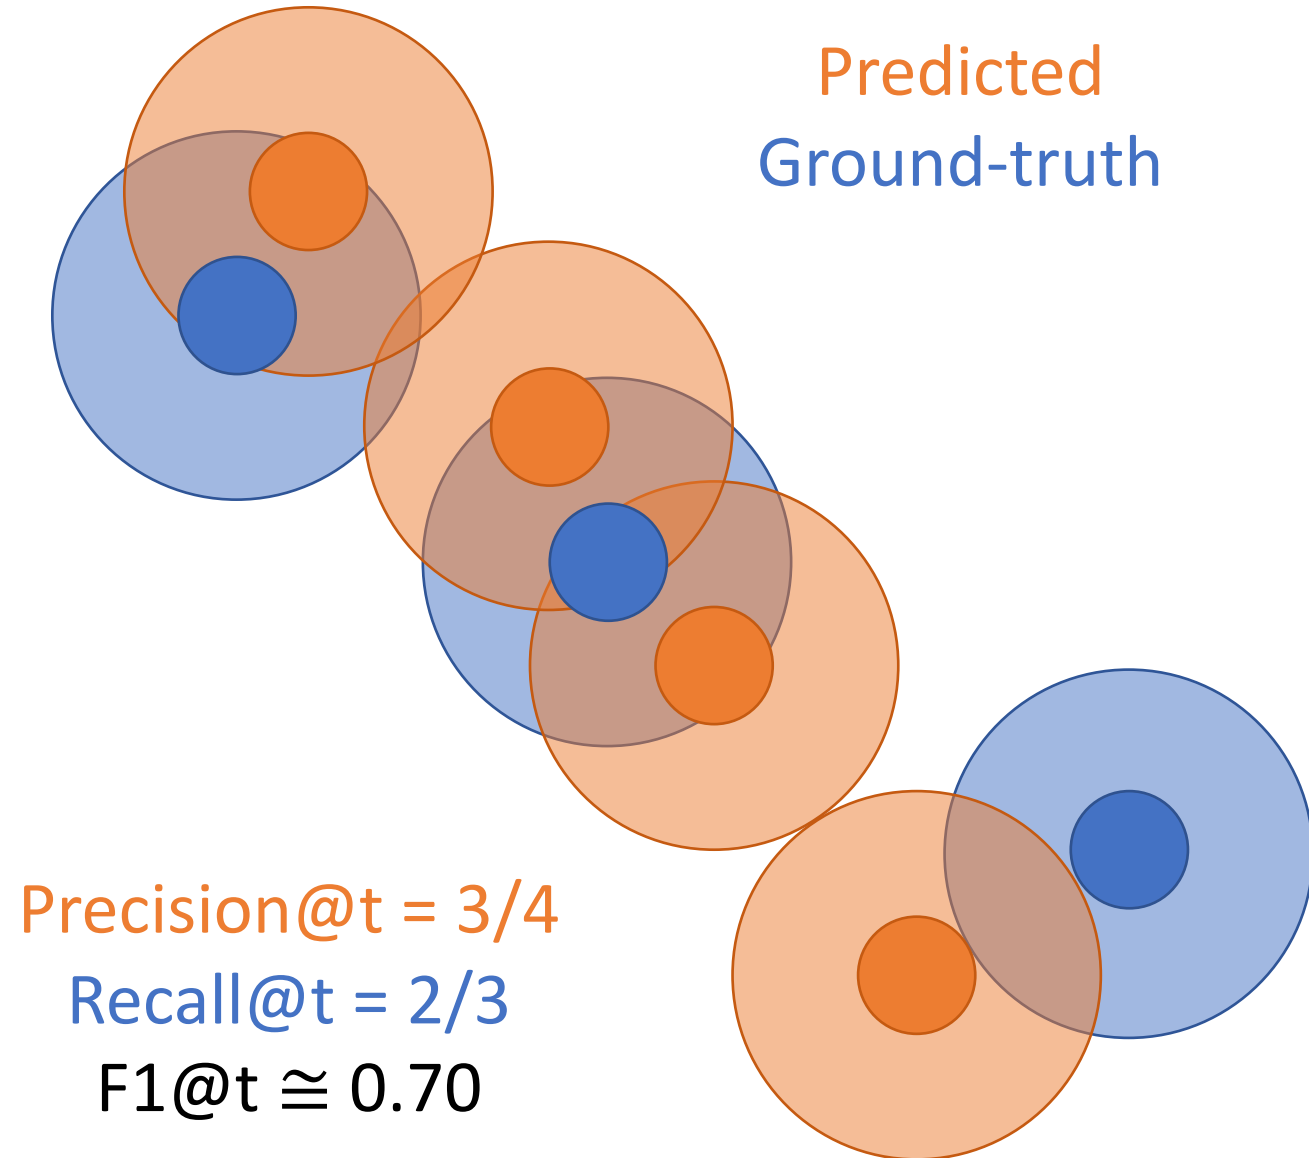
# Shape Comparison Metrics: F1 Score

Similar to Chamfer, sample points from the surface of the prediction and the ground-truth

Precision@t = fraction of predicted points within t of some ground-truth point

Recall@t = fraction of ground-truth points within t of some predicted point

$$F1@t = 2 * \frac{Precision@t * Recall@t}{Precision@t + Recall@t}$$



# Shape Comparison Metrics: F1 Score

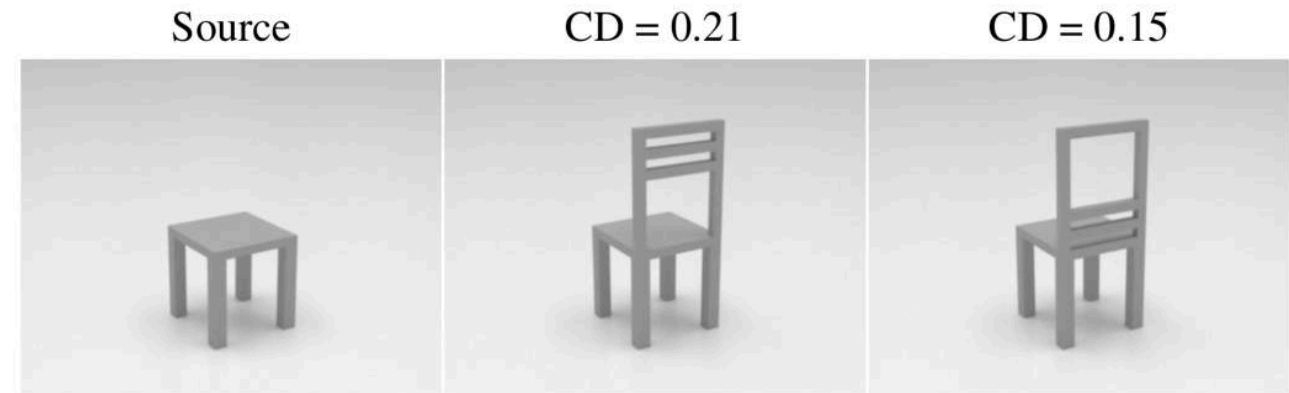
Similar to Chamfer, sample points from the surface of the prediction and the ground-truth

Precision@t = fraction of predicted points within t of some ground-truth point

Recall@t = fraction of ground-truth points within t of some predicted point

$$F1@t = 2 * \frac{Precision@t * Recall@t}{Precision@t + Recall@t}$$

F1 score is robust to outliers!



**Conclusion:** F1 score is probably the best shape prediction metric in common use

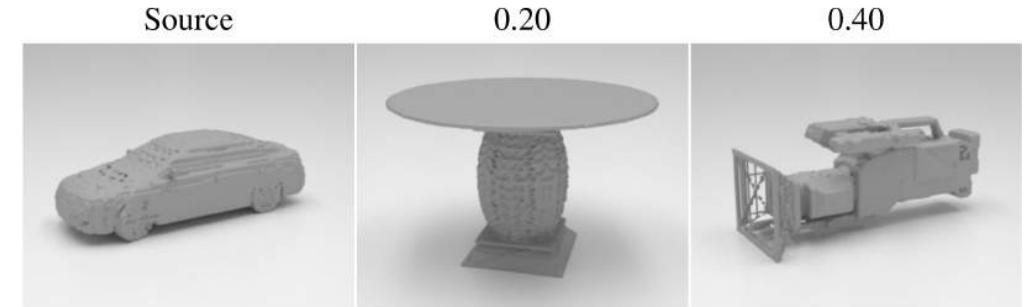
Figure credit: Tatarchenko et al, "What Do Single-view 3D Reconstruction Networks Learn?", CVPR 2019



# Shape Comparison Metrics: Summary

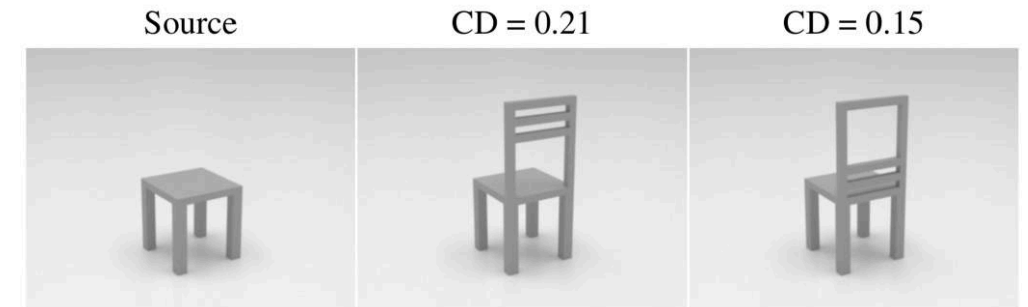
## Intersection over Union:

Doesn't capture fine structure,  
not meaningful at low values



## Chamfer Distance:

Very sensitive to outliers  
Can be directly optimized



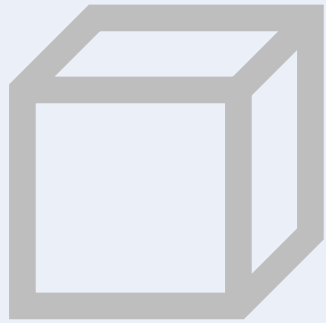
## F1 score:

Robust to outliers, but need to  
look at different threshold values  
to capture details at different scales



# 3D Shape Prediction

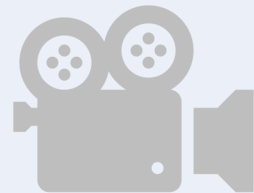
Shape Representations



Metrics



Camera Systems

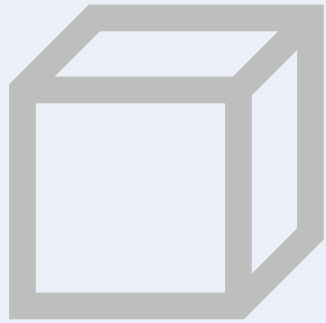


Datasets



# 3D Shape Prediction

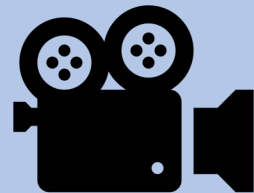
Shape Representations



Metrics



Camera Systems

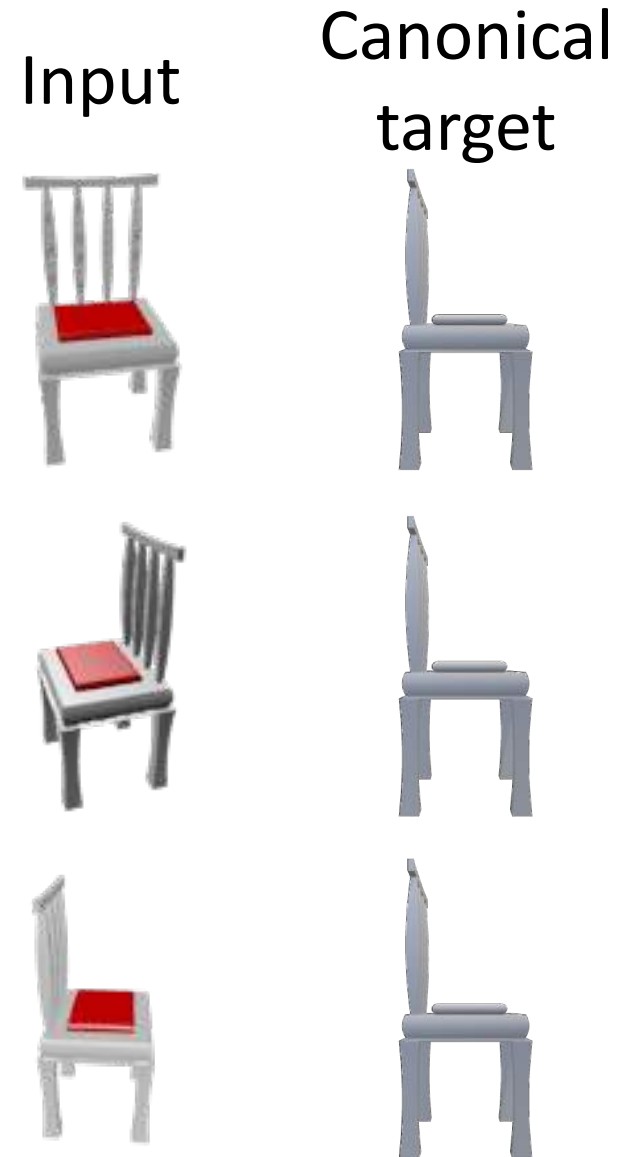


Datasets



# Cameras: Canonical vs View Coordinates

**Canonical Coordinates:** Predict 3D shape in a canonical coordinate system (e.g. front of chair is +z) regardless of the viewpoint of the input image

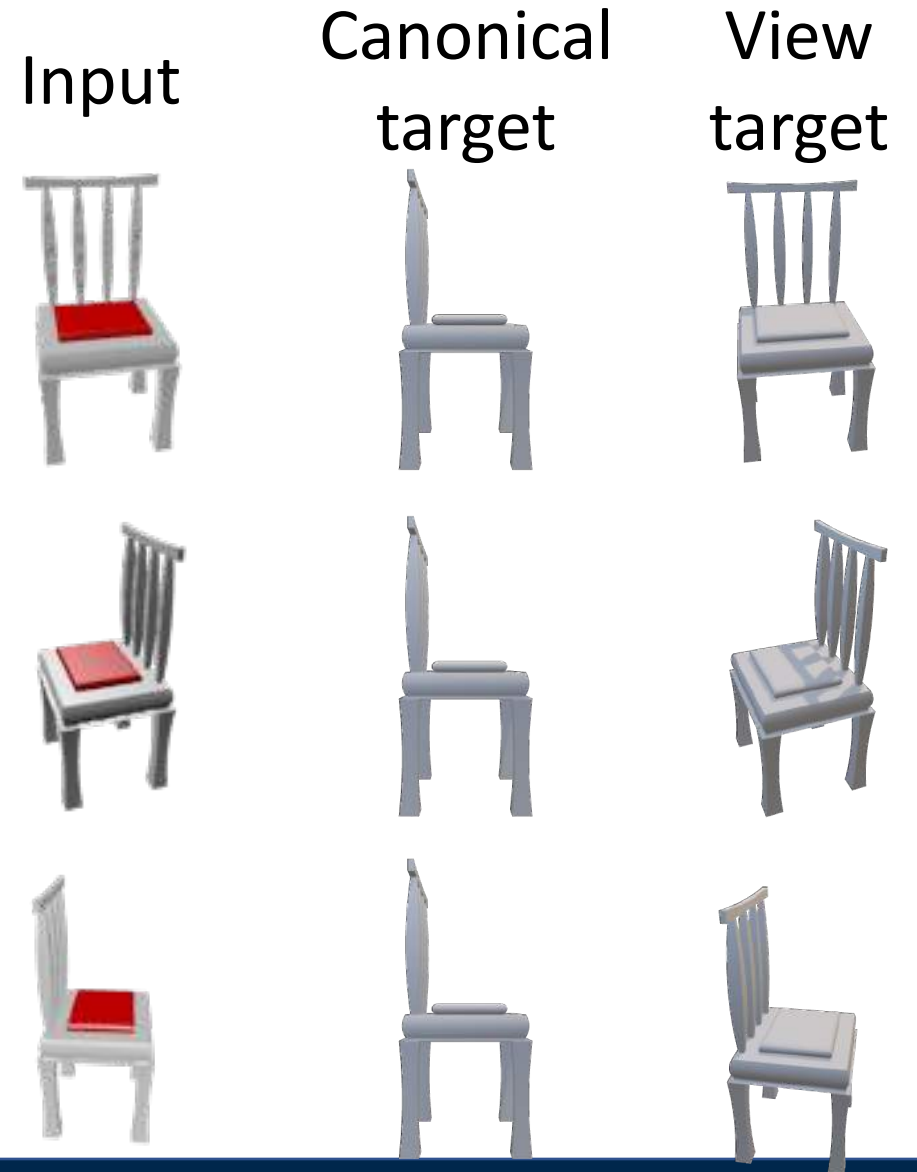


# Cameras: Canonical vs View Coordinates

**Canonical Coordinates:** Predict 3D shape in a canonical coordinate system (e.g. front of chair is +z) regardless of the viewpoint of the input image

**View Coordinates:** Predict 3D shape aligned to the viewpoint of the camera

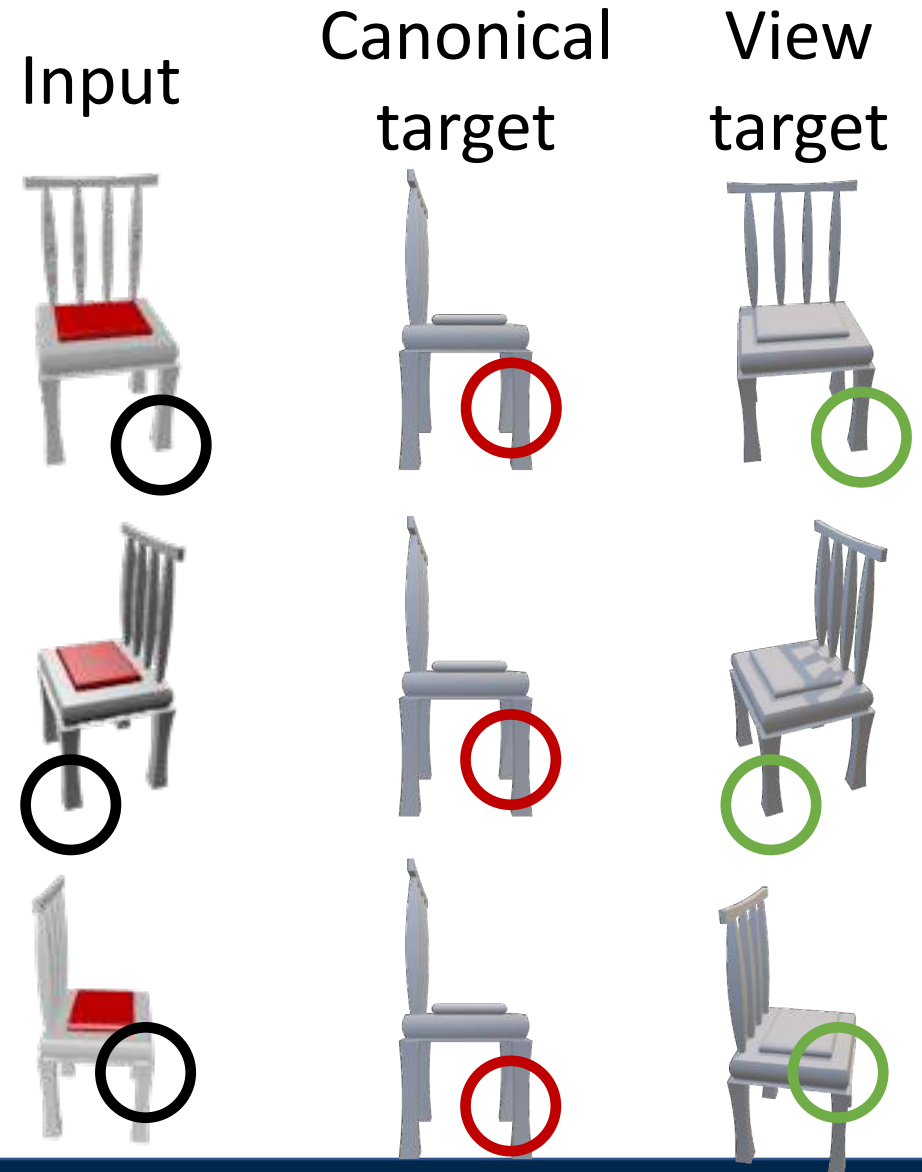
Many papers predict in canonical coordinates – easier to load data



# Cameras: Canonical vs View Coordinates

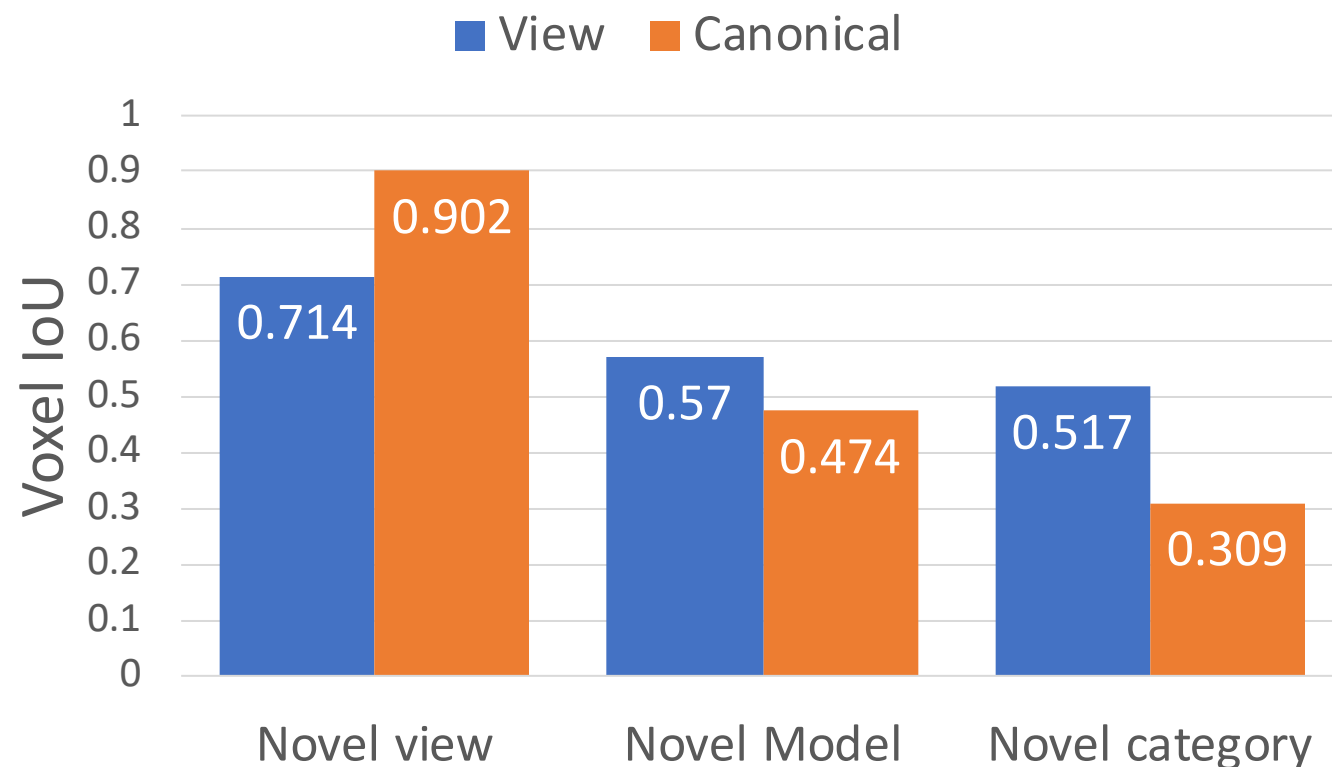
**Problem:** Canonical view breaks the “principle of feature alignment”: Predictions should be aligned to inputs

View coordinates maintain alignment between inputs and predictions!

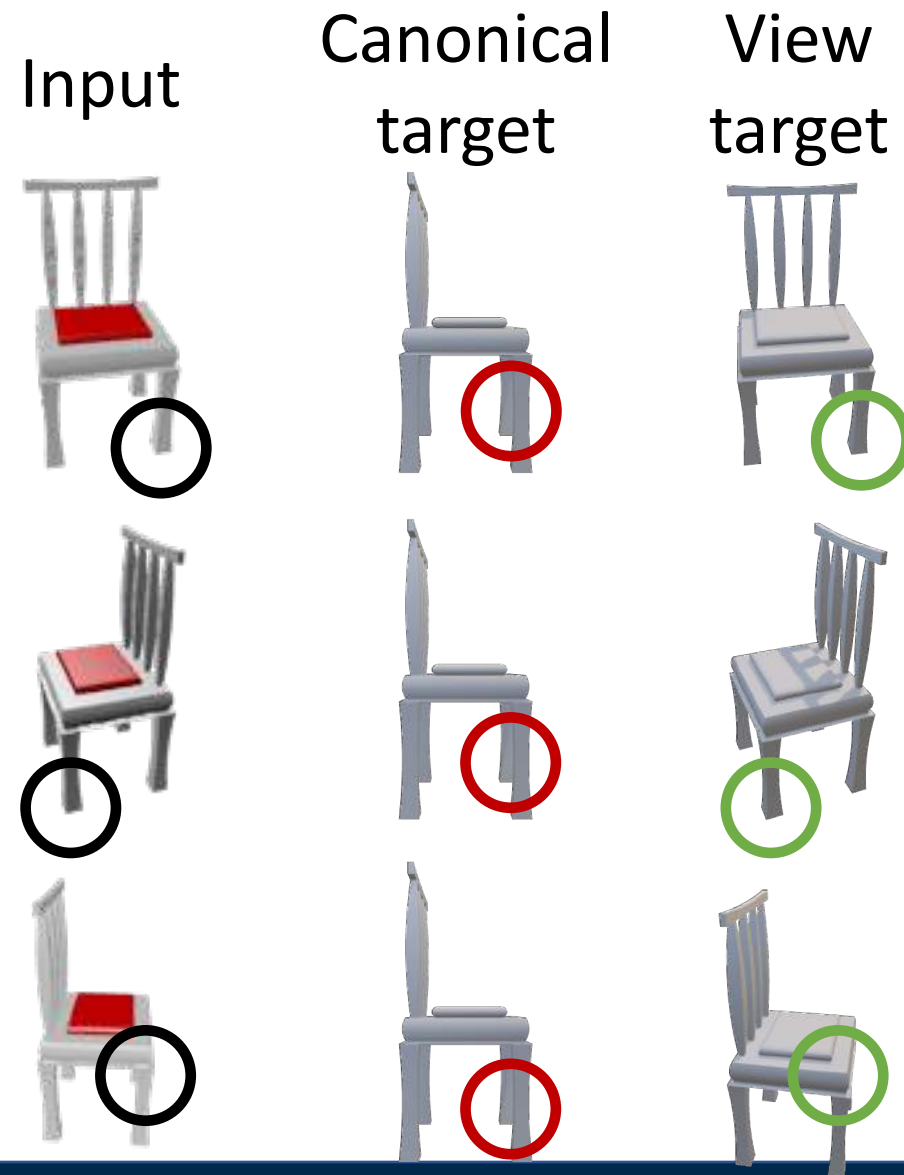


# Cameras: Canonical vs View Coordinates

**Problem:** Canonical view overfits to training shapes:  
Better generalization to new views of known shapes  
Worse generalization to new shapes or new categories



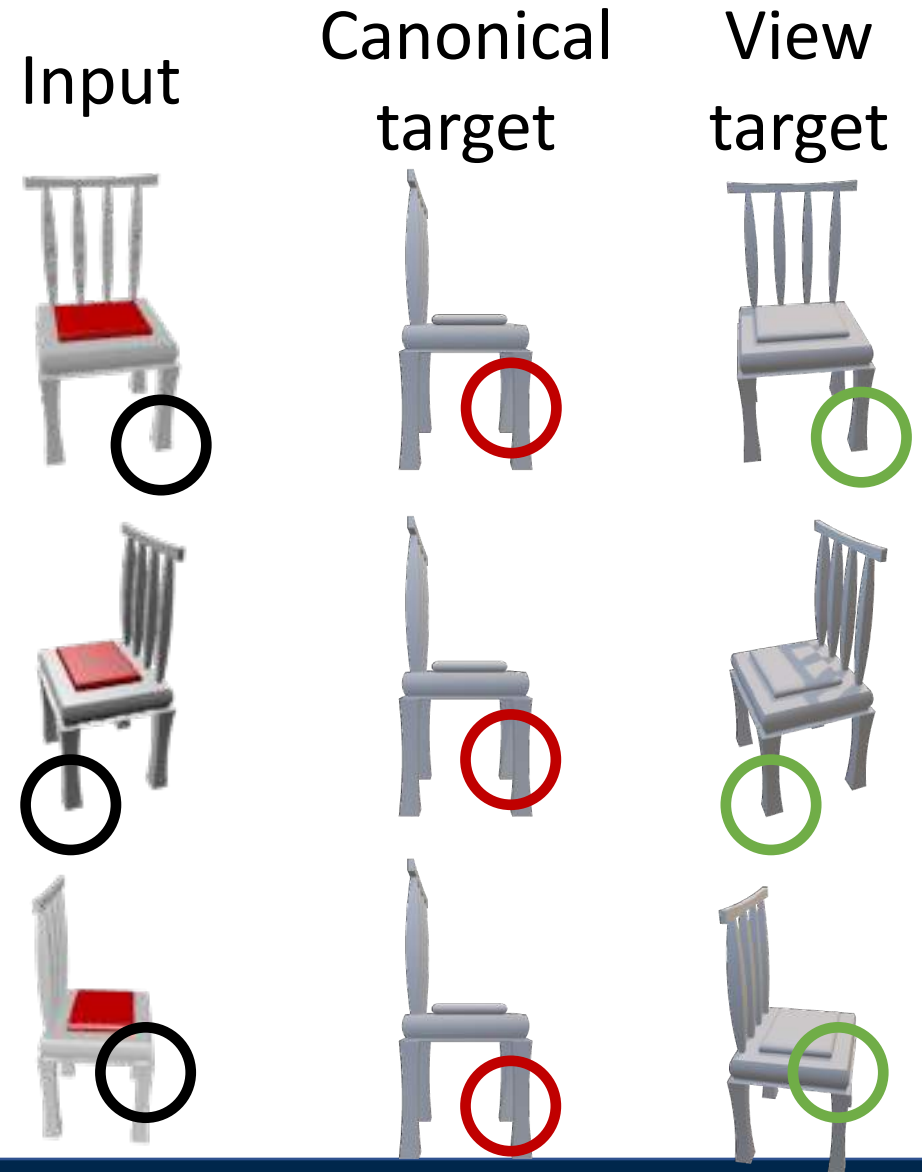
Shin et al, "Pixels, voxels, and views: A study of shape representations for single view 3D object shape prediction", CVPR 2018



# Cameras: Canonical vs View Coordinates

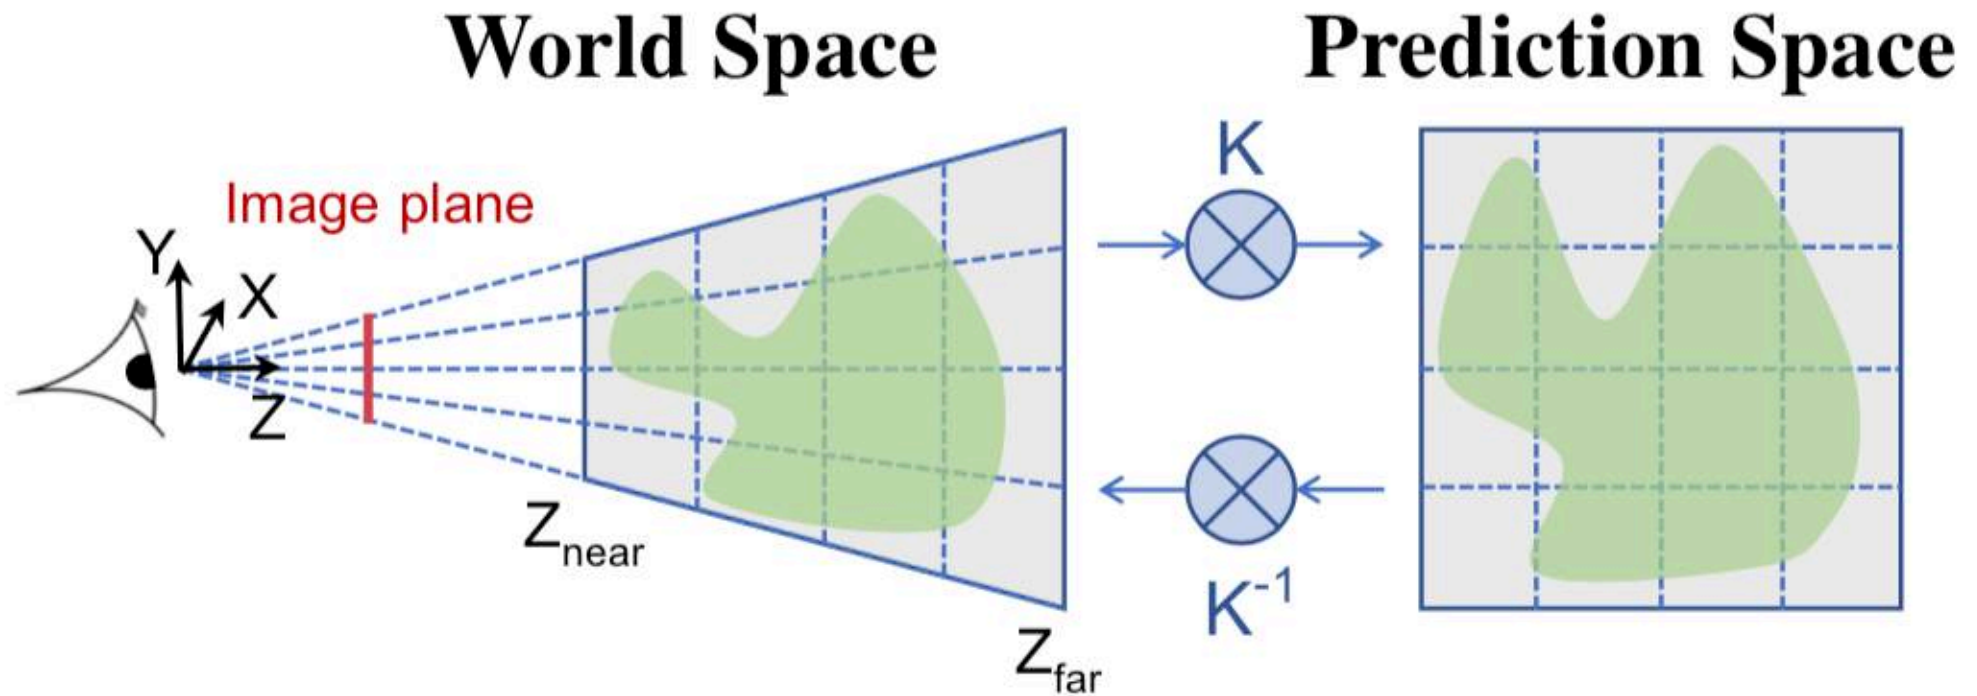
**Problem:** Canonical view overfits to training shapes:  
Better generalization to new views of known shapes  
Worse generalization to new shapes or new categories

**Conclusion:** Prefer view coordinate system





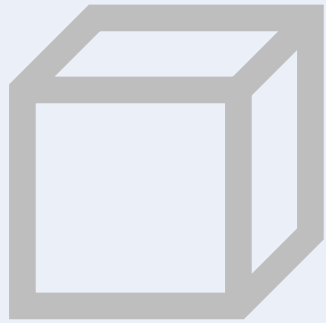
# View-Centric Voxel Predictions



**View-centric** predictions! Voxels take perspective camera into account, so our “voxels” are actually frustums

# 3D Shape Prediction

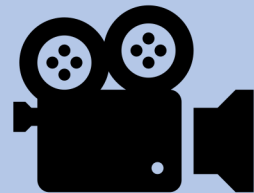
Shape Representations



Metrics



Camera Systems

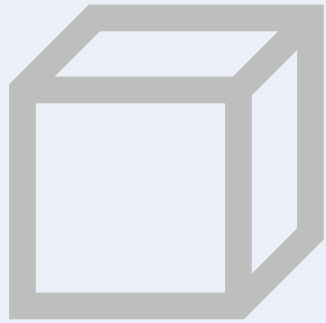


Datasets



# 3D Shape Prediction

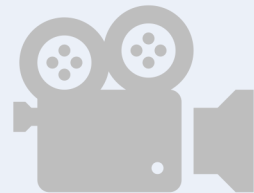
Shape Representations



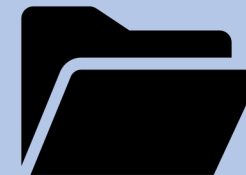
Metrics



Camera Systems



Datasets



# 3D Datasets: Object-Centric

## ShapeNet



~50 categories, ~50k 3D CAD models

Standard split has 13 categories, ~44k models, 25 rendered images per model

Many papers show results here

(-) Synthetic, isolated objects; no context

(-) Lots of chairs, cars, airplanes

Chang et al, "ShapeNet: An Information-Rich 3D Model Repository", arXiv 2015

Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016

# 3D Datasets: Object-Centric

## ShapeNet



~50 categories, ~50k 3D CAD models

Standard split has 13 categories, ~44k models, 25 rendered images per model

Many papers show results here

(-) Synthetic, isolated objects; no context

(-) Lots of chairs, cars, airplanes

Chang et al, "ShapeNet: An Information-Rich 3D Model Repository", arXiv 2015  
Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016

## Pix3D



9 categories, 219 3D models of IKEA furniture aligned to ~17k real images

Some papers train on ShapeNet and show qualitative results here, but use ground-truth segmentation masks

(+) Real images! Context!

(-) Small, partial annotations – only 1 obj/image

Sun et al, "Pix3D: Dataset and Methods for Single-Image 3D Shape Modeling", CVPR 2018

# 3D Shape Prediction: Mesh R-CNN

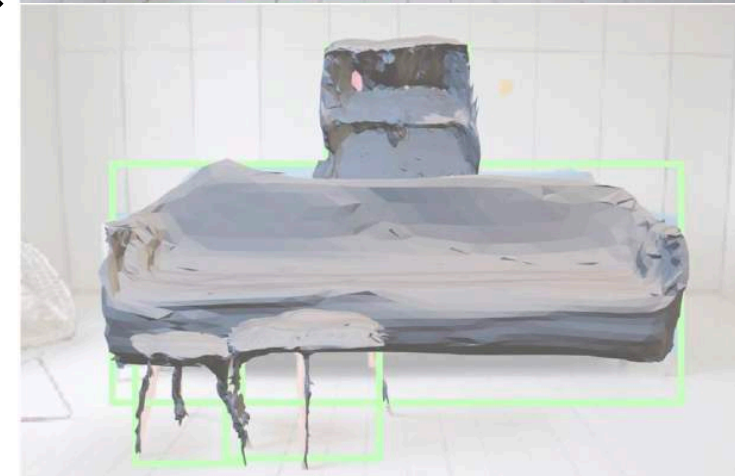
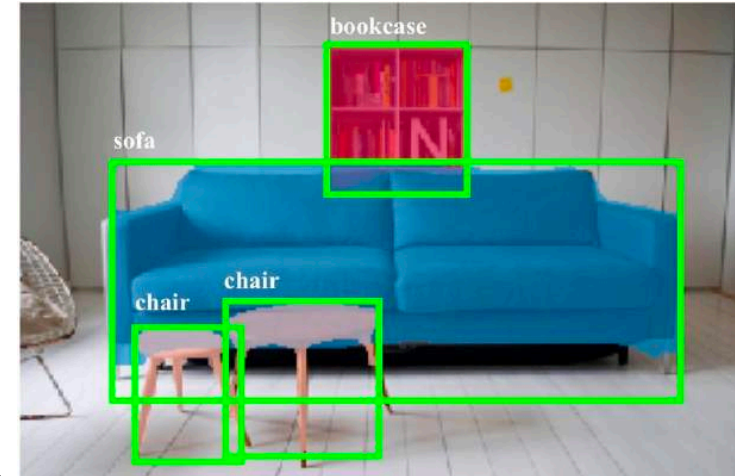
Mask R-CNN:

2D Image -> 2D shapes



**Mesh R-CNN:**

2D Image -> **Triangle Meshes**



He, Gkioxari, Dollár, and Girshick, "Mask R-CNN", ICCV 2017

Gkioxari, Malik, and Johnson, "Mesh R-CNN", ICCV 2019



# Mesh R-CNN: Task

**Input:** Single RGB image

**Output:**

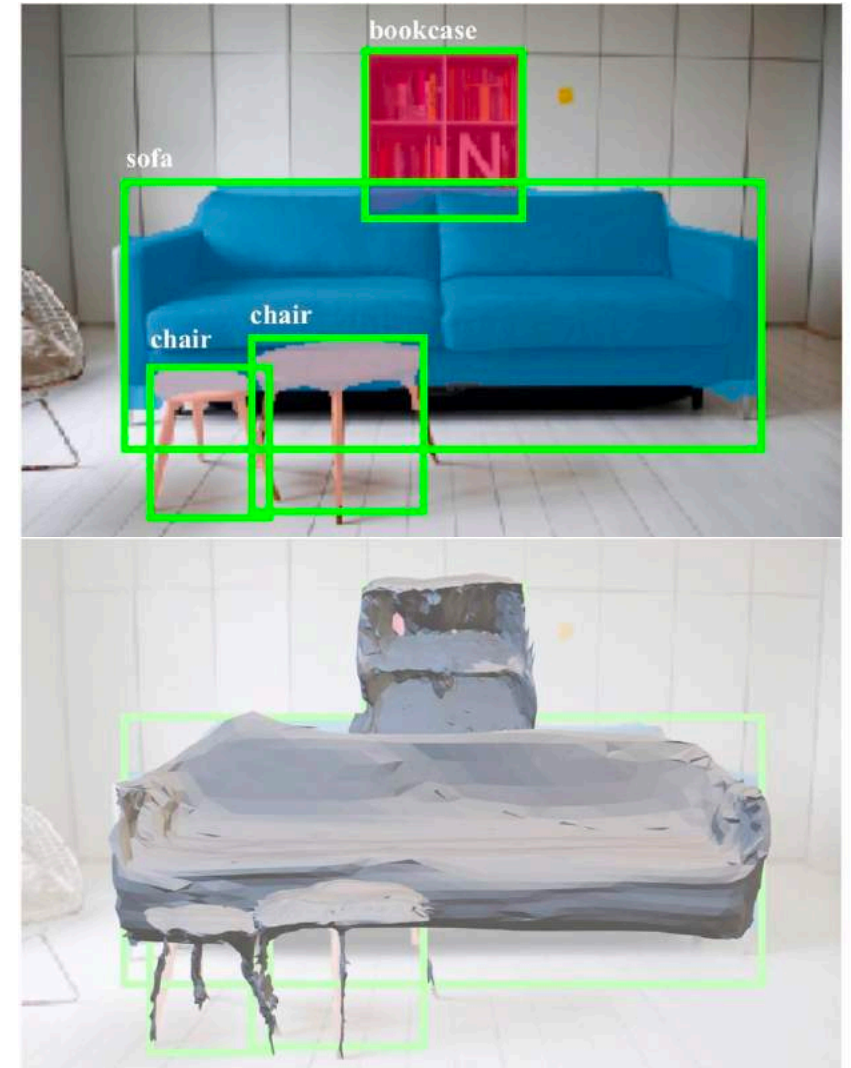
A set of detected objects

For each object:

- Bounding box
- Category label
- Instance segmentation
- 3D triangle mesh

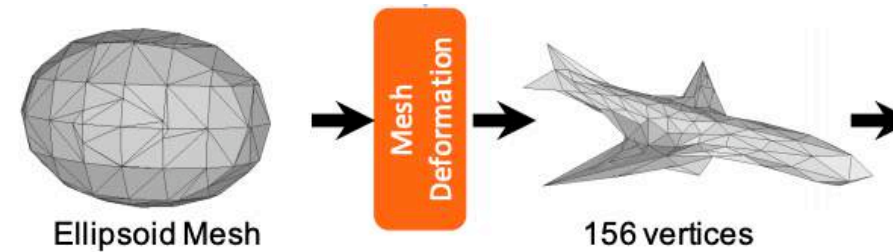
Mask R-CNN

Mesh head



# Mesh R-CNN: Hybrid 3D shape representation

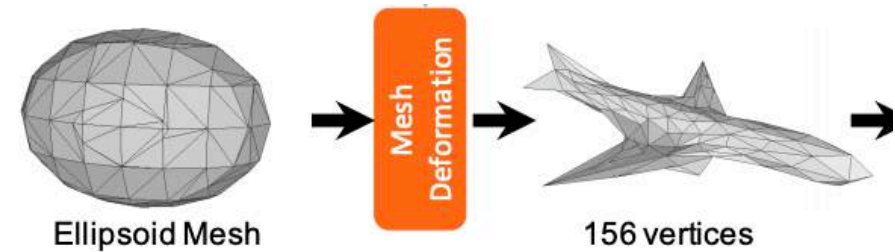
**Mesh deformation** gives good results, but the topology (verts, faces, genus, connected components) fixed by the initial mesh



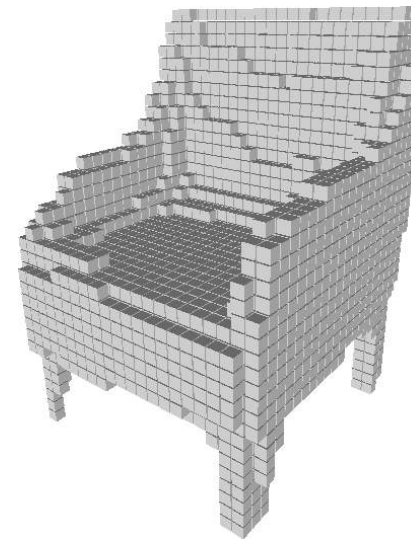


# Mesh R-CNN: Hybrid 3D shape representation

**Mesh deformation** gives good results, but the topology (verts, faces, genus, connected components) fixed by the initial mesh



**Our approach:** Use voxel predictions to create initial mesh prediction!



# Mesh R-CNN Pipeline

Input image



# Mesh R-CNN Pipeline

Input image



2D object recognition



# Mesh R-CNN Pipeline

Input image



2D object recognition



3D object voxels

# Mesh R-CNN Pipeline

Input image



2D object recognition



3D object meshes

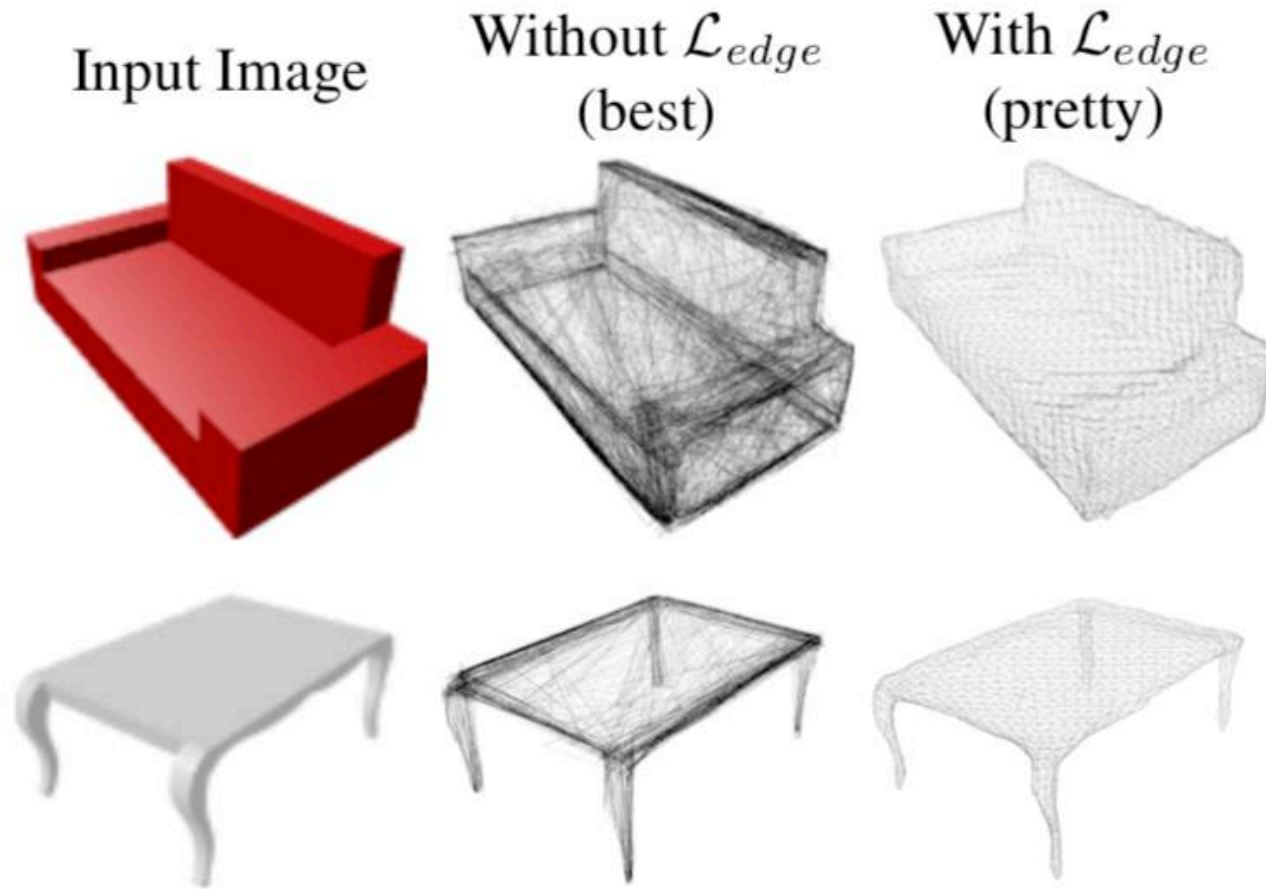
3D object voxels

# Mesh R-CNN: ShapeNet Results



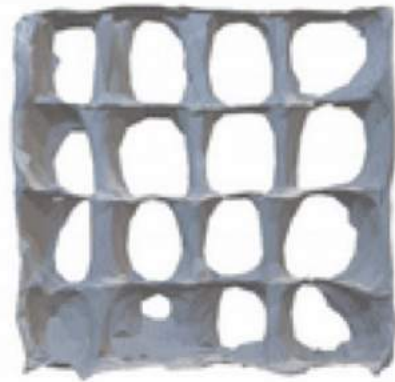


# Mesh R-CNN: Shape Regularizers



Using Chamfer as only mesh loss gives degenerate meshes. Also need "mesh regularizer" to encourage nice predictions:  
 $\mathcal{L}_{edge}$  = minimize L2 norm of edges in the predicted mesh

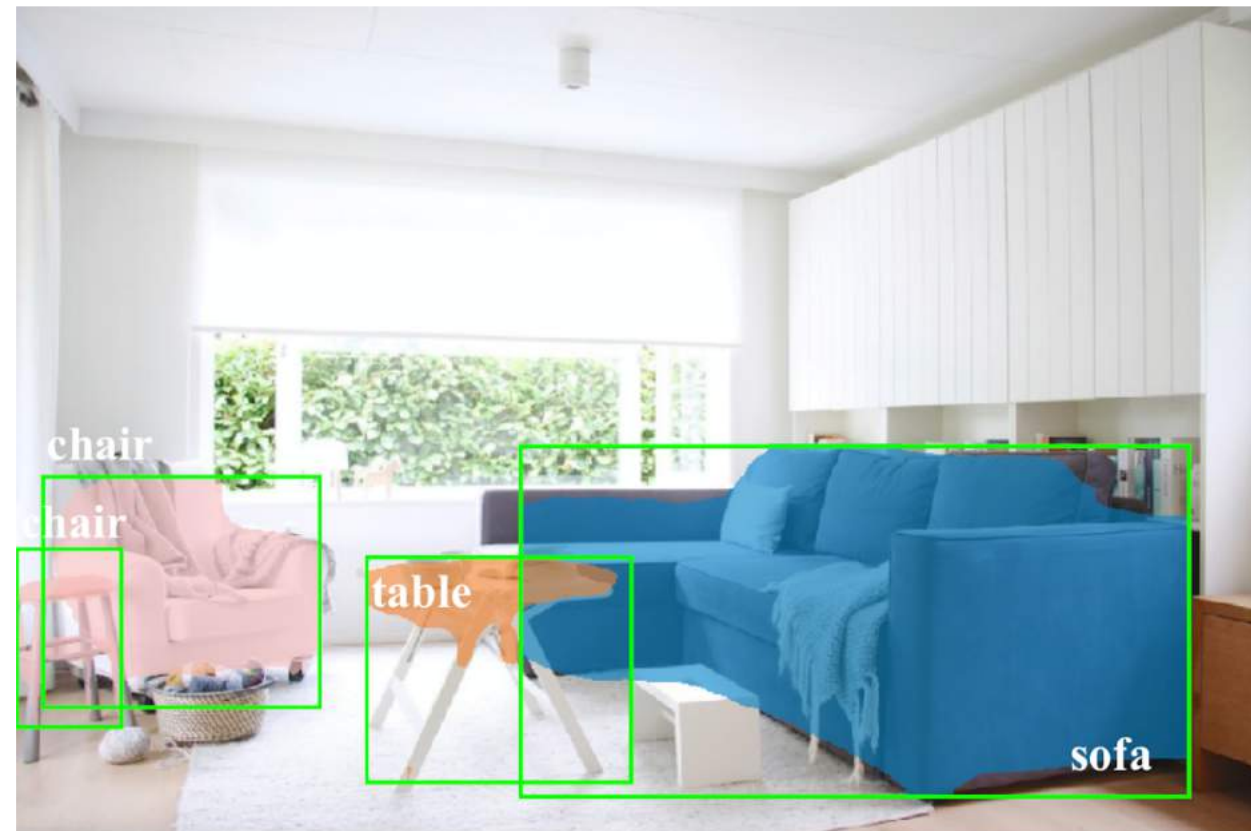
# Mesh R-CNN: Pix3D Results



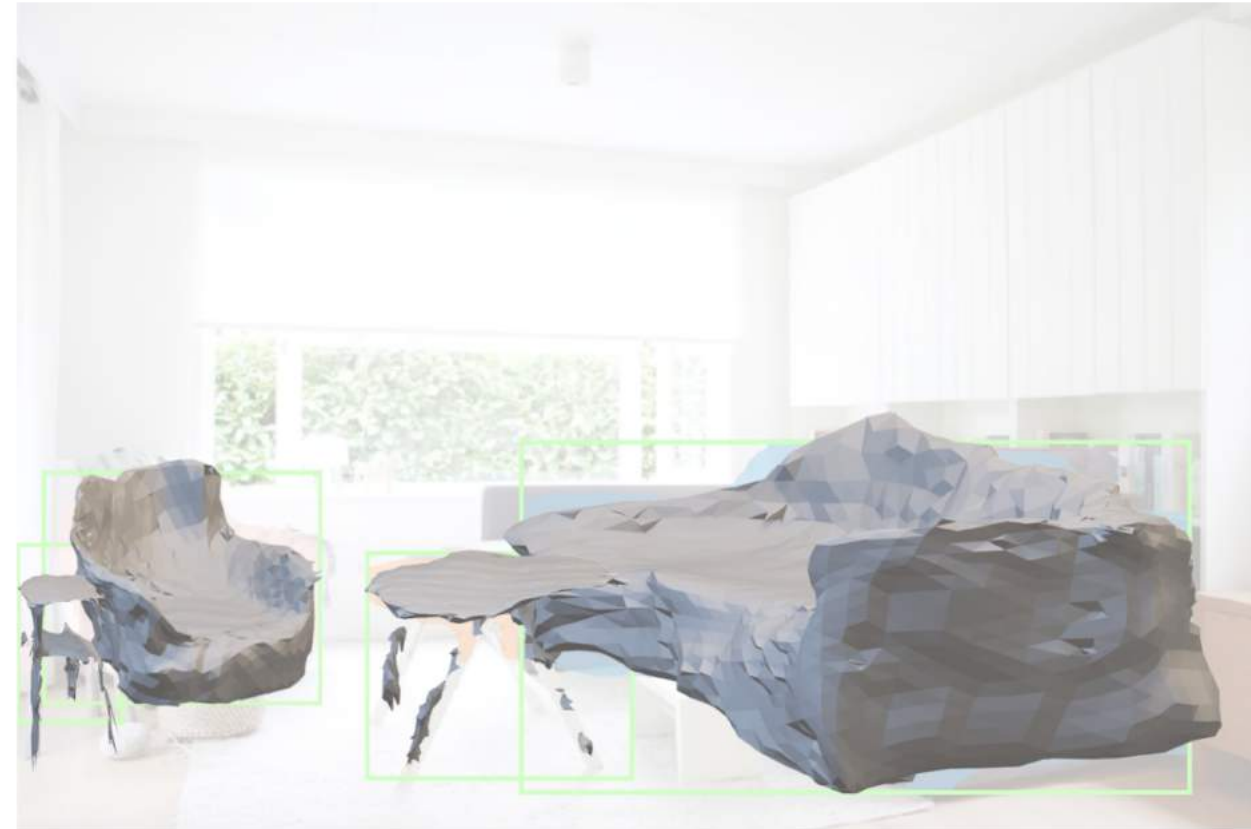


# Mesh R-CNN: Pix3D Results

Predicting many objects per scene



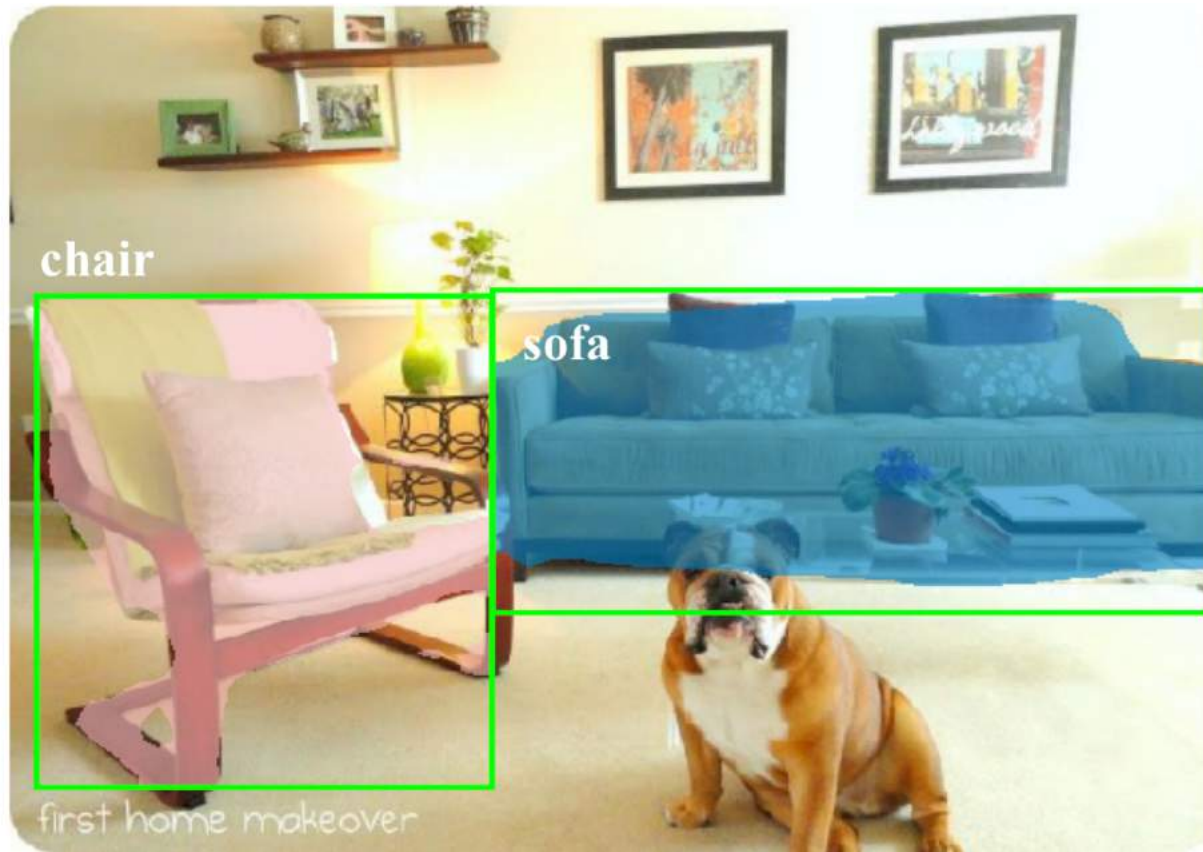
Box & Mask Predictions



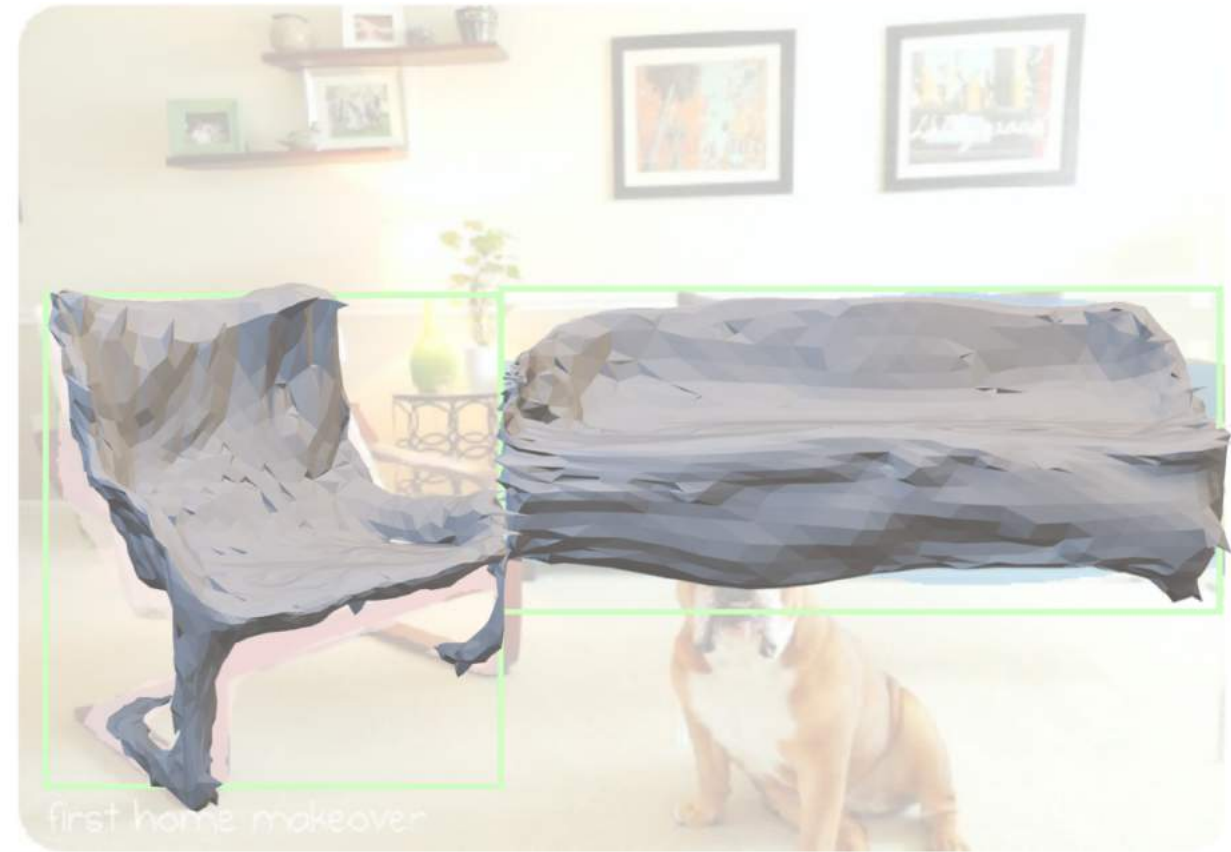
Mesh Predictions

# Mesh R-CNN: Pix3D Results

Amodal completion: predict occluded parts of objects



Box & Mask Predictions



Mesh Predictions

# Mesh R-CNN: Pix3D Results

Segmentation failures  
propagate to meshes



Box & Mask Predictions



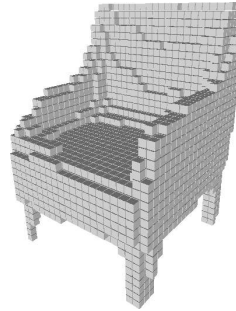
Mesh Predictions

# Recap

Predicting 3D Shapes  
from single image

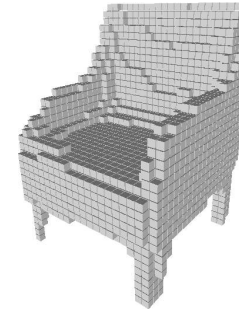


Input Image



3D Shape

Processing 3D  
input data

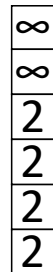


3D Shape

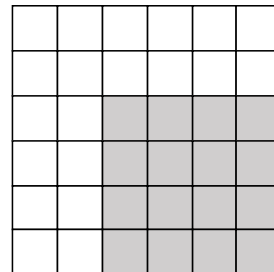


Chair

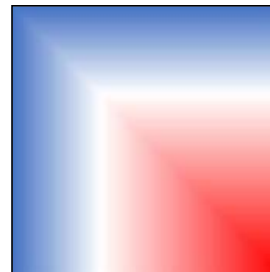
## 3D Shape Representations



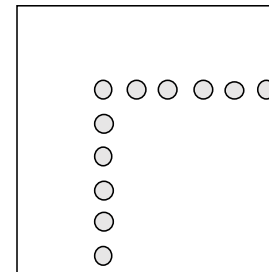
Depth  
Map



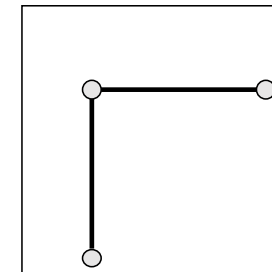
Voxel  
Grid



Implicit  
Surface



Pointcloud



Mesh

Next Time:  
Videos